

Subject Name Solutions

4351601 – Winter 2023

Semester 1 Study Material

Detailed Solutions and Explanations

Question 1(a) [3 marks]

Define the following terms: (1) Artificial Intelligence (2) Expert System.

Solution

Term	Definition
Artificial Intelligence	AI is a branch of computer science that creates machines capable of performing tasks that typically require human intelligence, such as learning, reasoning, and problem-solving.
Expert System	An expert system is a computer program that uses knowledge and inference rules to solve problems that normally require human expertise in a specific domain.

- **AI characteristics:** Learning, reasoning, perception
- **Expert system components:** Knowledge base, inference engine

Mnemonic

“AI Learns, Expert Advises”

Question 1(b) [4 marks]

Compare Biological Neural Network and Artificial Neural Network.

Solution

Aspect	Biological Neural Network	Artificial Neural Network
Processing	Parallel processing	Sequential/parallel processing
Speed	Slow (milliseconds)	Fast (nanoseconds)
Learning	Continuous learning	Batch/online learning
Storage	Distributed storage	Centralized storage

- **Biological:** Complex, fault-tolerant, self-repairing
- **Artificial:** Simple, precise, programmable

Mnemonic

“Bio is Complex, AI is Simple”

Question 1(c) [7 marks]

Explain types of AI with its applications.

Solution

Type of AI	Description	Applications
Narrow AI	AI designed for specific tasks	Voice assistants, recommendation systems
General AI	AI with human-level intelligence	Not yet achieved
Super AI	AI exceeding human intelligence	Theoretical concept

Mermaid Diagram (Code)

```
{Shaded}
{Highlighting} []
graph TD
    A[Types of AI] --> B[Narrow AI]
    A --> C[General AI]
    A --> D[Super AI]
    B --> E[Siri, Alexa]
    B --> F[Netflix Recommendations]
    C --> G[Human{-}level Tasks]
    D --> H[Beyond Human Intelligence]
{Highlighting}
{Shaded}
```

- **Current focus:** Narrow AI dominates today's applications
- **Future goal:** Achieving General AI safely

Mnemonic

“Narrow Now, General Goal, Super Scary”

Question 1(c) OR [7 marks]

Explain AI ethics and limitations.

Solution

Ethics Aspect	Description
Privacy	Protecting personal data and user information
Bias	Ensuring fairness across different groups
Transparency	Making AI decisions explainable
Accountability	Determining responsibility for AI actions

Limitations:

- **Data dependency:** Requires large, quality datasets
- **Computational power:** Needs significant processing resources
- **Lack of creativity:** Cannot truly create original concepts

Mnemonic

“Privacy, Bias, Transparency, Accountability”

Question 2(a) [3 marks]

Define the following terms: (1) Well posed Learning Problem (2) Machine Learning.

Solution

Term	Definition
Well posed Learning Problem	A learning problem with clearly defined task (T), performance measure (P), and experience (E) where performance improves with experience.
Machine Learning	A subset of AI that enables computers to learn and improve automatically from experience without being explicitly programmed.

- **Well posed formula:** $T + P + E = \text{Learning}$
- **ML advantage:** Automatic improvement from data

Mnemonic

“Task, Performance, Experience”

Question 2(b) [4 marks]

Explain Reinforcement Learning along with terms used in it.

Solution

Term	Description
Agent	The learner or decision maker
Environment	The world in which agent operates
Action	What agent can do in each state
State	Current situation of the agent
Reward	Feedback from environment

Mermaid Diagram (Code)

```
{Shaded}
{Highlighting} []
graph LR
    A[Agent] --> B[Action]
    B --> C[Environment]
    C --> D[State]
    D --> E[Reward]
    E --> A
{Highlighting}
{Shaded}
```

- **Learning process:** Trial and error approach
- **Goal:** Maximize cumulative reward

Mnemonic

“Agent Acts, Environment States and Rewards”

Question 2(c) [7 marks]

Compare Supervised, Unsupervised and Reinforcement Learning.

Solution

Aspect	Supervised	Unsupervised	Reinforcement
Data	Labeled data	Unlabeled data	Interactive data

Goal	Predict output	Find patterns	Maximize reward
Feedback	Immediate	None	Delayed
Examples	Classification	Clustering	Game playing

- **Supervised:** Teacher-guided learning
- **Unsupervised:** Self-discovery learning
- **Reinforcement:** Trial-and-error learning

Mnemonic

“Supervised has Teacher, Unsupervised Discovers, Reinforcement Tries”

Question 2(a) OR [3 marks]

Write Key features of Reinforcement Learning.

Solution

Feature	Description
Trial and Error	Learning through experimentation
Delayed Reward	Feedback comes after actions
Sequential Decision	Actions affect future states

- **No supervisor:** Agent learns independently
- **Exploration vs Exploitation:** Balance between trying new actions and using known good actions

Mnemonic

“Try, Delay, Sequence”

Question 2(b) OR [4 marks]

Explain Types of Reinforcement learning.

Solution

Type	Description
Positive RL	Adding positive stimulus to increase behavior
Negative RL	Removing negative stimulus to increase behavior

Based on Learning:

- **Model-based:** Agent learns environment model
- **Model-free:** Agent learns directly from experience

Mnemonic

“Positive Adds, Negative Removes”

Question 2(c) OR [7 marks]

Explain approaches to implement Reinforcement Learning.

Solution

Approach	Description	Example
Value-based	Learn value of states/actions	Q-Learning
Policy-based	Learn policy directly	Policy Gradient
Model-based	Learn environment model	Dynamic Programming

Mermaid Diagram (Code)

```
{Shaded}
{Highlighting} []
graph TD
    A[RL Approaches] --> B[Value{-}based]
    A --> C[Policy{-}based]
    A --> D[Model{-}based]
    B --> E[Q{-}Learning]
    C --> F[Policy Gradient]
    D --> G[Dynamic Programming]
{Highlighting}
{Shaded}
```

- **Value-based:** Estimates value functions
- **Policy-based:** Optimizes policy parameters
- **Model-based:** Uses environment model

Mnemonic

“Value, Policy, Model”

Question 3(a) [3 marks]

Describe the activation functions ReLU and sigmoid.

Solution

Function	Formula	Range
ReLU	$f(x) = \max(0, x)$	$[0, \infty)$
Sigmoid	$f(x) = 1/(1 + e^{-x})$	$(0, 1)$

- **ReLU advantage:** No vanishing gradient problem
- **Sigmoid advantage:** Smooth gradient, probabilistic output

Mnemonic

“ReLU Rectifies, Sigmoid Squashes”

Question 3(b) [4 marks]

Explain Multi-layer feed forward ANN.

Solution

Component	Description
Input Layer	Receives input data
Hidden Layers	Process information (multiple layers)
Output Layer	Produces final result
Connections	Forward direction only

- **Information flow:** Unidirectional from input to output
- **No cycles:** No feedback connections

Mnemonic

“Input → Hidden → Output(ForwardOnly)”

Question 3(c) [7 marks]

Draw the structure of ANN and explain functionality of each of its components.

Solution

Mermaid Diagram (Code)

```
{Shaded}
{Highlighting} []
graph LR
    A[Input Layer] --> B[Hidden Layer 1]
    B --> C[Hidden Layer 2]
    C --> D[Output Layer]

    subgraph "Components"
        E[Neurons]
        F[Weights]
        G[Bias]
        H[Activation Function]
    end
{Highlighting}
{Shaded}
```

Component	Functionality
Neurons	Processing units that receive inputs and produce outputs
Weights	Connection strengths between neurons
Bias	Additional parameter to shift activation function
Activation Function	Introduces non-linearity to the network

- **Input layer:** Receives and distributes input data
- **Hidden layers:** Extract features and patterns
- **Output layer:** Produces final classification or prediction
- **Connections:** Weighted links between neurons

Mnemonic

“Neurons with Weights, Bias, and Activation”

Question 3(a) OR [3 marks]

Write a short note on Backpropagation.

Solution

Aspect	Description
Purpose	Training algorithm for neural networks
Method	Gradient descent with chain rule
Direction	Backward error propagation

- **Process:** Calculate error gradients backwards through network
- **Update:** Adjust weights to minimize error

Mnemonic

“Back-ward Error Propagation”

Question 3(b) OR [4 marks]

Explain Single-layer feed forward network.

Solution

Feature	Description
Structure	Input layer directly connected to output layer
Layers	Only input and output layers
Limitations	Can only solve linearly separable problems
Example	Perceptron

- **Capability:** Limited to linear decision boundaries
- **Applications:** Simple classification tasks

Mnemonic

“Single Layer, Linear Limits”

Question 3(c) OR [7 marks]

Draw and explain the architecture of Recurrent neural network.

Solution

Mermaid Diagram (Code)

```
{Shaded}
{Highlighting} []
graph LR
    A[Input] --> B[Hidden State]
    B --> C[Output]
    B --> B[Self-loop]
    D[Previous State] --> B
{Highlighting}
{Shaded}
```

Component	Function
Hidden State	Maintains memory of previous inputs
Recurrent Connection	Feedback from hidden state to itself
Sequence Processing	Handles sequential data

- **Memory:** Retains information from previous time steps
- **Applications:** Language modeling, speech recognition
- **Advantage:** Can process variable-length sequences

Mnemonic

“Recurrent Remembers, Loops Back”

Question 4(a) [3 marks]

Define NLP and write down advantages of it.

Solution

Term	Definition
NLP	Natural Language Processing - enables computers to understand, interpret, and generate human language

Advantages:

- **Human-computer interaction:** Natural communication
- **Automation:** Automated text processing and analysis
- **Accessibility:** Voice interfaces for disabled users

Mnemonic

“Natural Language, Natural Interaction”

Question 4(b) [4 marks]

Compare NLU and NLG.

Solution

Aspect	NLU (Understanding)	NLG (Generation)
Purpose	Interpret human language	Generate human language
Input	Text/Speech	Structured data
Output	Structured data	Text/Speech
Examples	Sentiment analysis	Text summarization

- **NLU:** Converts unstructured text to structured data
- **NLG:** Converts structured data to natural text

Mnemonic

“NLU Understands, NLG Generates”

Question 4(c) [7 marks]

Explain word tokenization and frequency distribution of words with suitable example.

Solution

Process	Description	Example
Tokenization	Breaking text into individual words/tokens	“Hello world” → [“Hello”, “world”]
Frequency Distribution	Counting occurrence of each token	{“Hello”: 1, “world”: 1}

Example:

Text: "The cat sat on the mat"

Tokens: ["The", "cat", "sat", "on", "the", "mat"]

Frequency: {"The": 1, "cat": 1, "sat": 1, "on": 1, "the": 1, "mat": 1}

- **Case sensitivity:** "The" and "the" counted separately
- **Applications:** Text analysis, search engines
- **Preprocessing:** Essential step for NLP tasks

Mnemonic

"Tokenize then Count"

Question 4(a) OR [3 marks]

List disadvantages of NLP.

Solution

Disadvantage	Description
Ambiguity	Multiple meanings of words/sentences
Context dependency	Meaning changes with context
Language complexity	Grammar rules and exceptions

- **Cultural variations:** Different languages, dialects
- **Computational cost:** Resource-intensive processing

Mnemonic

"Ambiguous, Contextual, Complex"

Question 4(b) OR [4 marks]

Explain types of ambiguities in NLP.

Solution

Type	Description	Example
Lexical	Word has multiple meanings	"Bank" (financial/river)
Syntactic	Multiple parse trees possible	"I saw a man with a telescope"
Semantic	Multiple interpretations	"Flying planes can be dangerous"

- **Resolution:** Context analysis, statistical models
- **Challenge:** Major hurdle in NLP systems

Mnemonic

"Lexical words, Syntactic structure, Semantic meaning"

Question 4(c) OR [7 marks]

Explain stemming words and parts of speech(POS) tagging with suitable example.

Solution

Process	Description	Example
Stemming	Reducing words to root/stem form	"running" → "run", "flies" → "fli"
POS Tagging	Assigning grammatical categories	"The/DT cat/NN runs/VB fast/RB"

Stemming Example:

Original: ["running", "runs", "runner"]
Stemmed: ["run", "run", "runner"]

POS Tagging Example:

Sentence: "The quick brown fox jumps"
Tagged: "The/DT quick/JJ brown/JJ fox/NN jumps/VB"

- **Stemming purpose:** Reduce vocabulary size, group related words
- **POS purpose:** Understand grammatical structure
- **Applications:** Information retrieval, grammar checking

Mnemonic

"Stem to Root, Tag by Grammar"

Question 5(a) [3 marks]

Define the term word embedding and list various word embedding techniques.

Solution

Term	Definition
Word Embedding	Dense vector representations of words that capture semantic relationships

Techniques:

- **TF-IDF:** Term Frequency-Inverse Document Frequency
- **Bag of Words (BoW):** Simple word occurrence counting
- **Word2Vec:** Neural network-based embeddings

Mnemonic

"TF-IDF counts, BoW bags, Word2Vec vectorizes"

Question 5(b) [4 marks]

Explain about Challenges with TF-IDF and BoW.

Solution

Method	Challenges
TF-IDF	Sparse vectors, no semantic similarity, high dimensionality
BoW	Order ignored, context lost, sparse representation

Common Issues:

- **Sparsity:** Most vector elements are zero
- **No semantics:** Similar words have different vectors
- **High dimensions:** Memory and computation intensive

Mnemonic

“Sparse, No Semantics, High Dimensions”

Question 5(c) [7 marks]

Explain applications of NLP with suitable examples.

Solution

Application	Description	Example
Machine Translation	Translate between languages	Google Translate
Sentiment Analysis	Determine emotional tone	Product review analysis
Question Answering	Answer questions from text	Chatbots, virtual assistants
Spam Detection	Identify unwanted emails	Email filters
Spelling Correction	Fix spelling errors	Auto-correct in text editors

Mermaid Diagram (Code)

```
{Shaded}
{Highlighting} []
graph TD
    A[NLP Applications] --> B[Machine Translation]
    A --> C[Sentiment Analysis]
    A --> D[Question Answering]
    A --> E[Spam Detection]
    A --> F[Spelling Correction]
{Highlighting}
{Shaded}
```

- **Real-world impact:** Improves human-computer interaction
- **Business value:** Automates text processing tasks
- **Growing field:** New applications emerging constantly

Mnemonic

“Translate, Sentiment, Question, Spam, Spell”

Question 5(a) OR [3 marks]

Describe the Glove(Global Vector for word representation).

Solution

Aspect	Description
Purpose	Create word vectors using global corpus statistics
Method	Combines global matrix factorization and local context
Advantage	Captures both global and local statistical information

- **Global statistics:** Uses word co-occurrence information
- **Pre-trained:** Available trained vectors for common use

Mnemonic

“Global Vectors, Local Context”

Question 5(b) OR [4 marks]

Explain the Inverse Document Frequency (IDF).

Solution

Component	Formula	Purpose
IDF	$\log(N/df)$	Measure word importance across documents
N	Total documents	Corpus size
df	Document frequency	Documents containing the term

- **High IDF:** Rare words (more informative)
- **Low IDF:** Common words (less informative)
- **Application:** Part of TF-IDF weighting scheme

Mnemonic

“Inverse Document, Rare is Important”

Question 5(c) OR [7 marks]

Explain calculation of TF(Term Frequency) for a document with suitable example.

Solution

Method	Formula	Description
Raw TF	$f(t,d)$	Simple count of term in document
Normalized TF	$f(t,d)/\max(f(w,d))$	Normalized by maximum frequency
Log TF	$1 + \log(f(t,d))$	Logarithmic scaling

Example Document: “The cat sat on the mat. The mat was soft.”

Term	Count	Raw TF	Normalized TF	Log TF
“the”	3	3	1.0	1.48
“cat”	1	1	0.33	1.0
“mat”	2	2	0.67	1.30

Calculation Steps:

1. Count each term occurrence
2. Apply chosen TF formula
3. Use in TF-IDF calculation

Mnemonic

“Count, Normalize, Log”