

Subject Name Solutions

1333204 – Winter 2023

Semester 1 Study Material

Detailed Solutions and Explanations

Question 1(a) [3 marks]

Define: Field, Record, Metadata

Solution

Term	Definition
Field	A single unit of data representing a specific attribute in a database table (e.g., name, age, ID)
Record	A complete set of related fields that represents one entity instance (a row in a table)
Metadata	Data that describes the structure, properties, and relationships of other data (“data about data”)

Mnemonic

“FRM: Fields Row-up as Metadata”

Question 1(b) [4 marks]

Define (i) E-R model (ii) Entity (iii) Entity set and (iv) attributes

Solution

Term	Definition
E-R Model	A graphical approach to database design that models entities, their attributes, and relationships
Entity	A real-world object, concept, or thing that has an independent existence
Entity Set	A collection of similar entities that share the same attributes (represented as a table)
Attributes	Properties or characteristics that describe an entity (represented as columns in tables)

```
erDiagram
    ENTITY {
        string attribute1
        number attribute2
    }
    ENTITY_SET ||--o{ ENTITY : contains}
```

Mnemonic

“EEAA: Entities Exist As Attributes”

Question 1(c) [7 marks]

List the advantages and disadvantages of DBMS.

Solution

Advantages	Disadvantages
Data sharing: Multiple users can access simultaneously	Cost: Expensive hardware/software requirements
Data integrity: Maintains accuracy through constraints	Complexity: Requires specialized training
Data security: Controls access through permissions	Performance: Can be slow for large databases
Data independence: Changes to storage don't affect apps	Vulnerability: Central failure point risks data loss
Reduced redundancy: Eliminates duplicate data	Conversion costs: Migrating from file systems is expensive

Mnemonic

“SIDSR vs CCPVC” (Sharing, Integrity, Data independence, Security, Redundancy vs Cost, Complexity, Performance, Vulnerability, Conversion)

Question 1(c) OR [7 marks]

Write the full form of DBA. Explain the roles and responsibilities of DBA.

Solution

DBA: Database Administrator

Responsibilities of DBA

- Database design:** Creates efficient database schema
- Security management:** Sets up user access controls
- Performance tuning:** Optimizes queries and indexes
- Backup & recovery:** Implements data protection plans
- Maintenance:** Updates software and applies patches
- Troubleshooting:** Resolves database issues
- User support:** Trains and assists database users

```
flowchart TD
    A[Database Administrator] --> B[Database Design]
    A --> C[Security Management]
    A --> D[Performance Tuning]
    A --> E[Backup & Recovery]
    A --> F[Maintenance]
    A --> G[Troubleshooting]
    A --> H[User Support]
```

Mnemonic

“SPBT-MUS” (Security, Performance, Backup, Troubleshooting, Maintenance, User support)

Question 2(a) [3 marks]

Explain single valued v/s multi-valued attributes with suitable examples

Solution

Attribute Type	Description	Examples
Single-valued	Holds only one value for each entity instance	Employee ID, Birth Date, Name
Multi-valued	Can hold multiple values for the same entity	Phone Numbers, Skills, Email Addresses

```
erDiagram
    EMPLOYEE ^{
        string emp_id
        string name
        date birth_date
        string phone_numbers
        string skills
    }
```

Mnemonic

“SIM: Single Is Minimal, Multi Is Many”

Question 2(b) [4 marks]

Explain Key Constraints for E-R diagram

Solution

Key Constraint	Description
Primary Key	Uniquely identifies each entity in an entity set
Candidate Key	Any attribute that could serve as a primary key
Foreign Key	References primary key of another entity set
Super Key	Any set of attributes that uniquely identifies an entity

```
erDiagram
    STUDENT ^{
        int student_id PK
        string name
        string email
    }
    COURSE ^{
        int course_id PK
        string title
    }
    ENROLLMENT ^{
        int student_id FK
        int course_id FK
        date enroll_date
    }
    STUDENT ||{-{-}o}{ ENROLLMENT : has}
    COURSE ||{-{-}o}{ ENROLLMENT : includes}
```

Mnemonic

“PCFS: Primary Candidates Find Superkeys”

Question 2(c) [7 marks]

Construct an E-R diagram for banking management system.

Solution

```
erDiagram
    CUSTOMER{}{
        int customer\_id PK
        string name
        string address
        string phone
    }
    ACCOUNT{}{
        int account\_no PK
        string type
        float balance
        date open\_date
    }
    TRANSACTION{}{
        int trans\_id PK
        float amount
        string type
        date trans\_date
    }
    BRANCH{}{
        int branch\_id PK
        string name
        string location
    }
    CUSTOMER ||{-{-}o\{ ACCOUNT : has}
    ACCOUNT ||{-{-}o\{ TRANSACTION : includes}
    BRANCH ||{-{-}o\{ ACCOUNT : manages}
    ACCOUNT {-}{-}|| CUSTOMER : belongs\_to}
```

Key Entities and Relationships:

- **Customer:** Stores customer information
- **Account:** Different account types (savings, checking)
- **Transaction:** Records deposits, withdrawals
- **Branch:** Different bank locations
- **Relationships:** Customers have accounts, accounts have transactions, branches manage accounts

Mnemonic

“CATB: Customers Access Transactions at Branches”

Question 2(a) OR [3 marks]

Explain specialization v/s generalization with suitable examples

Solution

Concept	Direction	Description	Example
Specialization	Top-down	Breaking a general entity into more specific sub-entities	Person → <i>Student, Employee</i>
Generalization	Bottom-up	Combining similar entities into a higher-level entity	Car, Truck → <i>Vehicle</i>

```

erDiagram
    PERSON {
        int person_id
        string name
        string address
    }
    STUDENT {
        string major
        float gpa
    }
    EMPLOYEE {
        string department
        float salary
    }
    PERSON ||{-{-}}|| STUDENT : specializes}
    PERSON ||{-{-}}|| EMPLOYEE : specializes}

```

Mnemonic

“SG-TD-BU: Specialization Goes Top-Down, Generalization Builds Up”

Question 2(b) OR [4 marks]

Define Chasp trap. Explain when it occurs. Explain the solution for Chasp trap

Solution

Chasp trap: A problem that occurs in ER diagrams when there are multiple paths between entities, causing ambiguity in relationship interpretations.

Aspect	Description
Occurrence	When there are two or more distinct paths between entity types creating a cycle
Problem	Leads to incorrect or ambiguous query results
Solution	Break one of the relationships or add constraints to clarify the intended path

```

erDiagram
    STUDENT ||{-{-}}|| SECTION : enrolled_in}
    SECTION ||{-{-}}|| COURSE : part_of}
    STUDENT ||{-{-}}o{ COURSE : studies}

    %% Solution:
    %% Remove direct STUDENT to COURSE relationship
    %% Or add clear constraints

```

Mnemonic

“COP: Cycles Of Paths need breaking”

Question 2(c) OR [7 marks]

Construct an E-R diagram for college management system.

Solution

```

erDiagram
    STUDENT {
        int student_id PK
        string name
    }

```

```

        string address
        date dob
        string phone
    \}
    DEPARTMENT \{
        int dept\_id PK
        string name
        string location
        string hod
    \}
    FACULTY \{
        int faculty\_id PK
        string name
        string qualification
        date join\_date
    \}
    COURSE \{
        int course\_id PK
        string title
        int credits
        string description
    \}
    EXAM \{
        int exam\_id PK
        date date
        string type
    \}
STUDENT \|-{-}|- DEPARTMENT : belongs\_to}
FACULTY \|-{-}|- DEPARTMENT : works\_in}
DEPARTMENT ||{-{-}o\{ COURSE : offers}
FACULTY ||{-{-}o\{ COURSE : teaches}
STUDENT {-}{-}o\{ COURSE : enrolls}
STUDENT {-}{-}o\{ EXAM : takes}
COURSE ||{-{-}o\{ EXAM : has}

```

Key Entities and Relationships:

- **Student:** Stores student details
- **Department:** Academic divisions
- **Faculty:** Teachers and professors
- **Course:** Subjects taught
- **Exam:** Evaluation events
- **Relationships:** Students enroll in courses, faculty teach courses, departments offer courses

Mnemonic

“SDFCE: Students Delight Faculty by Completing Exams”

Question 3(a) [3 marks]

Explain GROUP BY clause with example.

Solution

GROUP BY clause groups rows that have the same values into summary rows.

Feature	Description
Purpose	Arranges identical data into groups for aggregate functions
Usage	Used with aggregate functions (COUNT, SUM, AVG, MAX, MIN)
Syntax	SELECT column1, COUNT(*) FROM table GROUP BY column1;

```

SELECT department, AVG(salary)
FROM employees
GROUP BY department;

```

Mnemonic

“GAS: Group And Summarize”

Question 3(b) [4 marks]

List Data Definition Language (DDL) commands. Explain any two DDL commands with examples.

Solution

DDL Commands: CREATE, ALTER, DROP, TRUNCATE, RENAME

Command	Description	Example
CREATE	Creates database objects like tables, views, indexes	CREATE TABLE students (id INT PRIMARY KEY, name VARCHAR(50));
ALTER	Modifies existing database objects	ALTER TABLE students ADD COLUMN email VARCHAR(100);
DROP	Removes database objects	DROP TABLE students;
TRUNCATE	Removes all records from a table	TRUNCATE TABLE students;

Mnemonic

“CADTR: Create, Alter, Drop, Truncate, Rename”

Question 3(c) [7 marks]

Perform the following Query on the “Students” table having the field’s enr_no, name, percent, branch in SQL.

Solution

```

{--1. Display all records in Students table}
SELECT * FROM Students;

{--2. Display only branch without duplicate value}
SELECT DISTINCT branch FROM Students;

{--3. Display all records sorted in descending order of name}
SELECT * FROM Students ORDER BY name DESC;

{--4. Add one new column to store address, named "address"}
ALTER TABLE Students ADD address VARCHAR(100);

{--5. Display all students belongs to branch "ICT"}
SELECT * FROM Students WHERE branch = {ICT};

{--6. Delete all students having percent less than 60}
DELETE FROM Students WHERE percent {} 60;

{--7. Display the students names starts with "S"}
SELECT * FROM Students WHERE name LIKE {S\%};

```

Query	Purpose
SELECT	Retrieves data from tables
DISTINCT	Eliminates duplicate values
ORDER BY	Sorts results in specified order
ALTER TABLE	Modifies table structure
WHERE	Filters records based on conditions
DELETE	Removes records matching conditions
LIKE	Pattern matching in string comparison

Mnemonic

“SDOAWDL: Select Distinct Order Alter Where Delete Like”

Question 3(a) OR [3 marks]

Explain GRANT command with syntax and example.

Solution

GRANT command gives specific privileges to users on database objects.

Component	Description
Syntax	<code>GRANT privilege(s) ON object TO user [WITH GRANT OPTION];</code>
Privileges	SELECT, INSERT, UPDATE, DELETE, ALL PRIVILEGES
Objects	Tables, views, sequences, etc.

```
GRANT SELECT, UPDATE ON employees TO user1;
GRANT ALL PRIVILEGES ON database\_name.* TO user2 WITH GRANT OPTION;
```

Mnemonic

“GPO: Grant Privileges to Others”

Question 3(b) OR [4 marks]

Compare Truncate command and Drop command.

Solution

Feature	TRUNCATE	DROP
Purpose	Removes all rows from table	Removes entire table structure
Structure	Keeps table structure intact	Deletes table definition completely
Recovery	Cannot be easily rolled back	Can be recovered until committed
Speed	Faster than DELETE	Quick operation
Triggers	Does not activate triggers	Does not activate triggers

```
{-- Truncate example}
```

```
TRUNCATE TABLE students;
```

```
{-- Drop example}
```

```
DROP TABLE students;
```

Mnemonic

"TRC-DST: Truncate Removes Contents, Drop Destroys Structure Totally"

Question 3(c) OR [7 marks]

Write the Output of Following Query.

Solution

Query	Output	Explanation
ABS(-23), ABS(49)	23, 49	Returns absolute value
SQRT(25), SQRT(81)	5, 9	Returns square root
POWER(3,2), POWER(-2,3)	9, -8	Returns x^y (first value raised to power of second)
MOD(15,4), MOD(21,3)	3, 0	Returns remainder after division
ROUND(123.446,1), ROUND(123.456,2)	123.4, 123.46	Rounds to specified decimal places
CEIL(234.45), CEIL(-234.45)	235, -234	Rounds up to nearest integer
FLOOR(-12.7), FLOOR(12.7)	-13, 12	Rounds down to nearest integer

```
SELECT ABS({-}23), ABS(49);      {-} 23, 49}
```

```
SELECT SQRT(25), SQRT(81);     {-} 5, 9
```

```
SELECT POWER(3,2), POWER({-}2,3); {-} 9, {-}8
```

```
SELECT MOD(15,4), MOD(21,3);  {-} 3, 0
```

```
SELECT ROUND(123.446,1), ROUND(123.456,2); {-} 123.4, 123.46
```

```
SELECT CEIL(234.45), CEIL({-}234.45); {-} 235, {-}234
```

```
SELECT FLOOR({-}12.7), FLOOR(12.7); {-} {-}13, 12
```

Mnemonic

"ASPMRCF: Absolute Square Power Modulo Round Ceiling Floor"

Question 4(a) [3 marks]

List data types in SQL. Explain any two data types with example.

Solution

SQL Data Types: INTEGER, FLOAT, VARCHAR, CHAR, DATE, DATETIME, BOOLEAN, BLOB

Data Type	Description	Example
INTEGER	Whole numbers without decimal points	id INTEGER = 101
VARCHAR	Variable-length character string	name VARCHAR(50) = 'John'
DATE	Stores date values (YYYY-MM-DD)	birth_date DATE = '2000-05-15'
FLOAT	Decimal numbers with floating point	salary FLOAT = 45000.50

```
CREATE TABLE employees (
    id INTEGER,
    name VARCHAR(50),
    salary FLOAT
);
```

Mnemonic

“IVDB: Integers & Varchars are Database Basics”

Question 4(b) [4 marks]

Explain Full function dependency with example.

Solution

Full Function Dependency: When Y is functionally dependent on X, but not on any subset of X.

Concept	Description	Example
Definition	Attribute B is fully functionally dependent on A if B depends on all of A	Student_ID → Name(<i>fulldependency</i>)
Non-example	When attribute depends only on part of composite key	{Student_ID, Course_ID} → Student_Name(<i>partial</i>)

```
flowchart TD
    A[Student\_ID] --{-{-}--> B[Student\_Name]]
    subgraph Full Function Dependency
        C[Course\_ID] --{-{-}--> D[Course\_Name]]
    end
    subgraph Partial Function Dependency
        E[Student\_ID, Course\_ID] --{-{-}--> F[Student\_Name]]
    end
```

Mnemonic

“FFD: Full, not Fraction of Dependency”

Question 4(c) [7 marks]

Define normalization. Explain 2NF (Second Normal Form) with example and solution.

Solution

Normalization: Process of organizing database to minimize redundancy and dependency by dividing large tables into smaller tables and defining relationships between them.

2NF (Second Normal Form):

- A table is in 2NF if it is in 1NF and no non-prime attribute is dependent on any proper subset of candidate key.

Before 2NF	Problem
Order(Order_ID, Product_ID, Product_Name, Quantity, Price)	Product_Name depends on only Product_ID, not full key

After 2NF	Solution
Order(Order_ID, Product_ID, Quantity) Product(Product_ID, Product_Name, Price)	Only full key dependencies Product details depend only on Product_ID

```

erDiagram
    ORDER {
        int order_id
        int product_id
        int quantity
    }
    PRODUCT {
        int product_id
        string product_name
        float price
    }
    ORDER {--}{-}|| PRODUCT : contains

```

Mnemonic

“2NF-PPD: Partial dependency Problems Divided”

Question 4(a) OR [3 marks]

Explain commands: 1) To_Number() 2) To_Char()

Solution

Function	Purpose	Syntax	Example
TO_NUMBER()	Converts string to number	TO_NUMBER(string, [format])	TO_NUMBER('123.45') = 123.45
TO_CHAR()	Converts number/date to string	TO_CHAR(value, [format])	TO_CHAR(1234, '9999') = '1234'

{-->} Convert string to number
SELECT TO_NUMBER({123.45}) FROM dual; {-->} 123.45

{-->} Convert date to formatted string
SELECT TO_CHAR(SYSDATE, {DD{-}MON{-}YYYY}) FROM dual; {-->} 20{-}JAN{-}2024

{-->} Convert number to formatted string
SELECT TO_CHAR(1234.56, {\$9,999.99}) FROM dual; {-->} \$1,234.56

Mnemonic

“NC: Numbers and Characters conversion”

Question 4(b) OR [4 marks]

Explain 1NF (First Normal Form) with example and solution.

Solution

1NF (First Normal Form): A relation is in 1NF if it contains no repeating groups or arrays.

Before 1NF	Problem
Student(ID, Name, Courses) Example: (101, John, “Math,Science,History”)	Courses column contains multiple values Multi-valued attribute

After 1NF	Solution
Student(ID, Name, Course) Examples: (101, John, Math), (101, John, Science), (101, John, History)	One course per row Atomic values
<pre> erDiagram STUDENT\._BEFORE \{ int id string name string courses } STUDENT\._AFTER \{ int id string name string course } </pre>	

Mnemonic

“1NF-ARM: Atomic values Remove Multivalues”

Question 4(c) OR [7 marks]

Explain function dependency in SQL. Explain Partial function dependency with example.

Solution

Functional Dependency: A relationship where one attribute determines the value of another attribute.

Notation: $X \rightarrow Y$ (X determines Y)

Partial Functional Dependency: When an attribute depends on only part of a composite primary key.

Concept	Example	Explanation
Composite Key	{Student_ID, Course_ID}	Together forms primary key
Partial Dependency	{Student_ID, Course_ID} \rightarrow Student_Name	Student_Name depends only on Student_ID
Problem	Update anomalies, data redundancy	Same student name repeated for multiple courses

```

flowchart TD
    A[Student\_ID] {-{->} B[Student\_Name]}
    C[Course\_ID] {-{->} D[Course\_Name]}
    E["Student\_ID, Course\_ID"] {-{->} F[Grade]}
    subgraph "Partial Dependency"
        A {-{->} B}
        end
    subgraph "Full Dependency"
        E {-{->} F}
        end

```

Solution: Decompose into separate tables where each non-key attribute is fully dependent on the key.

Mnemonic

“PD-CPK: Partial Dependency - Component of Primary Key”

Question 5(a) [3 marks]

Explain the properties of Transaction with example.

Solution

Transaction Properties (ACID):

Property	Description	Example
Atomicity	All operations complete successfully or none does	Bank transfer: debit and credit both happen or neither
Consistency	Database remains in valid state before and after	Account balance constraints remain valid
Isolation	Transactions execute as if they were the only one	Two users updating same record don't interfere
Durability	Committed changes survive system failure	Once confirmed, a deposit remains even after power loss

```
flowchart LR
    A[START TRANSACTION] --> B[Debit Account A]
    B --> C[Credit Account B]
    C --> D{Successful?}
    D -- Yes --> E[COMMIT]
    D -- No --> F[ROLLBACK]
```

Mnemonic

“ACID: Atomicity, Consistency, Isolation, Durability”

Question 5(b) [4 marks]

Write the Queries using set operators to find following using given “Student” and “CR” (Class Representative) tables.

Solution

```
{-- 1. List the name of the persons who are either a student or a CR}
SELECT Stnd\_Name FROM Student
UNION
SELECT CR\_Name FROM CR;

{-- 2. List the name of the persons who are a student as well as a CR}
SELECT Stnd\_Name FROM Student
INTERSECT
SELECT CR\_Name FROM CR;

{-- 3. List the name of the persons who are only a student and not a CR}
SELECT Stnd\_Name FROM Student
MINUS
SELECT CR\_Name FROM CR;

{-- 4. List the name of the persons who are only a CR and not a student}
SELECT CR\_Name FROM CR
MINUS
SELECT Stnd\_Name FROM Student;
```

Set Operator	Purpose	Result for Example
UNION	Combines all distinct rows	Manoj, Rahil, Jiya, Rina, Jitesh, Priya
INTERSECT	Returns only common rows	Manoj, Rina
MINUS	Returns rows in first set but not second	Rahil, Jiya

MINUS (reversed)

Returns rows in second set but not first

Jitesh, Priya

Mnemonic

“UIMD: Union Includes, Minus Divides”

Question 5(c) [7 marks]

Explain Conflict Serializability in detail.

Solution

Conflict Serializability: A schedule is conflict serializable if it can be transformed into a serial schedule by swapping non-conflicting operations.

Key Concepts	Description
Conflict operations	Two operations conflict if they access same data item and at least one is write
Precedence graph	Directed graph showing conflicts between transactions
Serializable	If precedence graph has no cycles, schedule is conflict serializable

```
flowchart LR
    subgraph "Transaction T1"
        A[Read X] --{-{-}} B[Write X]
    end
    subgraph "Transaction T2"
        C[Read X] --{-{-}} D[Write X]
    end
    subgraph "Conflicts"
        B --{-{-}} C
    end
```

Example:

- T1: R(X), W(X)
- T2: R(X), W(X)

Serializable schedules:

- T1 followed by T2: R1(X), W1(X), R2(X), W2(X)
- T2 followed by T1: R2(X), W2(X), R1(X), W1(X)

Non-serializable: R1(X), R2(X), W1(X), W2(X) - Creates cycle in precedence graph

Mnemonic

“COPS: Conflict Operations Produce Serializability”

Question 5(a) OR [3 marks]

Explain the concept of Transaction with example.

Solution

Transaction: A logical unit of work that must be either completely performed or completely undone.

Transaction Phases	Description	Example
BEGIN	Marks start of transaction	START TRANSACTION

Execute operations	Database operations (read/write)	UPDATE account SET balance = balance - 1000 WHERE id = 123
COMMIT/ROLLBACK	End transaction with success/failure	COMMIT or ROLLBACK

```
flowchart LR
    A[BEGIN TRANSACTION] --> B[Read account balance]
    B --> C[Check if sufficient funds]
    C --> D[Update account balance]
    D --> E[Create transaction record]
    E --> F[COMMIT]
    F --> G[ROLLBACK]
```

Example:

```
BEGIN TRANSACTION;
UPDATE accounts SET balance = balance - 1000 WHERE acc\_no = 123;
UPDATE accounts SET balance = balance + 1000 WHERE acc\_no = 456;
COMMIT;
```

Mnemonic

“BEC: Begin, Execute, Commit”

Question 5(b) OR [4 marks]

Explain equi-join with syntax and example.

Solution

Equi-join: A join operation that uses equality comparison operator.

Feature	Description
Syntax	SELECT columns FROM table1, table2 WHERE table1.column = table2.column;
Purpose	Combines rows from two tables based on matching column values
Alternative	SELECT columns FROM table1 INNER JOIN table2 ON table1.column = table2.column;

```
{-{-} Traditional syntax}
SELECT s.name, d.dept\_name
FROM students s, departments d
WHERE s.dept\_id = d.dept\_id;
```

```
{-{-} INNER JOIN syntax}
SELECT s.name, d.dept\_name
FROM students s INNER JOIN departments d
ON s.dept\_id = d.dept\_id;
```

Mnemonic

“EQ-ME: Equality Matches Entries”

Question 5(c) OR [7 marks]

Explain View Serializability in detail.

Solution

View Serializability: A schedule is view serializable if it is view equivalent to some serial schedule.

Condition	Description
Initial read	If T1 reads initial value of data item X in schedule S, it must also read initial value in schedule S'
Final write	If T1 performs final write of data item X in S, it must also perform final write in S'
Dependency preservation	If T1 reads value of X written by T2 in S, it must also read from T2 in S'

```
flowchart LR
    A["Schedule S"] -->|{-{-}}| B["View Equivalent?"]
    B -->|Yes| C["View Serializable"]
    B -->|No| D["Not View Serializable"]

    subgraph "Read-Write Analysis"
        E["Initial Read Check"]
        F["Final Write Check"]
        G["Read-from-Write Check"]
    end
```

Comparison:

- **Conflict serializability:** More restrictive, easier to test (precedence graph)
- **View serializability:** More general, harder to test (NP-complete)

Example of view serializable but not conflict serializable:

- T1: W(X)
- T2: W(X)
- T3: R(X)
- Schedule: W1(X), W2(X), R3(X) - View equivalent to serial schedule T2,T1,T3

Mnemonic

“VIR-FF: View preserves Initial Reads and Final writes”