

Database Management System (1333204) - Summer 2024 Solution

Milav Dabgar

June 14, 2024

Question 1(a) [3 marks]

Define: DBMS, Instance, Metadata

Solution

- **DBMS (Database Management System):** Software that enables users to create, maintain, and access databases by controlling data organization, storage, retrieval, security, and integrity.
- **Instance:** The actual data stored in a database at a particular moment in time. It's the current state or snapshot of a database.
- **Metadata:** Data about data that describes database structure, including tables, fields, relationships, constraints, and indexes.

Mnemonic

[title=DIM view]Database system, Instance snapshot, Metadata description

Question 1(b) [4 marks]

Define and Explain with example: 1.Entity 2. Attribute

Solution

Table: Entity vs Attribute

| Concept | Definition | Example |
|-----------|--|--|
| Entity | A real-world object or concept that can be distinctly identified | Student (John), Book (Harry Potter), Car (Toyota Camry) |
| Attribute | Characteristic or property that describes an entity | Student: roll_no, name, address Book: ISBN, title, author |

Diagram:



Mnemonic

[title=EA-PC]Entities Are Physical/Conceptual, Attributes Provide Characteristics

Question 1(c) [7 marks]

Write the full form of DBA. Explain the roles and responsibilities of DBA.

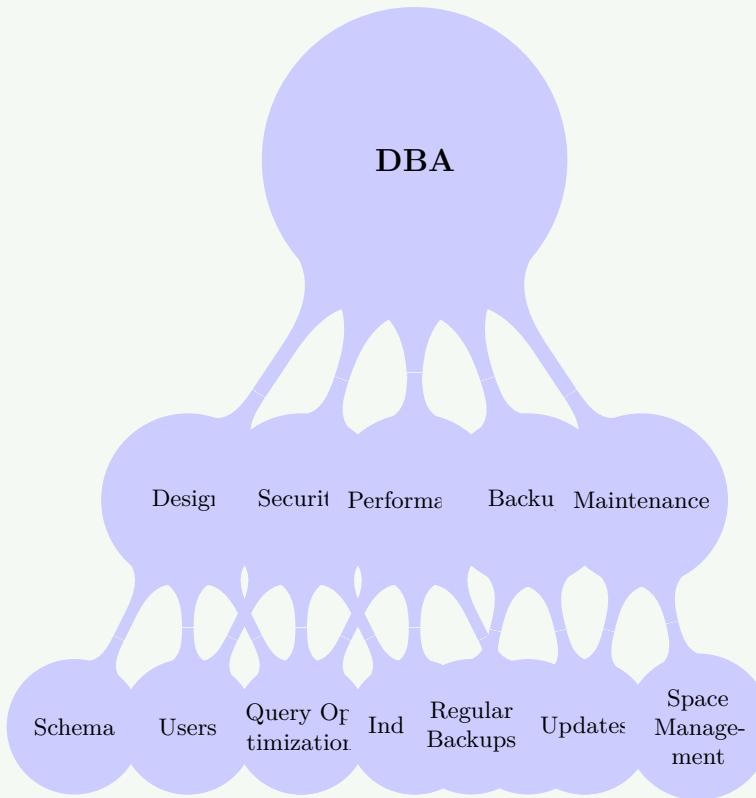
Solution

DBA stands for Database Administrator.

Table: DBA Responsibilities

| Role | Description |
|---------------------|--|
| Database Design | Creates logical/physical database structure and schema |
| Security Management | Controls access through user accounts and permissions |
| Performance Tuning | Optimizes queries, indexes for faster data retrieval |
| Backup & Recovery | Implements strategies to prevent data loss |
| Maintenance | Updates software, applies patches, monitors space |

Diagram:



Mnemonic

[title=SPMBU]Security, Performance, Maintenance, Backup, Updates

Question 1(c) OR [7 marks]

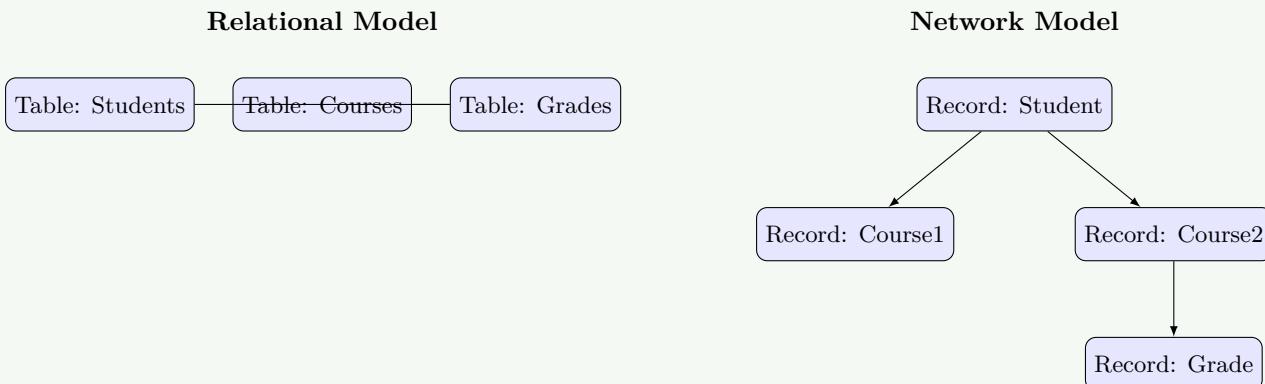
Explain relational and network data models in detail.

Solution

Table: Relational vs Network Data Models

| Feature | Relational Model | Network Model |
|----------------|--|--|
| Structure | Tables (relations) with rows and columns | Records connected by pointers forming complex networks |
| Relationship | Related through primary & foreign keys | Direct links between parent-child records |
| Flexibility | High - tables can be joined as needed | Limited - predefined physical connections |
| Examples | MySQL, Oracle, SQL Server | IDS, IDMS |
| Query Language | SQL (structured query language) | Procedural languages |

Diagram:



Mnemonic

[title=RSPEN]Relational uses Sets, Pointers Enable Networks

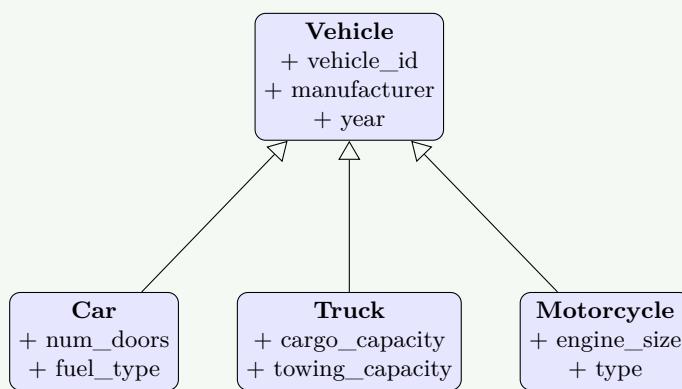
Question 2(a) [3 marks]

Draw figure and Explain Generalization.

Solution

Generalization: The process of extracting common characteristics from two or more entities to create a new higher-level entity.

Diagram:



Mnemonic

[title=BUSH]Bottom-Up Shared Hierarchy

Question 2(b) [4 marks]

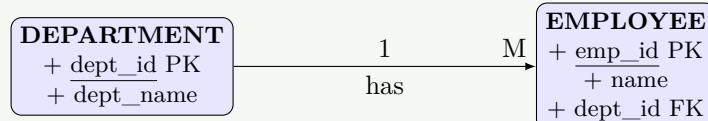
Explain Primary Key and Foreign Key Constraints.

Solution

Table: Primary Key vs Foreign Key

| Constraint | Definition | Properties | Example |
|-------------|--|---|---|
| Primary Key | Uniquely identifies each record in a table | Unique, Not Null, Only one per table | StudentID in Students table |
| Foreign Key | Links data between tables, references a primary key in another table | Can be NULL, Multiple allowed per table | DeptID in Employees table referencing Departments table |

Diagram:

**Mnemonic**

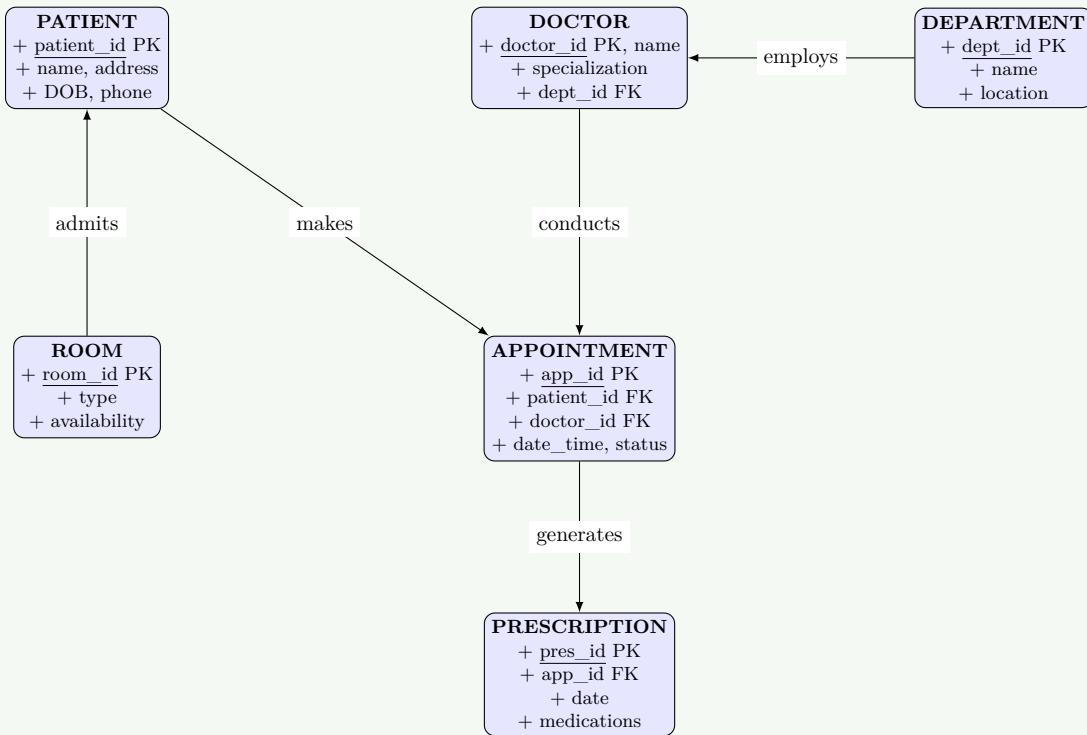
[title=PURE FIRE]Primary Uniquely References Entities, Foreign Imports Referenced Entities

Question 2(c) [7 marks]

Construct an E-R diagram for Hospital Management System.

Solution

E-R Diagram for Hospital Management System:

**Mnemonic**

[title=PADRE]Patients Appointments Doctors Rooms Entities

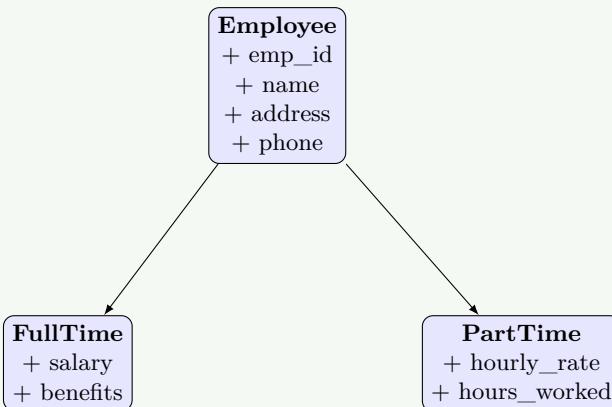
Question 2(a) OR [3 marks]

Draw figure and Explain Specialization.

Solution

Specialization: The process of creating new entities from an existing entity by adding unique attributes to distinguish them.

Diagram:

**Mnemonic**

[title=TDSB]Top-Down Specialized Breakdown

Question 2(b) OR [4 marks]

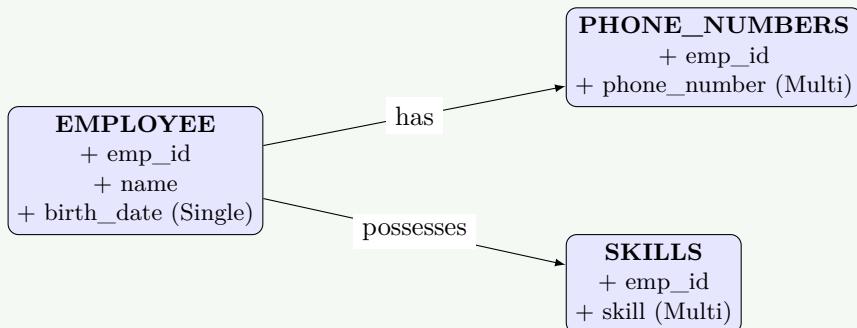
Explain single valued v/s multi-valued attributes with suitable examples.

Solution

Table: Single-valued vs Multi-valued Attributes

| Type | Definition | Example | Implementation |
|---------------|--|--------------------------------|---------------------------------------|
| Single-valued | Contains only one value for each entity instance | Person's birth date, SSN | Directly stored in table columns |
| Multi-valued | Can have multiple values for the same entity | Person's skills, phone numbers | Separate table or specialized formats |

Diagram:



Mnemonic

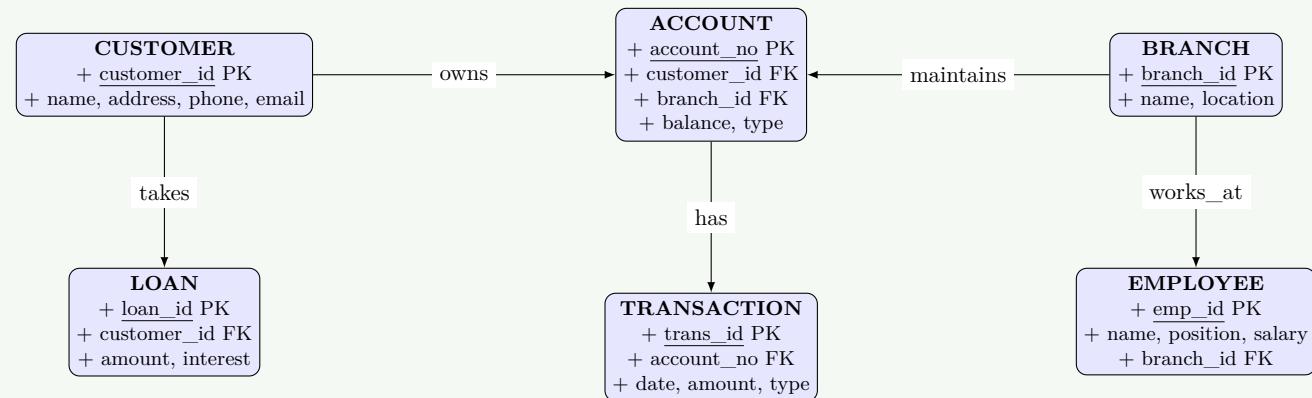
[title=SOME]Single One, Multiple Entries

Question 2(c) OR [7 marks]

Construct an E-R diagram for Banking Management System.

Solution

E-R Diagram for Banking Management System:



Mnemonic

[title=CABLE]Customers Accounts Branches Loans Employees

Question 3(a) [3 marks]

Explain WHERE and DESC clause with example.

Solution

Table: WHERE and DESC Clauses

| Clause | Purpose | Syntax | Example |
|--------|---|--------------------------------|--|
| WHERE | Filters rows based on specified condition | SE... FROM ... WHERE condition | SELECT * FROM employees WHERE salary > 50000 |
| DESC | Sorts results in descending order | SE... ORDER BY ... DESC | SELECT * FROM products ORDER BY price DESC |

Diagram (Data Flow):

Example Data Operation

```

1  -- Original
2  | ID | Name    | Marks |
3  | 1  | Alice   | 85   |
4  | 2  | Bob     | 92   |
5  | 3  | Carol   | 78   |
6
7  -- WHERE Marks > 80
8  | 1  | Alice   | 85   |
9  | 2  | Bob     | 92   |
10
11 -- ORDER BY Marks DESC
12 | 2  | Bob     | 92   |
13 | 1  | Alice   | 85   |
14 | 3  | Carol   | 78   |

```

Mnemonic

[title=WDF]Where filters Data, DESC orders First-highest

Question 3(b) [4 marks]

List DDL commands. Explain any two DDL commands with examples.

Solution

DDL (Data Definition Language) Commands:

1. CREATE
2. ALTER
3. DROP
4. TRUNCATE
5. RENAME

Table: CREATE and ALTER Commands

| Command | Purpose | Syntax | Example |
|---------|--------------------------|------------------|------------------------------------|
| CREATE | Creates database objects | CREATE TABLE ... | CREATE TABLE students (id INT...) |
| ALTER | Modifies structure | ALTER TABLE ... | ALTER TABLE students ADD COLUMN... |

```

1  -- CREATE example
2  CREATE TABLE employees (
3      emp_id INT PRIMARY KEY,
4      name VARCHAR(50) NOT NULL,
5      dept VARCHAR(30),
6      salary DECIMAL(10,2)
7 );
8
9  -- ALTER example
10 ALTER TABLE employees
11 ADD COLUMN hire_date DATE;

```

Mnemonic

[title=CADTR]Create Alter Drop Truncate Rename

Question 3(c) [7 marks]

Perform the following Query on the table "Company" having the field's eno, ename, salary, dept in SQL.

Queries:

1. Display all records in Company table.
2. Display only dept without duplicate value.
3. Display all records sorted in descending order of ename.
4. Add one new column "cityname" to store city.
5. Display name of all employees who do not stay in city "Mumbai".
6. Delete all employees having salary less than 10,000.
7. Display the employee names starts with "A".

Solution

```

1  -- 1. Display all records
2  SELECT * FROM Company;
3
4  -- 2. Display only dept without duplicates
5  SELECT DISTINCT dept FROM Company;
6
7  -- 3. Display records sorted by ename descending
8  SELECT * FROM Company ORDER BY ename DESC;
9
10 -- 4. Add new column "cityname"
11 ALTER TABLE Company ADD COLUMN cityname VARCHAR(50);
12
13 -- 5. Employees not in Mumbai
14 SELECT ename FROM Company WHERE cityname != 'Mumbai';
15
16 -- 6. Delete employees with salary < 10000
17 DELETE FROM Company WHERE salary < 10000;
18
19 -- 7. Employee names starting with "A"

```

20 | `SELECT ename FROM Company WHERE ename LIKE 'A%';`

Table: SQL Operations

| Operation | SQL Command | Purpose |
|-----------|---|--------------------|
| SELECT | <code>SELECT * FROM Company</code> | Retrieve all data |
| DISTINCT | <code>SELECT DISTINCT dept</code> | Remove duplicates |
| ORDER BY | <code>ORDER BY ename DESC</code> | Sort in descending |
| ALTER | <code>ALTER TABLE ADD COLUMN</code> | Add new column |
| WHERE | <code>WHERE cityname != 'Mumbai'</code> | Filter condition |
| DELETE | <code>DELETE FROM WHERE</code> | Remove records |
| LIKE | <code>WHERE ename LIKE 'A%'</code> | Pattern matching |

Mnemonic

[title=SODA-WDL]Select Order Distinct Alter - Where Delete Like

Question 3(a) OR [3 marks]

Explain SELECT and DISTINCT clause with example.

Solution**Table: SELECT and DISTINCT Clauses**

| Clause | Purpose | Syntax | Example |
|-----------|------------------------------|---|--|
| SE-LECT | Retrieves data from database | <code>SELECT columns FROM table</code> | <code>SELECT name, age FROM students</code> |
| DIS-TINCT | Eliminates duplicate values | <code>SELECT DISTINCT columns FROM table</code> | <code>SELECT DISTINCT department FROM employees</code> |

```

1 -- Original: Sales, IT, HR, IT, Sales
2
3 -- SELECT dept_name
4 Sales
5 IT
6 HR
7 IT
8 Sales
9
10 -- SELECT DISTINCT dept_name
11 Sales
12 IT
13 HR

```

Mnemonic

[title=SUD]Select Unique with Distinct

Question 3(b) OR [4 marks]

List DML commands. Explain any two DML commands with examples.

Solution

DML (Data Manipulation Language) Commands:

1. INSERT
2. UPDATE
3. DELETE
4. SELECT

Table: INSERT and UPDATE Commands

| Com-mand | Purpose | Syntax | Example |
|----------|------------------------------|-----------------------------|--|
| INSERT | Adds new records | INSERT INTO ... VALUES | INSERT INTO students VALUES (1, 'John', 85) |
| UPDATE | Modifies existing records | UPDATE ... SET ... WHERE | UPDATE students SET marks=90 WHERE id=1 |

```

1 -- INSERT example
2 INSERT INTO employees (emp_id, name, dept, salary)
3 VALUES (101, 'John Smith', 'IT', 65000);
4
5 -- UPDATE example
6 UPDATE employees
7 SET salary = 70000
8 WHERE emp_id = 101;

```

Mnemonic

[title=IUDS]Insert Update Delete Select

Question 3(c) OR [7 marks]

Write the Output of Following Query.

Solution

Table: SQL Function Outputs

| Function | Description | Output |
|------------------------|---------------------|--------|
| ABS(-34), ABS(16) | Absolute value | 34, 16 |
| SQRT(16), SQRT(64) | Square root | 4, 8 |
| POWER(5,2), POWER(2,4) | Power function | 25, 16 |
| MOD(15,3), MOD(13,3) | Modulus (remainder) | 0, 1 |
| ROUND(123.456,1) | Round to 1 decimal | 123.5 |
| ROUND(123.456,2) | Round to 2 decimals | 123.46 |
| CEIL(122.6) | Round up | 123 |
| CEIL(-122.6) | Round up (negative) | -122 |
| FLOOR(-157.5) | Round down | -158 |
| FLOOR(157.5) | Round down | 157 |

Mnemonic

[title=ASPRCF]Absolute Square Power Remainder Ceiling Floor

Question 4(a) [3 marks]

List data types in SQL. Explain 1.VARCHAR() and 2.INT() data types with example.

Solution**SQL Data Types Categories:**

1. Numeric (INT, FLOAT, DECIMAL)
2. Character (CHAR, VARCHAR)
3. Date/Time (DATE, TIME, DATETIME)
4. Binary (BLOB, BINARY)
5. Boolean (BOOL)

Table: VARCHAR and INT Data Types

| Data Type | Description | Size | Example |
|------------|----------------------------------|--------------------|-----------------------|
| VARCHAR(n) | Variable-length character string | Up to n characters | VARCHAR(50) for names |
| INT | Integer numeric data | Usually 4 bytes | INT for IDs, counts |

```

1 CREATE TABLE students (
2     student_id INT PRIMARY KEY,
3     name VARCHAR(50) NOT NULL,
4     age INT,
5     email VARCHAR(100)
6 );

```

Mnemonic

[title=VIA]Variable strings, Integers for Ages

Question 4(b) [4 marks]

Explain 2NF (Second Normal Form) with example and solution.

Solution

2NF Definition: A relation is in 2NF if it is in 1NF and no non-prime attribute is dependent on any proper subset of any candidate key.

Table: Before 2NF

| student_id | course_id | course_name | instructor |
|------------|-----------|-------------|--------------|
| S1 | C1 | Database | Prof. Smith |
| S1 | C2 | Networking | Prof. Jones |
| S2 | C1 | Database | Prof. Smith |
| S3 | C3 | Programming | Prof. Wilson |

Problem: Non-prime attributes (course_name, instructor) depend only on course_id, not the entire key (student_id, course_id).

Diagram: 2NF Solution

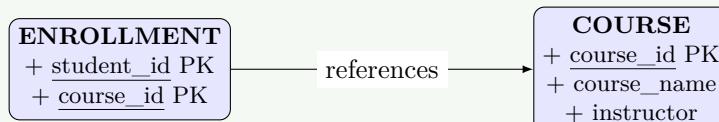


Table: After 2NF

Enrollment Table:

| student_id | course_id |
|------------|-----------|
| S1 | C1 |
| S1 | C2 |
| S2 | C1 |
| S3 | C3 |

Course Table:

| course_id | course_name | instructor |
|-----------|-------------|--------------|
| C1 | Database | Prof. Smith |
| C2 | Networking | Prof. Jones |
| C3 | Programming | Prof. Wilson |

Mnemonic

[title=PFPK]Partial Functional dependency on Primary Key

Question 4(c) [7 marks]

Explain function dependency. Explain Partial function dependency with example.

Solution

Functional Dependency: Relationship between attributes where one attribute's value determines another attribute's value.

Notation: $X \rightarrow Y$ (X determines Y)

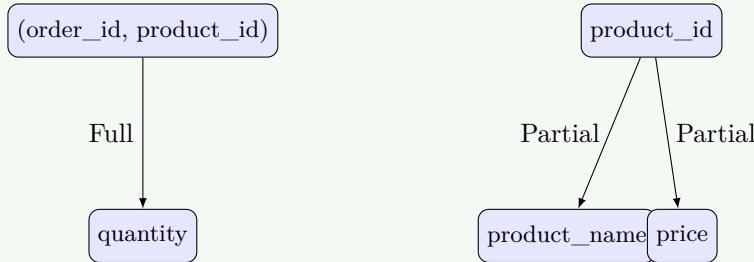
Partial Functional Dependency: When a non-prime attribute depends on part of a composite key rather than the whole key.

Table: Order Details (Before Normalization)

| order_id | product_id | quantity | product_name | price |
|----------|------------|----------|--------------|-------|
| O1 | P1 | 5 | Keyboard | 50 |
| O1 | P2 | 2 | Mouse | 25 |
| O2 | P1 | 1 | Keyboard | 50 |
| O3 | P3 | 3 | Monitor | 200 |

Functional Dependencies:

- $(\text{order_id}, \text{product_id}) \rightarrow \text{quantity}$
- $\text{product_id} \rightarrow \text{product_name}$
- $\text{product_id} \rightarrow \text{price}$

Diagram (Dependency Graph):**Solution (Normalized Tables):**

Orders Table:

| order_id | product_id | quantity |
|----------|------------|----------|
| O1 | P1 | 5 |
| O1 | P2 | 2 |
| O2 | P1 | 1 |
| O3 | P3 | 3 |

Products Table:

| product_id | product_name | price |
|------------|--------------|-------|
| P1 | Keyboard | 50 |
| P2 | Mouse | 25 |
| P3 | Monitor | 200 |

Mnemonic

[title=PDPK]Partial Dependency on Part of Key

Question 4(a) OR [3 marks]

Explain commands: 1) To_Char() 2) To_Date()

Solution**Table: Conversion Functions**

| Function | Purpose | Syntax | Example |
|-----------|--------------------------------|-------------------|-----------------------------|
| TO_CHAR() | Converts date/number to string | TO_CHAR(val, fmt) | TO_CHAR(SYSDATE, 'DD-MON') |
| TO_DATE() | Converts string to date | TO_DATE(str, fmt) | TO_DATE('14-JUN', 'DD-MON') |

```

1 SELECT TO_CHAR(SYSDATE, 'DD-MON-YYYY') FROM DUAL;
2 SELECT TO_DATE('2024-06-14', 'YYYY-MM-DD') FROM DUAL;

```

Mnemonic

[title=DCS]Date Conversion Strings

Question 4(b) OR [4 marks]

Explain Full function dependency with example.

Solution

Full Functional Dependency: When an attribute is functionally dependent on a composite key, and dependent on the entire key, not just part of it.

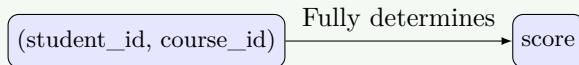
Table: Exam Results

| student_id | course_id | exam_date | score |
|------------|-----------|------------|-------|
| S1 | C1 | 2024-05-10 | 85 |
| S1 | C2 | 2024-05-15 | 92 |
| S2 | C1 | 2024-05-10 | 78 |
| S2 | C2 | 2024-05-15 | 88 |

Full Functional Dependency:

- $(\text{student_id}, \text{course_id}) \rightarrow \text{score}$ (score depends on both student and course)

Diagram:



Explanation: The score attribute fully depends on the composite key (student_id, course_id) because:

- Different students can have different scores for the same course
- Same student can have different scores for different courses
- We need both student_id and course_id to determine a specific score

Mnemonic

[title=FCEK]Fully dependent on Complete/Entire Key

Question 4(c) OR [7 marks]

Define normalization. Explain 1NF (First Normal Form) with example and solution.

Solution

Normalization: Process of organizing data to minimize redundancy, improve data integrity, and eliminate anomalies by dividing larger tables into smaller related tables.

1NF Definition: A relation is in 1NF if all attributes contain atomic (indivisible) values only.

Table: Before 1NF

| student_id | name | courses |
|------------|------|-----------------------------|
| S1 | John | Math, Physics |
| S2 | Mary | Chemistry, Biology, Physics |
| S3 | Tim | Computer Science |

Problems:

- Non-atomic values (multiple courses per cell)
- Cannot easily query or update specific courses

Diagram:



Table: After 1NF

| student_id | name | course |
|------------|------|------------------|
| S1 | John | Math |
| S1 | John | Physics |
| S2 | Mary | Chemistry |
| S2 | Mary | Biology |
| S2 | Mary | Physics |
| S3 | Tim | Computer Science |

Mnemonic

[title=ASAV]Atomic Single-value Attributes only Valid

Question 5(a) [3 marks]

Explain the concept of Transaction with example.

Solution

Transaction: A logical unit of work executed completely or not at all.

Table: Transaction Properties

| Property | Description |
|-------------|-------------------------|
| Atomicity | All or nothing |
| Consistency | Valid state transition |
| Isolation | Concurrent independence |
| Durability | Permanent persistence |

```

1 BEGIN TRANSACTION;
2     UPDATE accounts SET balance = balance - 500 WHERE id = 'A';
  
```

```

3   UPDATE accounts SET balance = balance + 500 WHERE id = 'B';
4   COMMIT;

```

Mnemonic

[title=ACID]Atomicity Consistency Isolation Durability

Question 5(b) [4 marks]**Explain equi join with syntax and example.****Solution****Equi Join:** Uses equality operator to match records.

```

1   SELECT e.name, d.dept_name
2   FROM employees e, departments d
3   WHERE e.dept_id = d.dept_id;

```

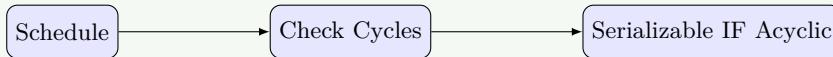
Diagram:**Mnemonic**

[title=MEET]Match Equal Elements Every Table

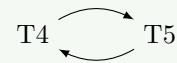
Question 5(c) [7 marks]**Explain Conflict Serializability in detail.****Solution****Conflict Serializability:** A way to ensure correctness of concurrent transactions by guaranteeing that the execution schedule is equivalent to some serial execution.**Table: Key Concepts in Conflict Serializability**

| Concept | Description |
|------------------------|---|
| Conflicting Operations | Two operations conflict if they access same data item and at least one is a write |
| Precedence Graph | Directed graph showing conflicts between transactions |
| Conflict Serializable | Schedule is conflict serializable if its precedence graph is acyclic |

Diagram:



$T_1 \xrightarrow{\text{Serializable}} T_2$



Cycle (Not Serial)

Example: Consider transactions T1 and T2:

- T1: Read(A), Write(A)
- T2: Read(A), Write(A)

Schedule S1: R1(A), W1(A), R2(A), W2(A) - Serializable (equivalent to $T_1 \rightarrow T_2$) Schedule S2: R1(A), R2(A), W1(A), W2(A) - Not serializable (contains cycle in precedence graph)

Steps to Determine Conflict Serializability:

1. Identify all pairs of conflicting operations
2. Construct the precedence graph
3. Check if the graph has cycles
4. If no cycles, the schedule is conflict serializable

Mnemonic

[title=COPS]Conflicts, Operations, Precedence, Serializability

Question 5(a) OR [3 marks]

Explain the properties of Transaction with example.

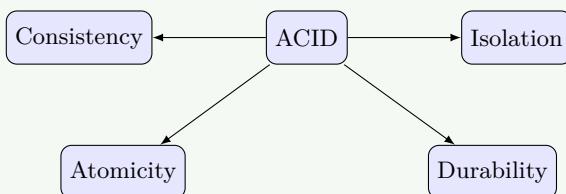
Solution

ACID Properties of Transactions:

Table: ACID Properties

| Prop- erty | Description | Example |
|------------------|---|---|
| Atomic- ity | All operations complete successfully or none do | Bank transfer - both debit and credit must succeed or fail together |
| Consis- tency | Database must be in a consistent state before and after transaction | After transferring \$100, total money in system remains unchanged |
| Isola- tion | Concurrent transactions don't interfere with each other | Transaction A doesn't see partial results of Transaction B |
| Durabil- ity | Once committed, changes are permanent | Power failure won't cause committed transaction to be lost |

Diagram (ACID):



Example:

```

1  -- ATM Withdrawal Transaction
2  BEGIN TRANSACTION;
3  -- Check balance
4  SELECT balance FROM accounts WHERE account_id = 'A123';
5
6  -- If sufficient, update balance
7  UPDATE accounts SET balance = balance - 100 WHERE account_id = 'A123';
8
9  -- Record the withdrawal
10 INSERT INTO transactions (account_id, type, amount, date)
11 VALUES ('A123', 'WITHDRAWAL', 100, SYSDATE);
12
13 -- If all operations successful
14 COMMIT;
15 -- If any operation fails
16 -- ROLLBACK;
17 END TRANSACTION;

```

Mnemonic

[title=ACID]Atomicity Consistency Isolation Durability

Question 5(b) OR [4 marks]

Write the Queries using set operators...

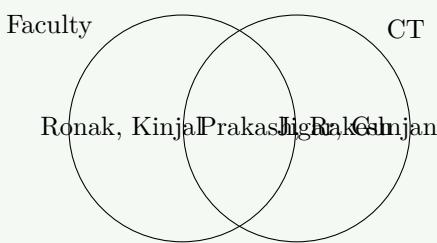
Solution**Queries:**

1. Either Faculty or CT (UNION)
2. Both Faculty and CT (INTERSECT)
3. Only Faculty (MINUS)
4. Only CT (MINUS)

```

1  -- 1. UNION
2  SELECT FacultyName FROM Faculty UNION SELECT CTName FROM CT;
3
4  -- 2. INTERSECT
5  SELECT FacultyName FROM Faculty INTERSECT SELECT CTName FROM CT;
6
7  -- 3. MINUS (Fac - CT)
8  SELECT FacultyName FROM Faculty MINUS SELECT CTName FROM CT;
9
10 -- 4. MINUS (CT - Fac)
11 SELECT CTName FROM CT MINUS SELECT FacultyName FROM Faculty;

```



Mnemonic

[title=UIMM]Union Intersect Minus Minus

Question 5(c) OR [7 marks]

Explain View Serializability in detail.

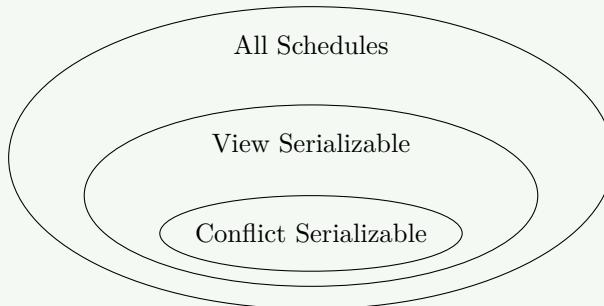
Solution

View Serializability: A schedule is view serializable if it is view equivalent to some serial schedule, meaning it produces the same "view" (or final state) of the database.

Table: Comparison with Conflict Serializability

| Aspect | View Serializability | Conflict Serializability |
|------------|--|--|
| Definition | Based on the final results of reads and writes | Based on conflicts between operations |
| Condition | Preserves initial read, final write, and read-write dependency | Preserves all conflicts between operations |
| Scope | Broader class of schedules | Subset of view serializable schedules |
| Testing | More complex to test | Can test with precedence graph |

Diagram (Subset):

**View Equivalence Conditions:**

- Initial Reads: If T1 reads an initial value of data item A in schedule S1, it must also read the initial value in S2.
- Final Writes: If T1 performs the final write on data item A in S1, it must also perform the final write in S2.
- Read-Write Dependency: If T1 reads a value of A written by T2 in S1, it must also read the value written by T2 in S2.

Example of View Serializable but not Conflict Serializable Schedule: Consider transactions with blind writes (writes without reading):

- T1: W1(A)
- T2: W2(A)

Schedule S: W1(A), W2(A) - View serializable to both T1→T2 and T2→T1 (final write is always T2) But W1(A) and W2(A) conflict, so a conflict graph would have an edge in both directions.

Mnemonic

[title=IRF]Initial reads, Result writes, Final view