

# Database Management System (1333204) - Summer 2025 Solution

Milav Dabgar

May 17, 2025

## Question 1(a) [3 marks]

Write a short note: Data Dictionary

### Solution

A **Data Dictionary** is a centralized repository that stores metadata about database structure, elements, and relationships.

**Table 1.** Data Dictionary Components

Component	Description
Table Names	List of all tables in database
Column Details	Data types, constraints, lengths
Relationships	Foreign key connections
Indexes	Performance optimization structures

### Key Features:

- **Metadata Storage:** Contains information about data structure
- **Data Integrity:** Maintains consistency rules and constraints
- **Documentation:** Provides comprehensive database documentation

### Mnemonic

**Mnemonic:** "Data Dictionary Delivers Details"

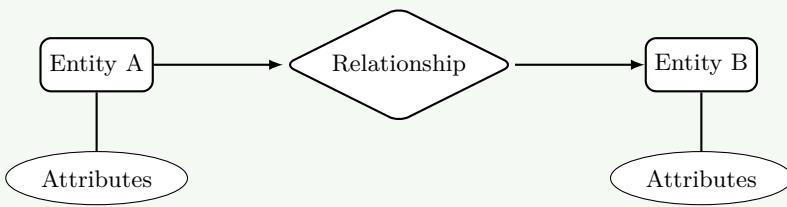
## Question 1(b) [4 marks]

Define (i) E-R model (ii) Entity (iii) Entity set and (iv) attributes

### Solution

**Table 2.** ER Model Definitions

Term	Definition
E-R Model	Conceptual data model using entities and relationships
Entity	Real-world object with independent existence
Entity Set	Collection of similar entities of same type
Attributes	Properties that describe entity characteristics

**Figure 1.** ER Model Components**Key Points:**

- **Conceptual Design:** High-level database design approach
- **Visual Representation:** Uses diagrams for clear understanding

**Mnemonic**

**Mnemonic:** "Entities Relate Meaningfully"

**Question 1(c) [7 marks]****Explain Advantages of DBMS****Solution****Table 3.** DBMS Advantages

Advantage	Benefit
<b>Data Independence</b>	Applications isolated from data structure changes
<b>Data Sharing</b>	Multiple users access same data simultaneously
<b>Data Security</b>	Access control and authentication mechanisms
<b>Data Integrity</b>	Consistency maintained through constraints
<b>Backup &amp; Recovery</b>	Automatic data protection and restoration
<b>Reduced Redundancy</b>	Eliminates duplicate data storage

**Key Benefits:**

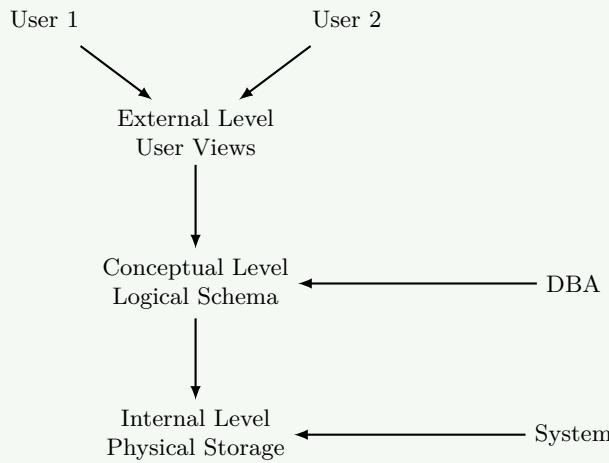
- **Centralized Control:** Single point of data management
- **Cost Effectiveness:** Reduces development and maintenance costs
- **Data Consistency:** Ensures uniform data across applications
- **Concurrent Access:** Multiple users can work simultaneously
- **Query Optimization:** Efficient data retrieval mechanisms

**Mnemonic**

**Mnemonic:** "Database Benefits Business Better"

**Question 1(c) OR [7 marks]****Explain Architecture of DBMS**

### Solution



**Figure 2.** Three-Level DBMS Architecture

**Table 4.** Architecture Levels

Level	Purpose	Users
External	Individual user views	End users, Applications
Conceptual	Complete logical structure	Database Administrator
Internal	Physical storage details	System programmers

#### Key Features:

- **Data Independence:** Changes at one level don't affect others
- **Security:** Different access levels for different users
- **Abstraction:** Hides complexity from users

#### Mnemonic

**Mnemonic:** "External Conceptual Internal Architecture"

## Question 2(a) [3 marks]

Explain UNIQUE KEY and PRIMARY KEY

### Solution

**Table 5.** Key Comparison

Feature	PRIMARY KEY	UNIQUE KEY
Null Values	Not allowed	One null allowed
Number per Table	Only one	Multiple allowed
Index Creation	Automatic clustered	Automatic non-clustered
Purpose	Entity identification	Data uniqueness

#### Key Differences:

- **Primary Key:** Uniquely identifies each record, cannot be null
- **Unique Key:** Ensures uniqueness but allows one null value

**Mnemonic**

**Mnemonic:** "Primary Prevents Nulls, Unique Understands Nulls"

**Question 2(b) [4 marks]**

Write a short note on Participation of Entity in ER diagram

**Solution**

**Table 6.** Participation Types

Type	Description	Symbol
<b>Total Participation</b>	Every entity must participate	Double line
<b>Partial Participation</b>	Some entities may not participate	Single line



**Figure 3.** Participation Example

**Key Concepts:**

- **Mandatory Participation:** Every instance must be involved
- **Optional Participation:** Some instances may not be involved
- **Business Rules:** Reflects real-world constraints

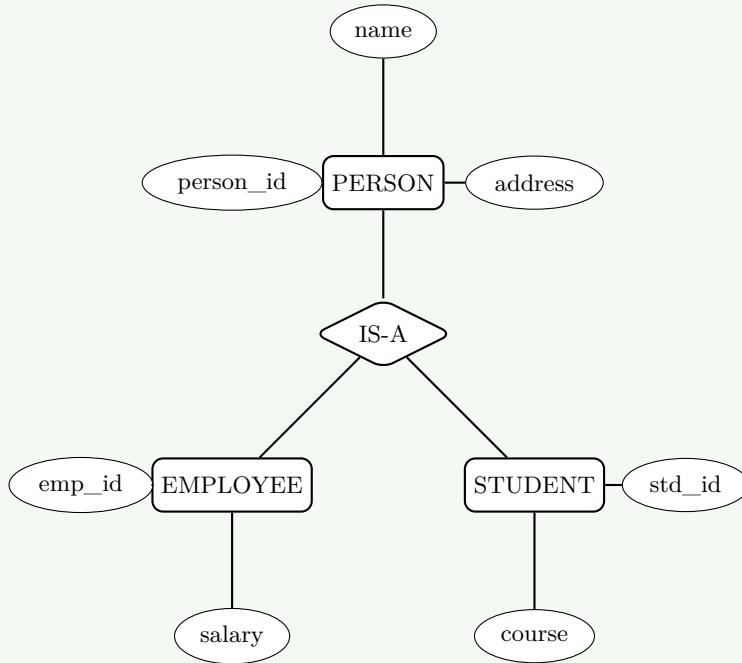
**Mnemonic**

**Mnemonic:** "Total Participation Requires All"

**Question 2(c) [7 marks]**

Describe Generalization concept in Detail for ER diagram

### Solution



**Figure 4.** Generalization Example

**Table 7.** Generalization Characteristics

Aspect	Description
<b>Bottom-up Process</b>	Combines similar entities into superclass
<b>Inheritance</b>	Subclasses inherit superclass attributes
<b>Specialization</b>	Reverse process of generalization
<b>Overlap Constraints</b>	Disjoint or overlapping subclasses

#### Key Features:

- **Attribute Inheritance:** Common attributes moved to superclass
- **Relationship Inheritance:** Relationships also inherited
- **Constraint Types:** Total/partial, disjoint/overlapping
- **ISA Relationship:** Represents "is-a" connection

#### Mnemonic

**Mnemonic:** "Generalization Groups Similar Entities"

## Question 2(a) OR [3 marks]

Explain Mapping Cardinality in ER diagram

**Solution****Table 8.** Cardinality Types

Type	Description	Example
<b>One-to-One (1:1)</b>	One entity relates to one other	Person-Passport
<b>One-to-Many (1:M)</b>	One entity relates to many others	Department-Employee
<b>Many-to-One (M:1)</b>	Many entities relate to one	Employee-Department
<b>Many-to-Many (M:N)</b>	Many entities relate to many	Student-Course

**Key Concepts:**

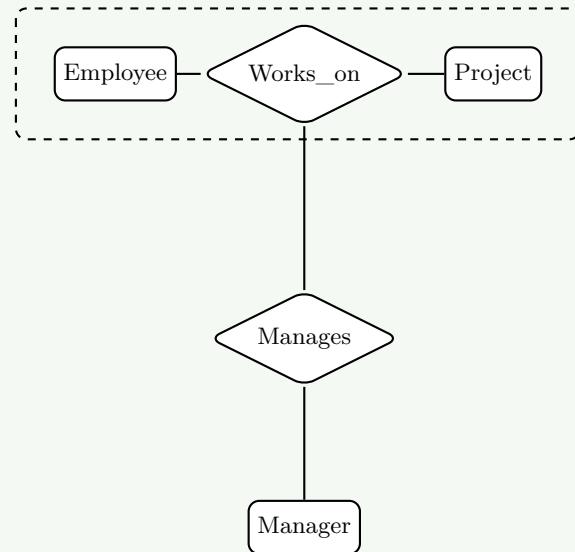
- **Relationship Constraints:** Defines how entities can be related
- **Business Rules:** Reflects real-world relationship limits

**Mnemonic**

**Mnemonic:** "One Or Many Mappings Matter"

**Question 2(b) OR [4 marks]**

Explain Aggregation in E-R diagram

**Solution****Figure 5.** Aggregation Example**Key Features:**

- **Relationship as Entity:** Treats relationship set as entity
- **Higher-level Relationships:** Allows relationships between relationships
- **Complex Modeling:** Handles advanced business scenarios
- **Abstraction Mechanism:** Simplifies complex relationships

**Table 9.** Aggregation Benefits

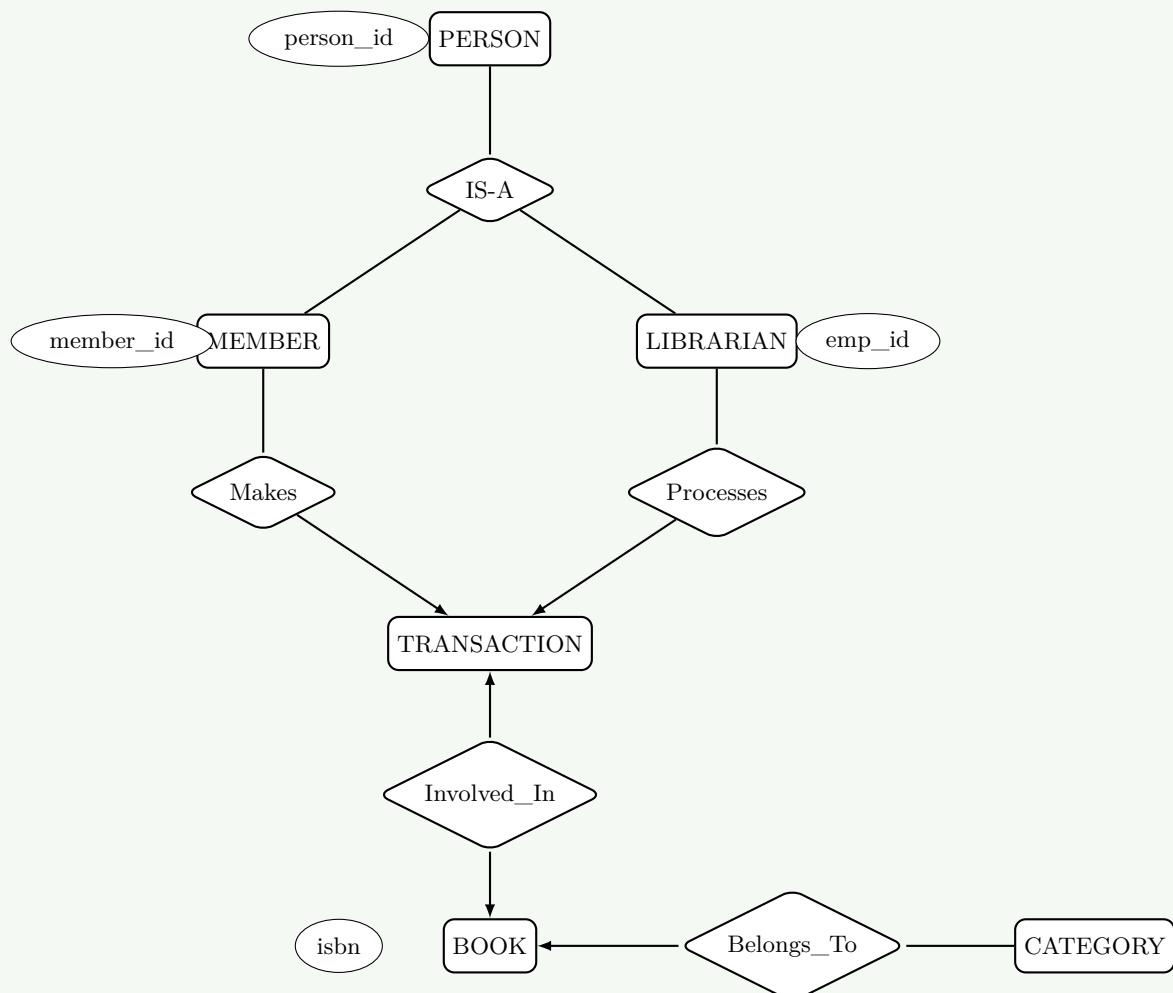
Benefit	Description
<b>Modeling Flexibility</b>	Handles complex relationships
<b>Semantic Clarity</b>	Clear representation of business rules
<b>Design Simplicity</b>	Reduces model complexity

**Mnemonic**

**Mnemonic:** "Aggregation Abstracts Advanced Associations"

**Question 2(c) OR [7 marks]**

Draw ER diagram of Library Management system using Enhanced ER model

**Solution****Figure 6.** Library Management System**Enhanced ER Features Used:**

- **Generalization:** Person superclass with Member and Librarian subclasses
- **Specialization:** Different attributes for different person types

- **Aggregation:** Transaction relationship involving multiple entities
- **Multiple Inheritance:** Complex relationship handling

**Mnemonic**

**Mnemonic:** "Library Links Literature Logically"

**Question 3(a) [3 marks]**

Explain SQL data types

**Solution**

**Table 10.** Common SQL Data Types

Category	Data Type	Description
Numeric	INT, DECIMAL, FLOAT	Store numbers
Character	CHAR, VARCHAR, TEXT	Store text
Date/Time	DATE, TIME, DATETIME	Store temporal data
Boolean	BOOLEAN	Store true/false

**Key Points:**

- **Data Integrity:** Ensures correct data storage
- **Storage Optimization:** Appropriate size allocation
- **Validation:** Automatic data type checking

**Mnemonic**

**Mnemonic:** "Data Types Define Storage"

**Question 3(b) [4 marks]**

Compare DROP and TRUNCATE commands

**Solution**

**Table 11.** DROP vs TRUNCATE Comparison

Feature	DROP	TRUNCATE
<b>Operation</b>	Removes table structure	Removes all data only
<b>Rollback</b>	Cannot rollback	Can rollback (in transaction)
<b>Speed</b>	Slower	Faster
<b>Triggers</b>	Fires triggers	Does not fire triggers
<b>Where Clause</b>	Not applicable	Not supported
<b>Auto-increment</b>	Resets	Resets to initial value

**Code Examples:**

```

1   -- DROP command
2   DROP TABLE student;
3
4   -- TRUNCATE command
5   TRUNCATE TABLE student;

```

**Key Differences:**

- Structure Impact:** DROP removes everything, TRUNCATE keeps structure
- Performance:** TRUNCATE is faster for large tables

**Mnemonic**

**Mnemonic:** "DROP Destroys, TRUNCATE Trims"

**Question 3(c) [7 marks]**

Consider a following Relational Schema and give Relational Algebra Expression for the following Queries  
**Students** (Name, SPI, DOB, Enrollment No)

**Solution****Relational Algebra Expressions:**

- i) List out all students whose SPI is lower than 6.0:

$$\sigma_{SPI < 6.0}(Students)$$

- ii) List name of student whose enrollment number contains 006:

$$\pi_{Name}(\sigma_{Enrollment\_No \text{ LIKE } '%006%'}(Students))$$

- iii) List all students with same DOB:

$$Students \bowtie_{Students.DOB=S2.DOB \wedge Students.Enrollment\_No \neq S2.Enrollment\_No} (\rho_{S2}(Students))$$

- iv) Display students name starting from same letter:

$$\pi_{Name}(Students \bowtie_{SUBSTR(Students.Name,1,1)=SUBSTR(S2.Name,1,1) \wedge Students.Enrollment\_No \neq S2.Enrollment\_No} (\rho_{S2}(Students)))$$

**Table 12.** Relational Algebra Operators Used

Operator	Symbol	Purpose
Selection	$\sigma$	Filter rows based on condition
Projection	$\pi$	Select specific columns
Join	$\bowtie$	Combine related tuples
Rename	$\rho$	Rename relations/attributes

**Mnemonic**

**Mnemonic:** "Select Project Join Rename"

## Question 3(a) OR [3 marks]

Explain use of Grant and Revoke command with example

### Solution

#### Code Examples:

```

1  -- GRANT command
2  GRANT SELECT, INSERT ON student TO user1;
3  GRANT ALL PRIVILEGES ON database1 TO user2;
4
5  -- REVOKE command
6  REVOKE INSERT ON student FROM user1;
7  REVOKE ALL PRIVILEGES ON database1 FROM user2;
```

#### Key Features:

- Access Control:** Manages user permissions
- Security:** Prevents unauthorized access
- Granular Control:** Specific privilege assignment

Table 13. Common Privileges

Privilege	Description
<b>SELECT</b>	Read data
<b>INSERT</b>	Add new records
<b>UPDATE</b>	Modify existing data
<b>DELETE</b>	Remove records
<b>ALL</b>	Complete access

#### Mnemonic

**Mnemonic:** "Grant Gives, Revoke Removes"

## Question 3(b) OR [4 marks]

Describe DML commands with Example

### Solution

Table 14. DML Commands

Command	Purpose	Example
<b>INSERT</b>	Add new records	INSERT INTO student...
<b>UPDATE</b>	Modify existing data	UPDATE student SET...
<b>DELETE</b>	Remove records	DELETE FROM student...
<b>SELECT</b>	Retrieve data	SELECT * FROM student...

#### Code Examples:

```

1  -- INSERT command
2  INSERT INTO Students (name, spi, dob)
```

```

3 | VALUES ('Alice', 8.5, '2000-05-15');
4 |
5 | -- UPDATE command
6 | UPDATE Students SET spi = 9.0
7 | WHERE name = 'Alice';
8 |
9 | -- DELETE command
10| DELETE FROM Students
11| WHERE spi < 6.0;
12|
13| -- SELECT command
14| SELECT name, spi FROM Students
15| WHERE spi > 8.0;

```

**Key Features:**

- **Data Manipulation:** Core database operations
- **Transaction Support:** Can be rolled back
- **Conditional Operations:** WHERE clause support

**Mnemonic**

**Mnemonic:** "Insert Update Delete Select"

**Question 3(c) OR [7 marks]**

List all Conversion function of DBMS and explain any three of them in detail

**Solution**

**Table 15.** Conversion Functions

Function	Purpose	Example
TO_CHAR	Convert to character	TO_CHAR(sysdate)
TO_DATE	Convert to date	TO_DATE('15-05-2025')
TO_NUMBER	Convert to number	TO_NUMBER('123.45')
CAST	General conversion	CAST('123' AS INT)
CONVERT	Data type conversion	CONVERT(varchar, 123)

**Detailed Explanation of Three Functions:****1. TO\_CHAR Function:**

- **Purpose:** Converts dates and numbers to character strings
- **Syntax:** TO\_CHAR(value, format)
- **Usage:** Date formatting, number formatting with specific patterns

**2. TO\_DATE Function:**

- **Purpose:** Converts character strings to date values
- **Syntax:** TO\_DATE(string, format)
- **Usage:** String to date conversion with specified format

**3. TO\_NUMBER Function:**

- **Purpose:** Converts character strings to numeric values
- **Syntax:** TO\_NUMBER(string, format)
- **Usage:** String to number conversion for calculations

**Key Benefits:**

- **Data Type Flexibility:** Seamless conversion between types

- **Format Control:** Specific formatting options
- **Error Handling:** Validation during conversion

**Mnemonic**

**Mnemonic:** "Convert Characters Dates Numbers"

**Question 4(a) [3 marks]**

Write short note: Domain Integrity Constraint

**Solution**

**Domain Integrity Constraints** ensure that data values fall within acceptable ranges and formats for specific attributes.

**Table 16.** Domain Constraint Types

Constraint	Purpose	Example
<b>CHECK</b>	Value range validation	<code>CHECK (age &gt;= 0 AND age &lt;= 100)</code>
<b>NOT NULL</b>	Prevents null values	<code>name VARCHAR(50) NOT NULL</code>
<b>DEFAULT</b>	Sets default values	<code>status VARCHAR(10) DEFAULT 'Active'</code>

**Key Features:**

- **Data Validation:** Ensures data quality at entry
- **Business Rules:** Implements domain-specific rules
- **Automatic Checking:** Validation occurs during DML operations

**Mnemonic**

**Mnemonic:** "Domain Defines Data Boundaries"

**Question 4(b) [4 marks]**

List all JOIN in DBMS and explain any two

**Solution**

**Table 17.** Types of JOINS

JOIN Type	Description
<b>INNER JOIN</b>	Returns matching records from both tables
<b>LEFT JOIN</b>	Returns all records from left table
<b>RIGHT JOIN</b>	Returns all records from right table
<b>FULL OUTER JOIN</b>	Returns all records from both tables
<b>CROSS JOIN</b>	Cartesian product of both tables
<b>SELF JOIN</b>	Table joined with itself

**Detailed Explanation:**

### 1. INNER JOIN:

```

1 SELECT s.name, c.course_name
2 FROM students s
3 INNER JOIN courses c ON s.course_id = c.course_id;

```

- Returns only matching records from both tables
- Most commonly used join type

### 2. LEFT JOIN:

```

1 SELECT s.name, c.course_name
2 FROM students s
3 LEFT JOIN courses c ON s.course_id = c.course_id;

```

- Returns all students, even if no course assigned
- NULL values for unmatched records

#### Mnemonic

**Mnemonic:** "Join Tables Together Thoughtfully"

## Question 4(c) [7 marks]

Explain Concept of Functional Dependency in detail

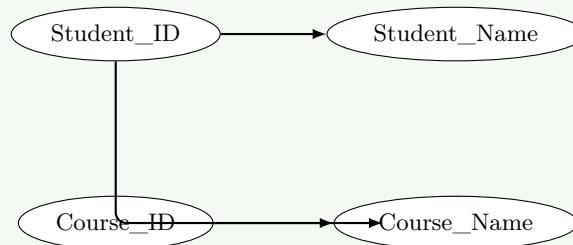
#### Solution

**Functional Dependency** occurs when the value of one attribute uniquely determines the value of another attribute.

**Notation:**  $A \rightarrow B$  (A functionally determines B)

**Table 18.** Types of Functional Dependencies

Type	Definition	Example
<b>Full FD</b>	All attributes in LHS needed	$\{Student\_ID, Course\_ID\} \rightarrow Grade$
<b>Partial FD</b>	Some LHS attributes redundant	$\{Student\_ID, Course\_ID\} \rightarrow Student\_Name$
<b>Transitive FD</b>	Indirect dependency through another attribute	$Student\_ID \rightarrow Dept\_ID \rightarrow Dept\_Name$



**Figure 7.** Functional Dependency Example

#### Key Properties:

- **Reflexivity:**  $A \rightarrow A$  (trivial dependency)
- **Augmentation:** If  $A \rightarrow B$ , then  $AC \rightarrow BC$
- **Transitivity:** If  $A \rightarrow B$  and  $B \rightarrow C$ , then  $A \rightarrow C$
- **Decomposition:** If  $A \rightarrow BC$ , then  $A \rightarrow B$  and  $A \rightarrow C$

**Mnemonic**

**Mnemonic:** "Functions Determine Dependencies Directly"

**Question 4(a) OR [3 marks]**

Write short note: Referential integrity Constraints

**Solution**

**Referential Integrity** ensures that foreign key values in one table correspond to existing primary key values in referenced table.

**Table 19.** Referential Integrity Rules

Rule	Description	Action
<b>INSERT Rule</b>	Foreign key must exist in parent	Reject invalid inserts
<b>DELETE Rule</b>	Handle parent record deletion	CASCADE, RESTRICT, SET NULL
<b>UPDATE Rule</b>	Handle primary key updates	CASCADE, RESTRICT

**Code Example:**

```

1 ALTER TABLE Orders
2 ADD CONSTRAINT FK_Customer
3 FOREIGN KEY (customer_id)
4 REFERENCES Customers(customer_id);

```

**Key Features:**

- **Foreign Key Constraint:** Links related tables
- **Data Consistency:** Prevents orphaned records
- **Relationship Maintenance:** Preserves table relationships

**Mnemonic**

**Mnemonic:** "References Require Related Records"

**Question 4(b) OR [4 marks]**

Explain union and intersection operations of relational algebra

**Solution**

**Table 20.** Set Operations Comparison

Operation	Symbol	Description	Requirement
<b>UNION</b>	$\cup$	Combines all tuples from both relations	Union compatible
<b>INTERSECTION</b>	$\cap$	Common tuples in both relations	Union compatible

**Union Operation:**

- **Syntax:**  $R \cup S$

- Result:** All tuples from R and S (duplicates removed)
- Requirement:** Same number and types of attributes

#### Intersection Operation:

- Syntax:**  $R \cap S$
- Result:** Tuples that exist in both R and S
- Requirement:** Union compatible relations

#### Example:

```

1 Students_CS U Students_IT = All students from both departments
2 Students_CS n Students_IT = Students in both departments

```

#### Mnemonic

**Mnemonic:** "Union Unites, Intersection Identifies Common"

## Question 4(c) OR [7 marks]

Explain Concept of Normalization in DBMS in detail

#### Solution

**Normalization** is the process of organizing database tables to minimize data redundancy and improve data integrity.

**Table 21.** Normal Forms

Normal Form	Requirements	Eliminates
<b>1NF</b>	Atomic values, no repeating groups	Multivalued attributes
<b>2NF</b>	1NF + No partial dependencies	Partial functional dependencies
<b>3NF</b>	2NF + No transitive dependencies	Transitive dependencies
<b>BCNF</b>	3NF + Every determinant is candidate key	Remaining anomalies

#### Normalization Process:

##### Step 1 - First Normal Form (1NF):

- Eliminate repeating groups
- Each cell contains single value
- Each record is unique

##### Step 2 - Second Normal Form (2NF):

- Must be in 1NF
- Remove partial dependencies
- Non-key attributes fully dependent on primary key

##### Step 3 - Third Normal Form (3NF):

- Must be in 2NF
- Remove transitive dependencies
- Non-key attributes not dependent on other non-key attributes

#### Benefits of Normalization:

- Reduced Redundancy:** Eliminates duplicate data
- Data Integrity:** Maintains consistency
- Storage Efficiency:** Minimizes storage space
- Update Anomalies:** Prevents inconsistent updates

**Mnemonic**

**Mnemonic:** "Normalize to Neat, Non-redundant Tables"

**Question 5(a) [3 marks]**

Describe Need of Normalization in DBMS

**Solution**

**Table 22.** Problems Solved by Normalization

Problem	Description	Solution
<b>Insertion Anomaly</b>	Cannot insert data without complete info	Separate tables
<b>Update Anomaly</b>	Multiple updates for single change	Remove redundancy
<b>Deletion Anomaly</b>	Loss of important data when deleting	Preserve dependencies

**Key Needs:**

- **Data Consistency:** Ensures uniform data across database
- **Storage Optimization:** Reduces redundant storage
- **Maintenance Simplicity:** Easier database updates

**Benefits:**

- **Improved Data Quality:** Reduces errors and inconsistencies
- **Flexible Design:** Easier to modify and extend
- **Better Performance:** For update operations

**Mnemonic**

**Mnemonic:** "Normalization Needs Neat Organization"

**Question 5(b) [4 marks]**

Explain properties of Transaction in DBMS

**Solution**

**Table 23.** ACID Properties

Property	Description	Purpose
<b>Atomicity</b>	All operations succeed or all fail	Ensures completeness
<b>Consistency</b>	Database remains in valid state	Maintains integrity
<b>Isolation</b>	Concurrent transactions don't interfere	Prevents conflicts
<b>Durability</b>	Committed changes are permanent	Ensures persistence

**Detailed Explanation:**

- **Atomicity:** Transaction is indivisible unit. Either all operations complete or none.
- **Consistency:** Database transitions from one valid state to another. All integrity constraints maintained.

- **Isolation:** Concurrent transactions appear to run sequentially. Intermediate states not visible to other transactions.
- **Durability:** Once committed, changes survive system failures. Data permanently stored.

**Mnemonic****Mnemonic:** "ACID Assures Correct Database"**Question 5(c) [7 marks]****Explain View Serializability in detail****Solution**

**View Serializability** determines if a concurrent schedule produces the same result as some serial schedule by examining read and write operations.

**Table 24.** View Equivalence Conditions

Condition	Description
<b>Initial Reads</b>	Same transactions read initial values
<b>Final Writes</b>	Same transactions perform final writes
<b>Intermediate Reads</b>	Read values from same writing transactions

**Key Concepts:**

**View Equivalent Schedules:** Two schedules are view equivalent if:

1. For each data item, if transaction T reads initial value in one schedule, it reads initial value in other
2. For each read operation, if T reads value written by T' in one schedule, same holds in other
3. For each data item, if T performs final write in one schedule, it performs final write in other

**Testing View Serializability:**

- **Precedence Graph:** Create directed graph
- **Cycle Detection:** Check for cycles in graph
- **Conflict Analysis:** Examine read-write conflicts

**Example Analysis:**

1 Schedule S1: R1(X) W1(X) R2(X) W2(X)  
 2 Schedule S2: R1(X) R2(X) W1(X) W2(X)

**Comparison with Conflict Serializability:**

- View serializability is less restrictive
- Some view serializable schedules are not conflict serializable
- More complex to test

**Mnemonic****Mnemonic:** "View Verifies Valid Schedules"**Question 5(a) OR [3 marks]****Perform 2NF on any Database**

## Solution

### Example: Student Course Database

Original Table (Not in 2NF):

```
1 Student_Course (Student_ID, Student_Name, Course_ID, Course_Name, Grade, Instructor)
2 Primary Key: {Student_ID, Course_ID}
```

#### Functional Dependencies:

- $\text{Student\_ID} \rightarrow \text{Student\_Name}$  (Partial dependency)
- $\text{Course\_ID} \rightarrow \text{Course\_Name}, \text{Instructor}$  (Partial dependency)
- $\{\text{Student\_ID}, \text{Course\_ID}\} \rightarrow \text{Grade}$

#### 2NF Decomposition:

Table 1: Students

```
1 Students (Student_ID, Student_Name)
2 Primary Key: Student_ID
```

Table 2: Courses

```
1 Courses (Course_ID, Course_Name, Instructor)
2 Primary Key: Course_ID
```

Table 3: Enrollments

```
1 Enrollments (Student_ID, Course_ID, Grade)
2 Primary Key: {Student_ID, Course_ID}
3 Foreign Keys: Student_ID -> Students, Course_ID -> Courses
```

**Result:** All partial dependencies eliminated, now in 2NF.

### Mnemonic

**Mnemonic:** "Second Normal Form Separates Dependencies"

## Question 5(b) OR [4 marks]

### Explain States of Transaction

## Solution

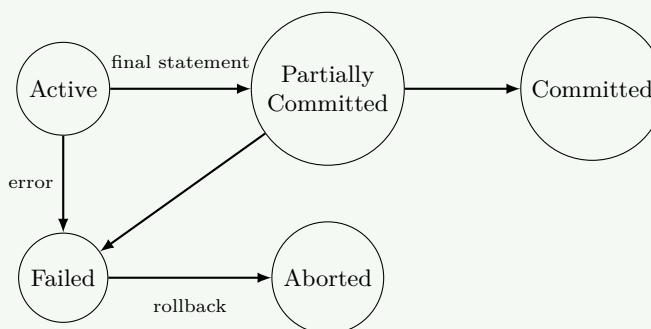


Figure 8. Transaction State Diagram

**Table 25.** Transaction States

State	Description	Actions
<b>Active</b>	Transaction is executing	Read/Write operations
<b>Partially Committed</b>	Final statement executed	Waiting for commit
<b>Committed</b>	Transaction completed successfully	Changes permanent
<b>Failed</b>	Cannot proceed normally	Error occurred
<b>Aborted</b>	Transaction rolled back	All changes undone

**Key Points:**

- Recovery:** System can recover from failed states
- Durability:** Committed changes are permanent
- Atomicity:** Aborted transactions leave no trace

**Mnemonic**

**Mnemonic:** "Transactions Travel Through States"

**Question 5(c) OR [7 marks]****Explain Conflict Serializability in detail****Solution**

**Conflict Serializability** ensures that a concurrent schedule is equivalent to some serial schedule by analyzing conflicting operations.

**Table 26.** Conflicting Operations

Operation Pair	Conflict Type	Reason
Read-Write	RW Conflict	Read before write
Write-Read	WR Conflict	Write before read
Write-Write	WW Conflict	Multiple writes

**Testing Conflict Serializability:****Step 1: Identify Conflicts**

- Find pairs of operations on same data item
- Check if operations belong to different transactions
- Determine if operations conflict

**Step 2: Create Precedence Graph**

- Nodes represent transactions
- Directed edges represent conflicts
- Edge from  $T_i$  to  $T_j$  if  $T_i$  conflicts with  $T_j$

**Step 3: Check for Cycles**

- If graph has no cycles  $\rightarrow$  Conflict serializable
- If graph has cycles  $\rightarrow$  Not conflict serializable

**Example Analysis:**

1 Schedule: R1(A) W1(A) R2(A) W2(B) R1(B) W1(B)  
 2 Conflicts:  
 3 - W1(A) conflicts with R2(A)  $\rightarrow$  T1 before T2  
 4 - W2(B) conflicts with R1(B)  $\rightarrow$  T2 before T1

5 - W2(B) conflicts with W1(B)  $\rightarrow$  T2 before T1

**Precedence Graph:**

**Figure 9.** Precedence Graph (Cycle)

**Result:** Contains cycle, therefore NOT conflict serializable.

**Comparing with View Serializability:**

- Conflict serializability is more restrictive
- All conflict serializable schedules are view serializable
- Easier to test than view serializability

**Mnemonic**

**Mnemonic:** "Conflicts Create Cycles, Check Carefully"