

# ER Model and Relational Algebra

Database Management System - Unit II

DIPLOMA ENGINEERING

INFORMATION & COMMUNICATION TECHNOLOGY

# Course Overview

**Program**

Diploma Engineering

**Branch**

Information & Communication Technology

**Subject Code**

DI04032011

**Subject**

Database Management System

This unit focuses on Entity-Relationship modeling and the transformation of conceptual designs into functional database structures. Understanding these concepts is fundamental to designing efficient, scalable database systems.

# Unit II Learning Objectives

01

## Master ER Model Fundamentals

Understand entities, attributes, and relationships that form the foundation of database design

02

## Apply Mapping Cardinality

Learn to define and implement one-to-one, one-to-many, and many-to-many relationships

03

## Design ER Diagrams

Create comprehensive visual representations of database structures using standard notation

04

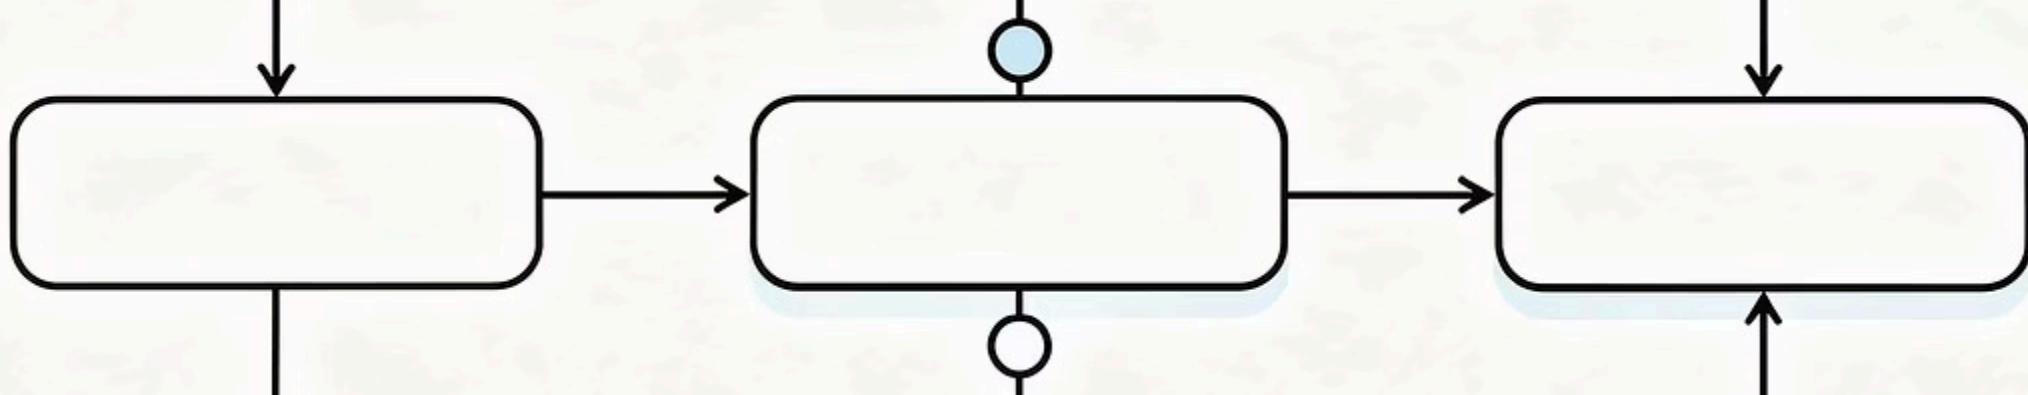
## Implement Enhanced Features

Work with advanced concepts including generalization, specialization, and aggregation

05

## Convert to Database Schema

Transform ER diagrams into functional relational database implementations



CHAPTER 2.1

## Basic Concepts of Entity-Relationship Model

The Entity-Relationship (ER) model is a high-level conceptual data model that helps define the data elements and relationships for a specified system. Developed by Peter Chen in 1976, it provides a graphical representation of the logical structure of databases.

The ER model serves as a blueprint for database design, allowing developers and stakeholders to visualize how data is organized and related before implementation. This approach reduces errors, improves communication, and ensures the database meets business requirements effectively.

# The Three Pillars of ER Modeling



## Entities

Objects or things in the real world that are distinguishable from other objects



## Attributes

Properties or characteristics that describe entities



## Relationships

Associations or connections between two or more entities

These three fundamental components work together to create a complete representation of the data structure. Mastering each pillar is essential for effective database design and ensures that all aspects of the real-world scenario are accurately captured in the model.

# Understanding Entities

## What is an Entity?

An entity is a real-world object or concept that can be distinctly identified and about which data can be stored. Entities represent the "nouns" of your database—the things you want to track.

### Key Characteristics:

- Must be distinguishable from other objects
- Has a set of associated attributes
- Can be physical or conceptual
- Represented by rectangles in ER diagrams

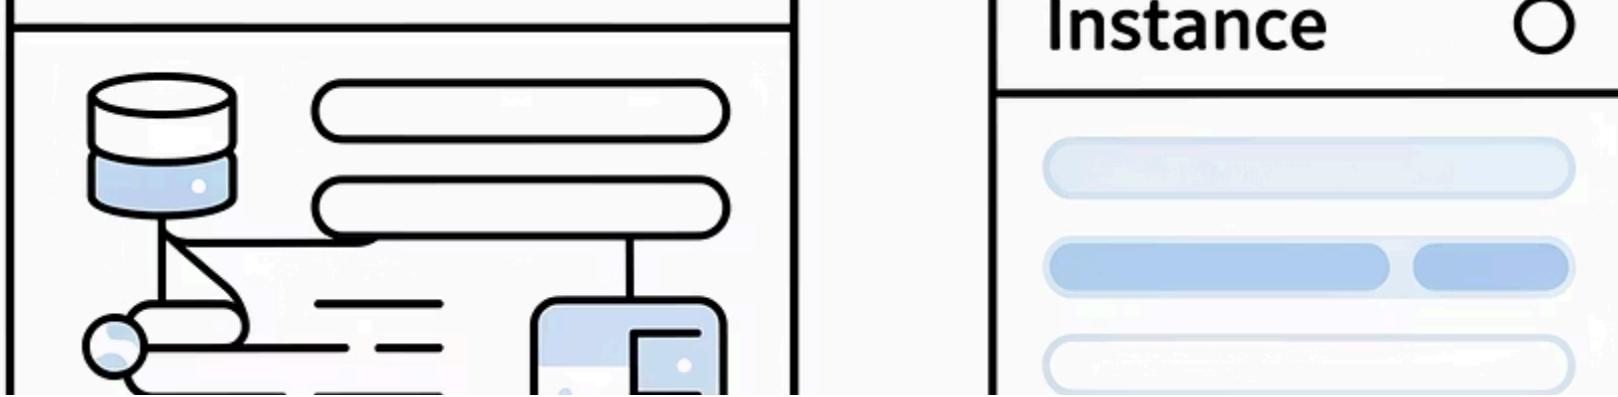
## Entity Examples

### Physical Entities:

- Student
- Employee
- Product
- Vehicle

### Conceptual Entities:

- Course
- Department
- Order
- Account



# Entity Types vs Entity Instances

## Entity Type

A category or classification of entities that share common properties. It defines the structure and attributes that all instances will have.

**Example:** STUDENT is an entity type that represents all students in a college database.

## Entity Instance

A specific occurrence or example of an entity type. Each instance has unique values for its attributes.

**Example:** "John Smith, ID: 12345" is a specific instance of the STUDENT entity type.

Think of entity types as the template or blueprint, while entity instances are the actual data entries that follow that template. A single entity type can have thousands or millions of instances in a real database.

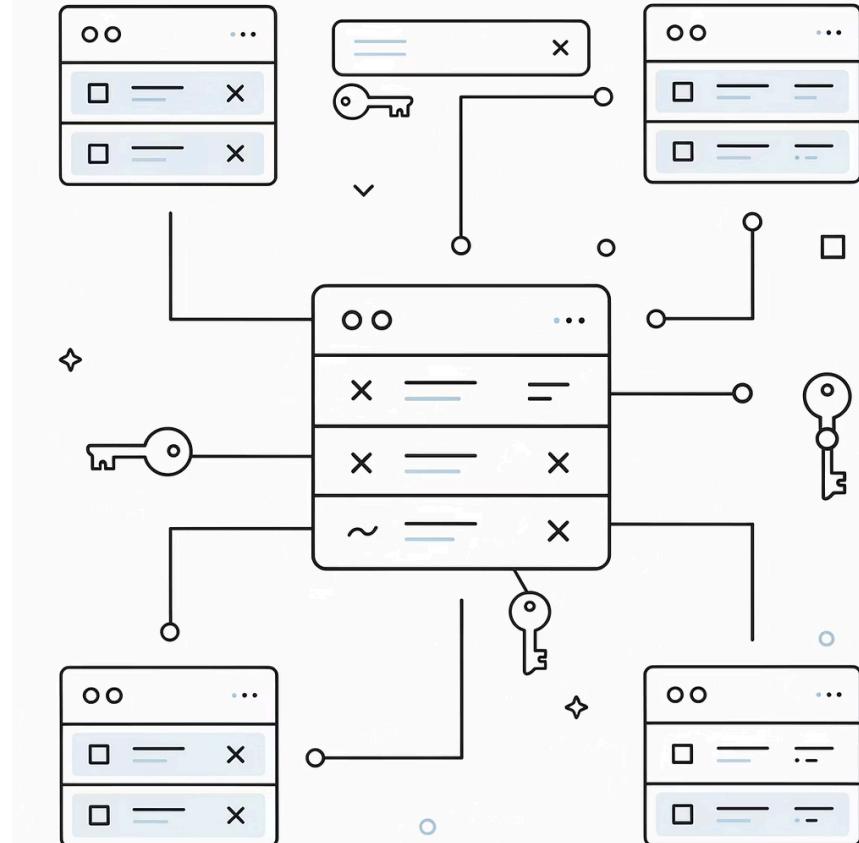
SECTION 2.1.2

# Attributes: Describing Entity Properties

Attributes are the properties or characteristics that describe an entity. They represent the data we want to store about each entity instance. Every entity type has a set of attributes that define what information is captured.

For example, a STUDENT entity might have attributes such as Student\_ID, Name, Date\_of\_Birth, Email, and Phone\_Number. Each attribute holds a specific piece of information that helps fully describe the student.

Choosing the right attributes is crucial for database design. Include too few, and you won't capture necessary information. Include too many, and you create unnecessary complexity and storage overhead.



# Types of Attributes



## Simple Attributes

Atomic attributes that cannot be divided further. Examples: Age, Gender, Student\_ID



## Composite Attributes

Attributes that can be divided into smaller sub-parts. Example: Full\_Name (First\_Name, Middle\_Name, Last\_Name) or Address (Street, City, State, ZIP)



## Multi-valued Attributes

Attributes that can have multiple values for a single entity. Examples: Phone\_Numbers, Email\_Addresses, Skills



## Derived Attributes

Attributes whose values can be calculated from other attributes. Examples: Age (from Date\_of\_Birth), Total\_Price (from Quantity  $\times$  Unit\_Price)

# Attribute Properties and Domain

## Key Concepts

Each attribute has specific properties that define how it stores and represents data:

- **Domain:** The set of allowed values for an attribute (e.g., Age domain might be 0-120)
- **Null Values:** Whether an attribute can be empty or must have a value
- **Default Values:** Predefined values assigned when no specific value is provided
- **Data Type:** The kind of data stored (integer, string, date, boolean, etc.)

Understanding these properties ensures data integrity and prevents invalid data entry.

For instance, defining Age as an integer with domain 0-120 prevents nonsensical values like -5 or 500.



### Domain Example

**Attribute:** Grade

**Domain:** {'A', 'B', 'C', 'D', 'F'}

**Data Type:** Character

**Null Allowed:** No

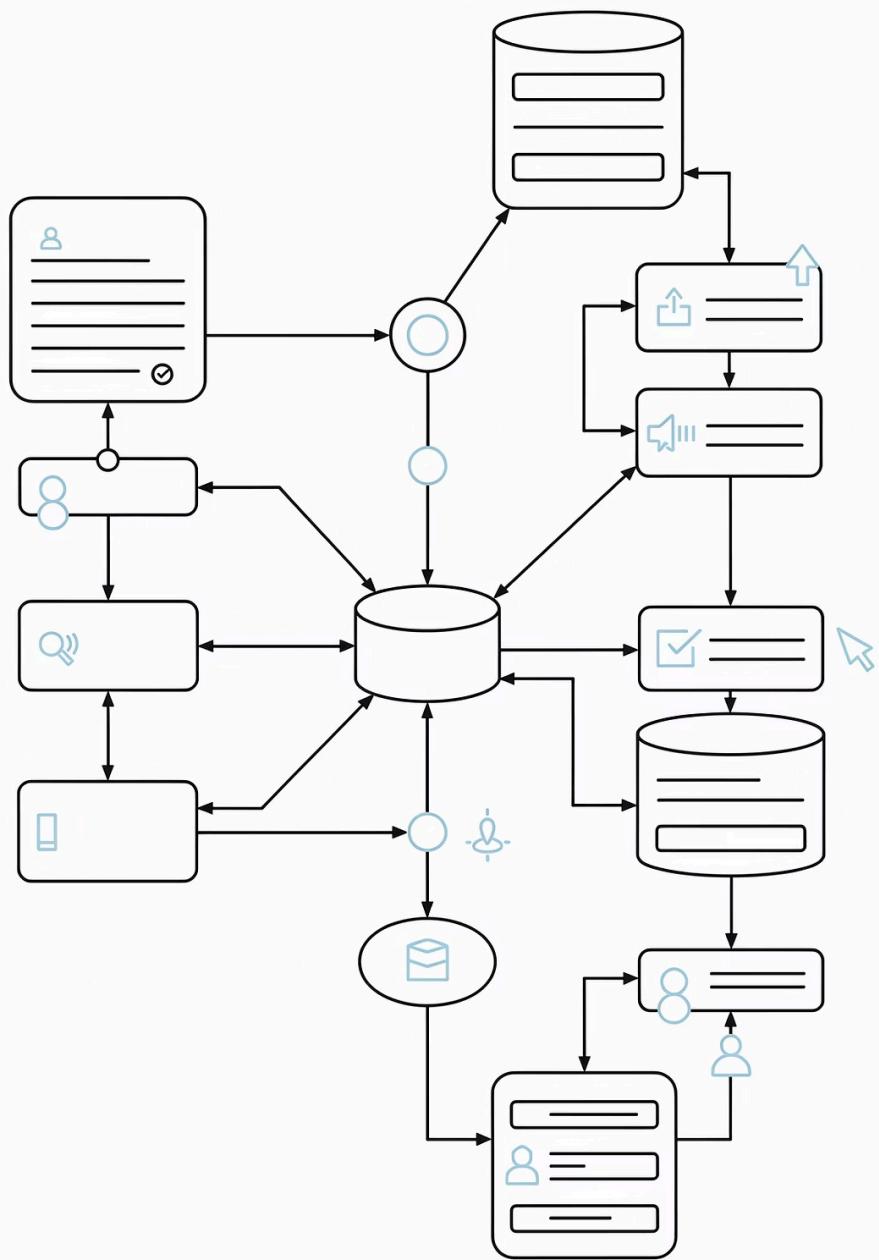
🔗 SECTION 2.1.3

# Relationships: Connecting Entities

A relationship is an association or connection between two or more entities. Relationships capture how entities interact with each other in the real world and are fundamental to understanding the structure of your data.

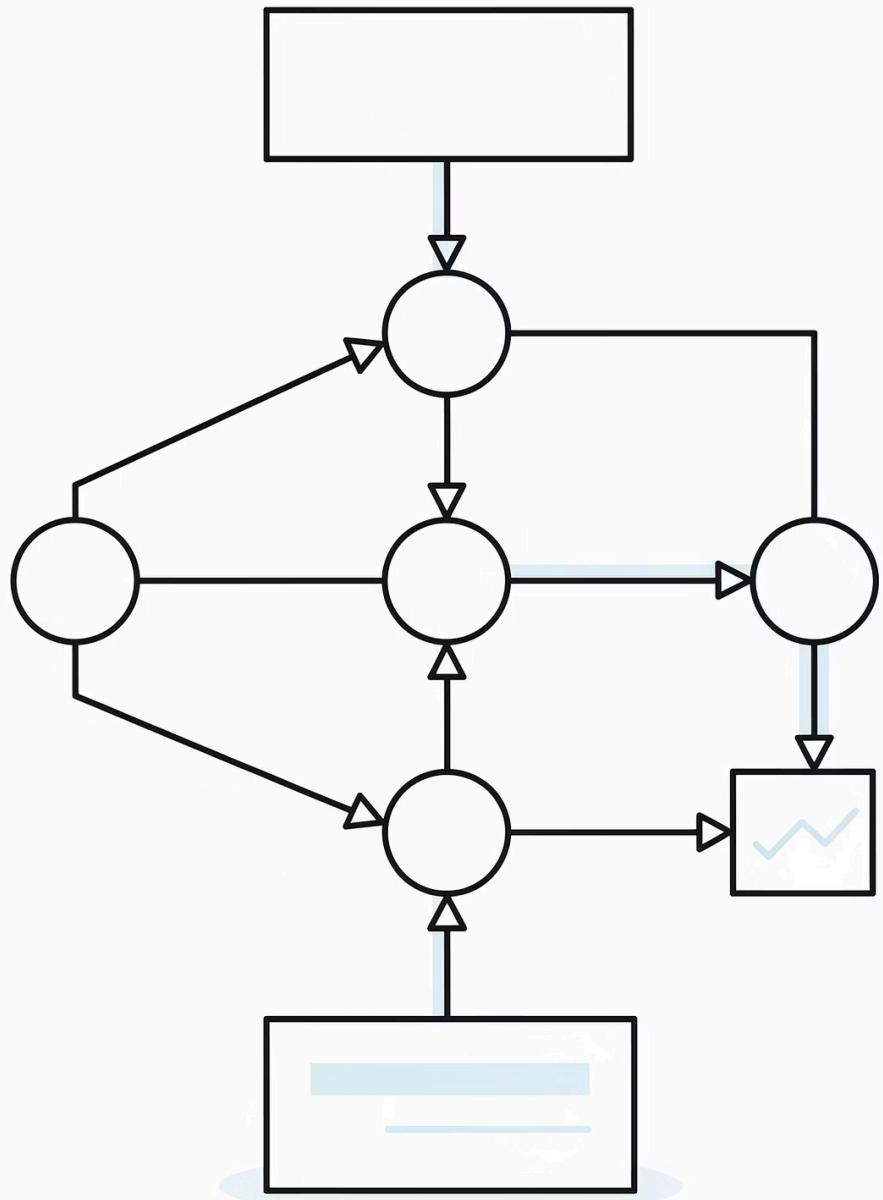
For example, a STUDENT entity may be related to a COURSE entity through an "enrolls in" relationship. Similarly, an EMPLOYEE entity might be related to a DEPARTMENT entity through a "works in" relationship.

Relationships are represented by diamond shapes in ER diagrams and are named using verbs that describe the association. The relationship name should clearly convey the nature of the connection between entities.



# Relationship Examples





SECTION 2.2.3.1

# Participation Constraints

Participation constraints define whether all instances of an entity must participate in a relationship or if participation is optional. This concept is crucial for maintaining data integrity and understanding mandatory versus optional associations.

## Total Participation

Every instance of the entity must participate in the relationship. Represented by a double line in ER diagrams.

**Example:** Every STUDENT must be enrolled in at least one COURSE. No student can exist without course enrollment.

## Partial Participation

Some instances may not participate in the relationship. Represented by a single line in ER diagrams.

**Example:** Not every EMPLOYEE has to manage a PROJECT. Some employees are individual contributors.

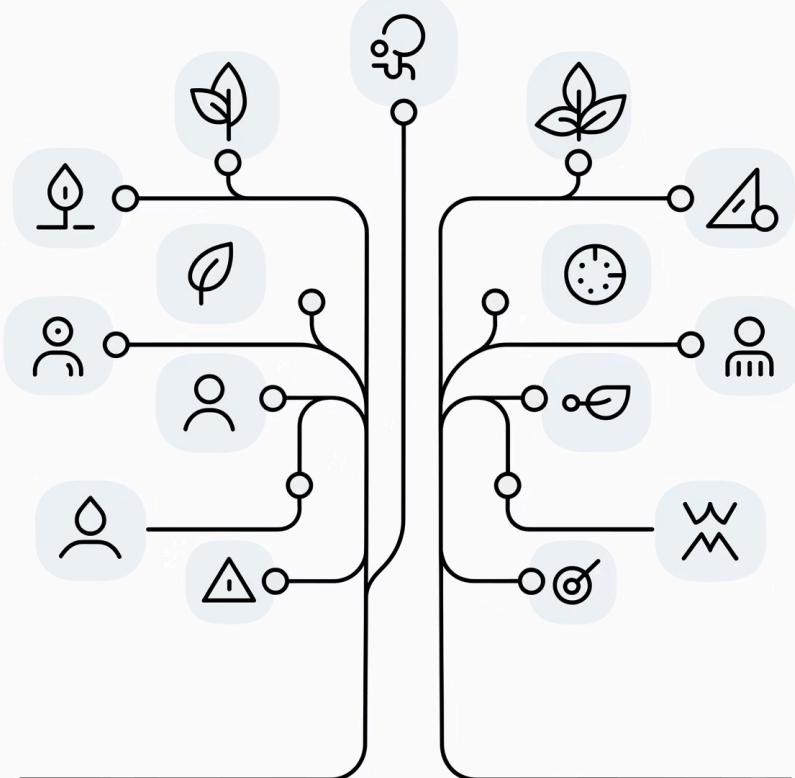
# Recursive Relationships

## Definition

A recursive relationship (also called a self-referencing relationship) occurs when an entity has a relationship with itself. The same entity type participates in the relationship in different roles.

## Common Examples

- **Employee-Supervisor:** An EMPLOYEE supervises other EMPLOYEES
- **Course Prerequisites:** A COURSE may be a prerequisite for other COURSES
- **Social Networks:** A PERSON is friends with other PERSONS
- **Organizational Structure:** A DEPARTMENT may be part of another DEPARTMENT



### Role Names

In recursive relationships, role names are essential to distinguish the different roles the entity plays. For example, in the Employee-Supervisor relationship, roles would be "supervisor" and "subordinate".

# Degree of Relationship Set

The degree of a relationship refers to the number of entity types that participate in the relationship. Understanding relationship degree is essential for accurate database modeling and determines the complexity of associations between entities.



## Binary (Degree 2)

Most common type. Involves two entity types.

**Example:** STUDENT enrolls in COURSE



## Ternary (Degree 3)

Involves three entity types simultaneously.

**Example:** SUPPLIER supplies PART to PROJECT



## N-ary (Degree N)

Involves more than three entity types (rare in practice).

**Example:** Complex multi-entity business processes

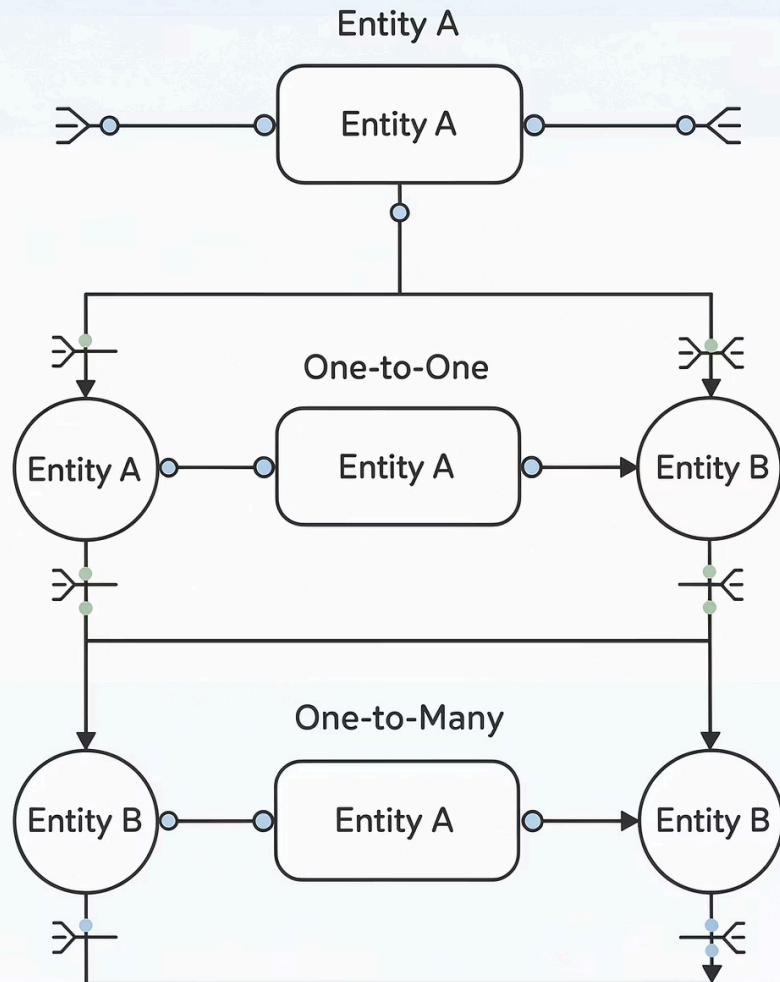
While binary relationships are most common and easiest to understand, ternary and higher-degree relationships are sometimes necessary to accurately model complex real-world scenarios. However, many ternary relationships can be decomposed into multiple binary relationships for simplicity.

# Ternary Relationship Example

Consider a scenario where we need to track which SUPPLIER supplies which PART to which PROJECT. This requires a ternary relationship because all three entities are needed to fully describe the business rule.



Breaking this into binary relationships would lose important information. We need to know specifically which supplier provides which part for which project—all three components together form a single business transaction.



SECTION 2.2

## Mapping Cardinality

Mapping cardinality, also called cardinality ratio, specifies the number of entity instances to which another entity instance can be associated through a relationship. It defines the numerical constraints on relationships and is one of the most important aspects of database design.

Cardinality constraints help enforce business rules and ensure data integrity. They determine whether relationships are one-to-one, one-to-many, or many-to-many, which directly impacts how tables are structured in the physical database.

Understanding and correctly implementing cardinality is crucial for preventing data anomalies and maintaining consistency across the database.

# Types of Cardinality

1

## One-to-One (1:1)

Each instance of Entity A is associated with at most one instance of Entity B, and vice versa.

**Example:** PERSON and PASSPORT - Each person has one passport, and each passport belongs to one person.

2

## One-to-Many (1:N)

Each instance of Entity A can be associated with multiple instances of Entity B, but each instance of B is associated with only one instance of A.

**Example:** DEPARTMENT and EMPLOYEE - One department has many employees, but each employee belongs to one department.

3

## Many-to-One (N:1)

Multiple instances of Entity A can be associated with one instance of Entity B.

**Example:** STUDENT and ADVISOR - Many students can have the same advisor.

4

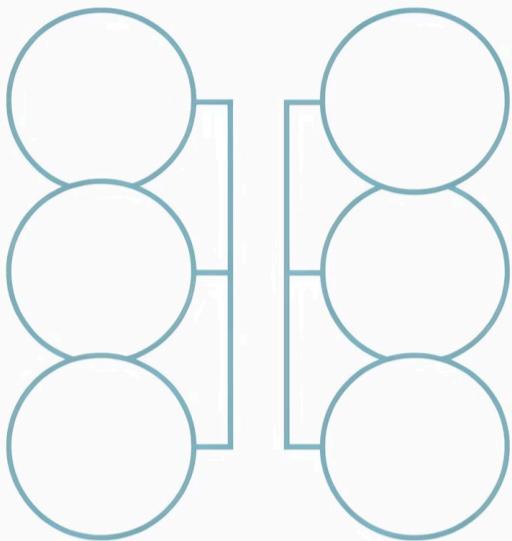
## Many-to-Many (M:N)

Each instance of Entity A can be associated with multiple instances of Entity B, and vice versa.

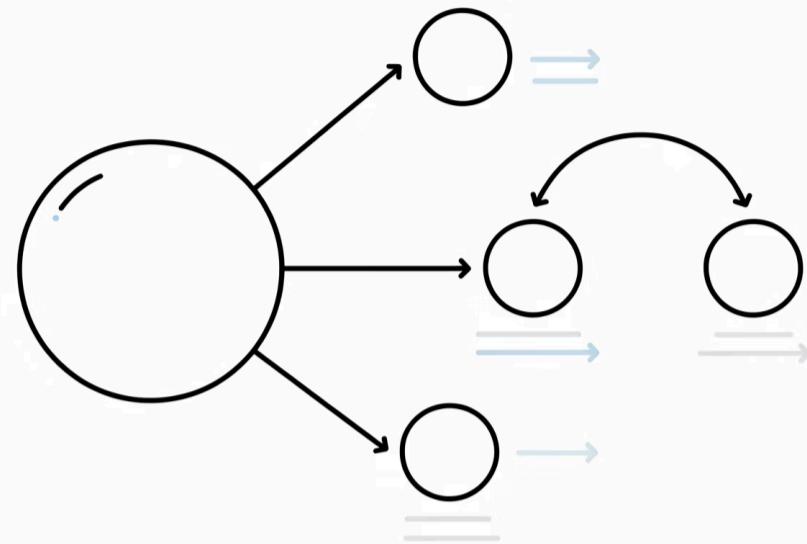
**Example:** STUDENT and COURSE - Students can enroll in many courses, and courses can have many students.

# Cardinality in Practice

## 1:1 Relationship



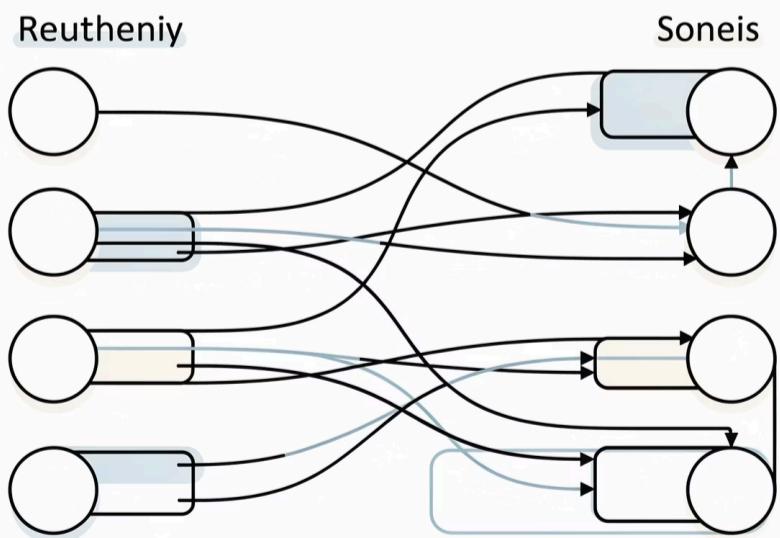
## 1:N Relationship



Implementation typically uses a foreign key in either table with a unique constraint, or combines both entities into a single table.

Implementation places a foreign key in the "many" side table referencing the "one" side's primary key.

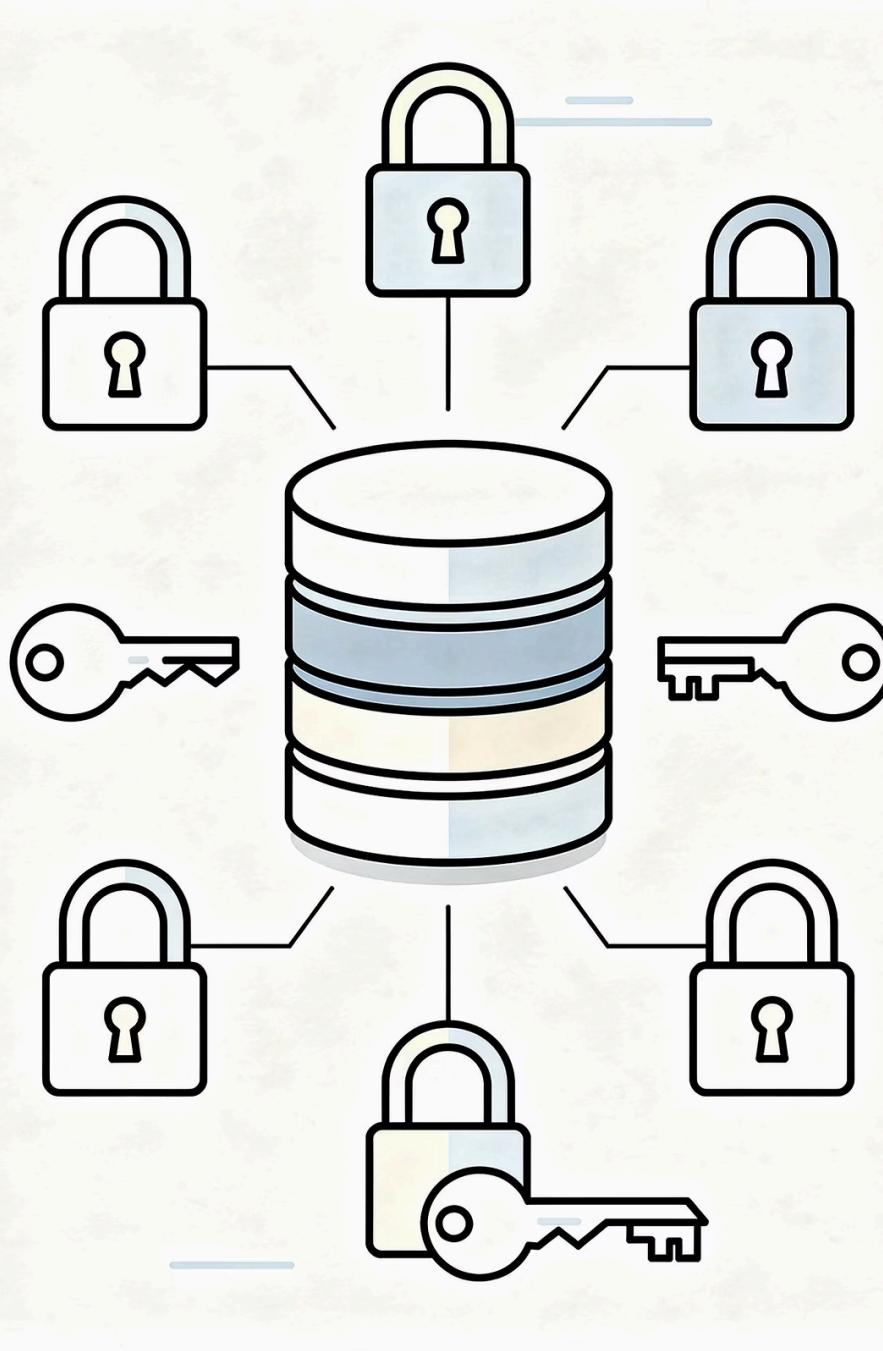
## M:N Relationship



### Junction Table Example

STUDENT\_COURSE table with Student\_ID and Course\_ID creates the many-to-many relationship between STUDENT and COURSE tables.

Implementation requires a junction table (also called bridge table or associative entity) containing foreign keys from both related tables.



KEY SECTION 2.3

## Database Keys

Keys are attributes or sets of attributes that uniquely identify entity instances and establish relationships between tables. They are fundamental to maintaining data integrity, enabling efficient data retrieval, and enforcing constraints in relational databases.

Proper key design ensures that each record can be uniquely identified, prevents duplicate entries, and maintains referential integrity across related tables. Understanding different types of keys is essential for effective database design and normalization.

# Types of Keys



## Primary Key

A unique identifier for each record in a table. Cannot contain NULL values. Each table has exactly one primary key.

**Example:** Student\_ID in STUDENT table



## Foreign Key

An attribute that creates a link between two tables by referencing the primary key of another table.

**Example:** Department\_ID in EMPLOYEE table references DEPARTMENT table



## Candidate Key

Any attribute or minimal set of attributes that can uniquely identify records. A table can have multiple candidate keys.

**Example:** Student\_ID, Email, or SSN could all be candidate keys



## Super Key

Any set of attributes that can uniquely identify records. May contain additional attributes beyond what's necessary.

**Example:** {Student\_ID, Name, Email} is a super key

# Key Relationships and Hierarchy



## Primary Key

Selected from candidate keys as the main identifier



## Candidate Keys

Minimal super keys that uniquely identify records



## Super Keys

Any combination of attributes that ensures uniqueness

The hierarchy shows how different key types relate to each other. All primary keys are candidate keys, all candidate keys are super keys, but not all super keys are candidate keys (they may contain redundant attributes).

# Primary Key Selection Criteria

Choosing the right primary key is a critical decision in database design. The primary key should be selected from available candidate keys based on specific criteria that ensure long-term stability and efficiency.



## Uniqueness

Must uniquely identify each record without any possibility of duplication



## Minimality

Should contain the minimum number of attributes necessary for unique identification



## Stability

Value should not change over time; avoid attributes that might be updated



## Not NULL

Must always have a value; NULL values are not allowed in primary keys



## Simplicity

Prefer simple data types (integer) over complex ones (composite attributes)

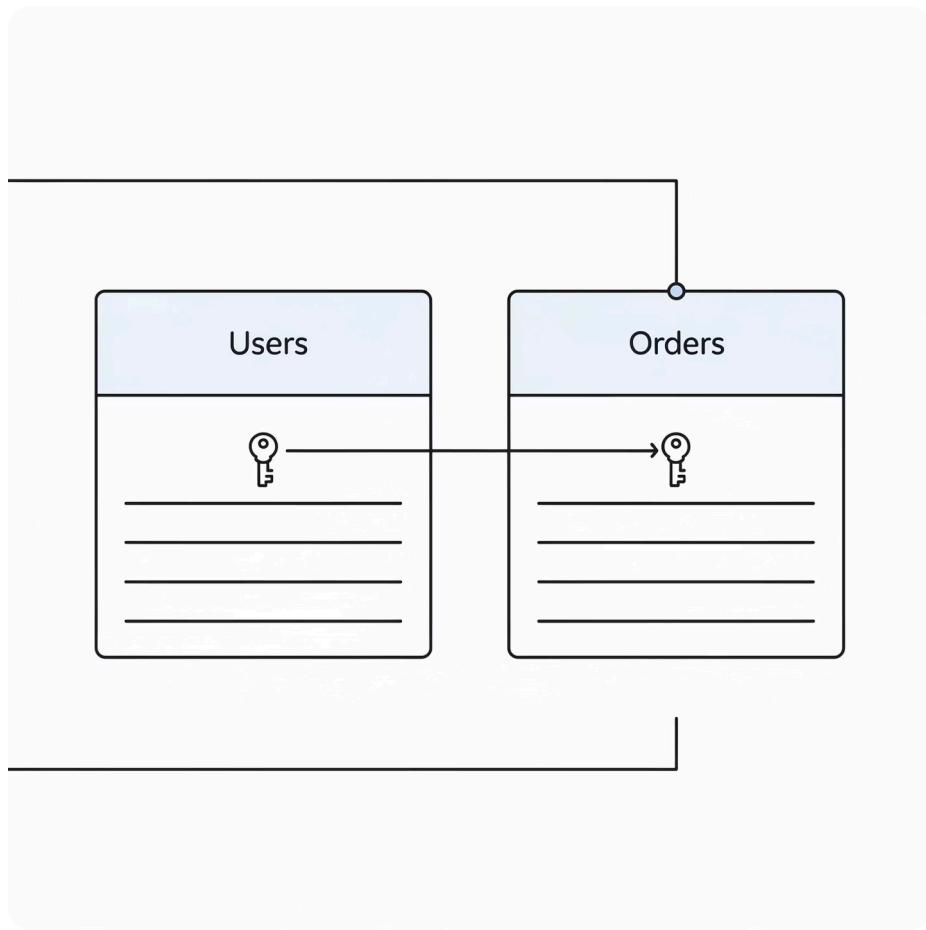
Many modern databases use surrogate keys (auto-incremented integers) as primary keys because they satisfy all these criteria perfectly and are independent of business data.

# Foreign Key Constraints

Foreign keys enforce referential integrity between tables, ensuring that relationships remain valid and consistent. They prevent orphaned records and maintain the logical connections between related data.

## Referential Integrity Rules

- **Insert Rule:** Cannot insert a record with a foreign key value that doesn't exist in the referenced table
- **Update Rule:** Changes to primary keys must be reflected in all related foreign keys (CASCADE) or prevented
- **Delete Rule:** Deletion of a primary key record requires handling of related foreign key records (CASCADE, SET NULL, or RESTRICT)



### CASCADE Operations

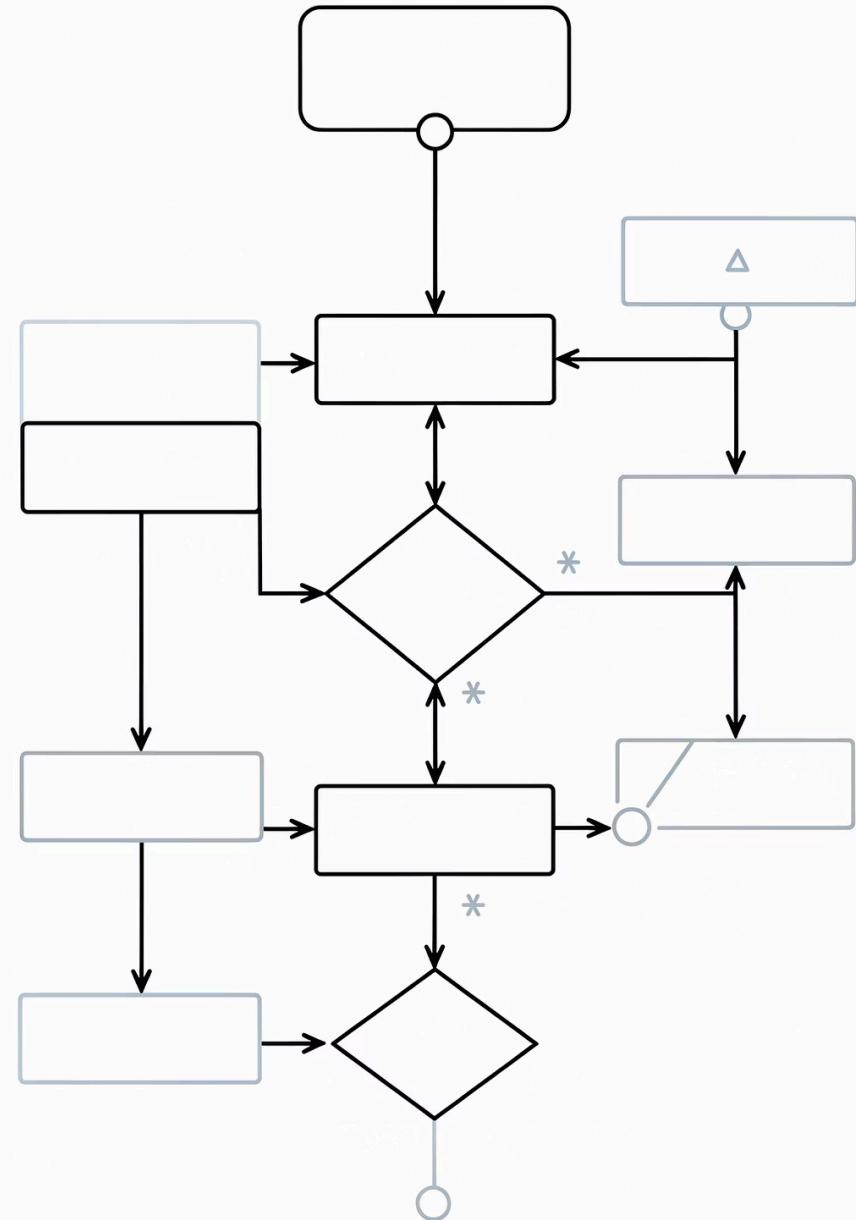
ON DELETE CASCADE automatically deletes related records. ON UPDATE CASCADE automatically updates foreign key values when primary keys change.

SECTION 2.4

# Entity-Relationship Diagrams

Entity-Relationship (ER) Diagrams are graphical representations of entities, attributes, and relationships in a database system. They serve as a visual blueprint that communicates database structure to both technical and non-technical stakeholders.

ER diagrams use standardized symbols and notations to represent different components, making them universally understood by database professionals. They are essential tools in the database design process, helping identify potential issues before implementation.



# ER Diagram Notation Symbols

## Rectangle

Represents an entity type

**Example:** STUDENT, EMPLOYEE, PRODUCT

## Diamond

Represents a relationship between entities

**Example:** Enrolls, Works\_In, Manages

## Double Ellipse

Represents multi-valued attributes

**Example:** Phone\_Numbers, Skills

## Double Rectangle

Represents weak entity types

## Ellipse / Oval

Represents an attribute of an entity

**Example:** Name, Age, Address

## Line

Connects entities to relationships and attributes to entities

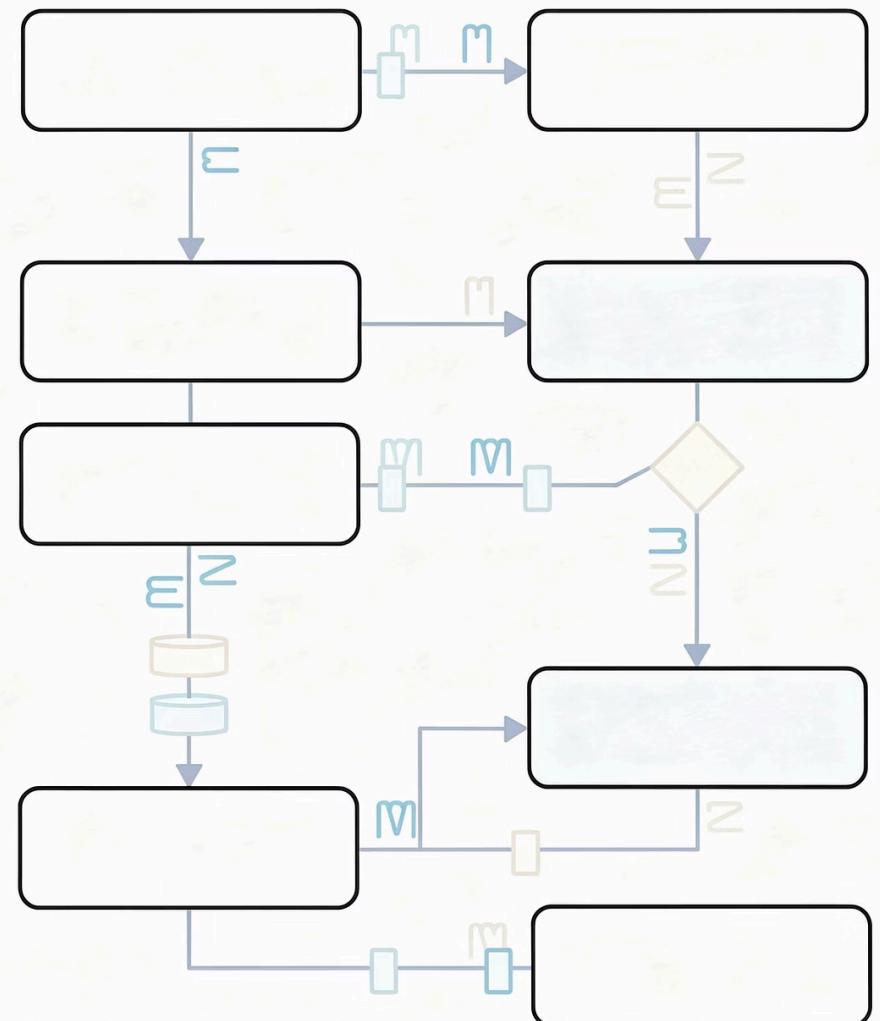
## Dashed Ellipse

Represents derived attributes

**Example:** Age (derived from DOB)

## Double Diamond

Represents identifying relationship for weak entities



# Cardinality Notation in ER Diagrams

Different notations exist for representing cardinality in ER diagrams. The two most common are Chen notation and Crow's Foot notation.

## Chen Notation

Uses numbers to show cardinality:

- 1 indicates "one"
- N or M indicates "many"
- Placed near the entity on the relationship line

**Example:** DEPARTMENT (1) ----<  
employs >---- (N) EMPLOYEE

## Crow's Foot Notation

Uses graphical symbols:

- Single line: one
- Crow's foot (three prongs): many
- Circle: optional (zero)
- Perpendicular line: mandatory (one or more)

Choose a notation style and use it consistently throughout your diagram. Crow's Foot notation is more popular in modern database design tools.

# Creating Effective ER Diagrams

01

## Identify Entities

List all the objects or concepts that need to be stored in the database

02

## Define Attributes

Determine what information needs to be stored about each entity

03

## Identify Relationships

Determine how entities are related to each other

04

## Determine Cardinality

Specify the number constraints for each relationship

05

## Identify Keys

Select primary keys and define foreign keys

06

## Review and Refine

Check for completeness, accuracy, and potential improvements

Creating an ER diagram is an iterative process. Expect to revise your diagram multiple times as you better understand the requirements and identify edge cases.

# Best Practices for ER Diagrams

- **Use clear, descriptive names**

Entity names should be nouns (singular form), relationship names should be verbs that clearly describe the association

- **Keep it simple and readable**

Avoid crossing lines when possible, organize entities logically, and don't overcrowd the diagram

- **Document assumptions**

Include notes about business rules, constraints, or special considerations that aren't obvious from the diagram

- **Validate with stakeholders**

Review the diagram with users and business analysts to ensure it accurately represents their requirements

- **Use consistent notation**

Stick to one notation style throughout the entire diagram to avoid confusion

- **Include all relevant constraints**

Show participation constraints, cardinality, and key attributes clearly

# Common ER Diagram Mistakes

## Redundant Relationships

Creating multiple relationships that represent the same association between entities, leading to data inconsistency

## Missing Relationships

Failing to identify all necessary connections between entities, resulting in incomplete data models

## Incorrect Cardinality

Misidentifying relationship constraints, which can lead to data integrity issues in implementation

## Attributes as Entities

Mistakenly representing simple attributes as separate entities when they should be properties of existing entities

## Missing Primary Keys

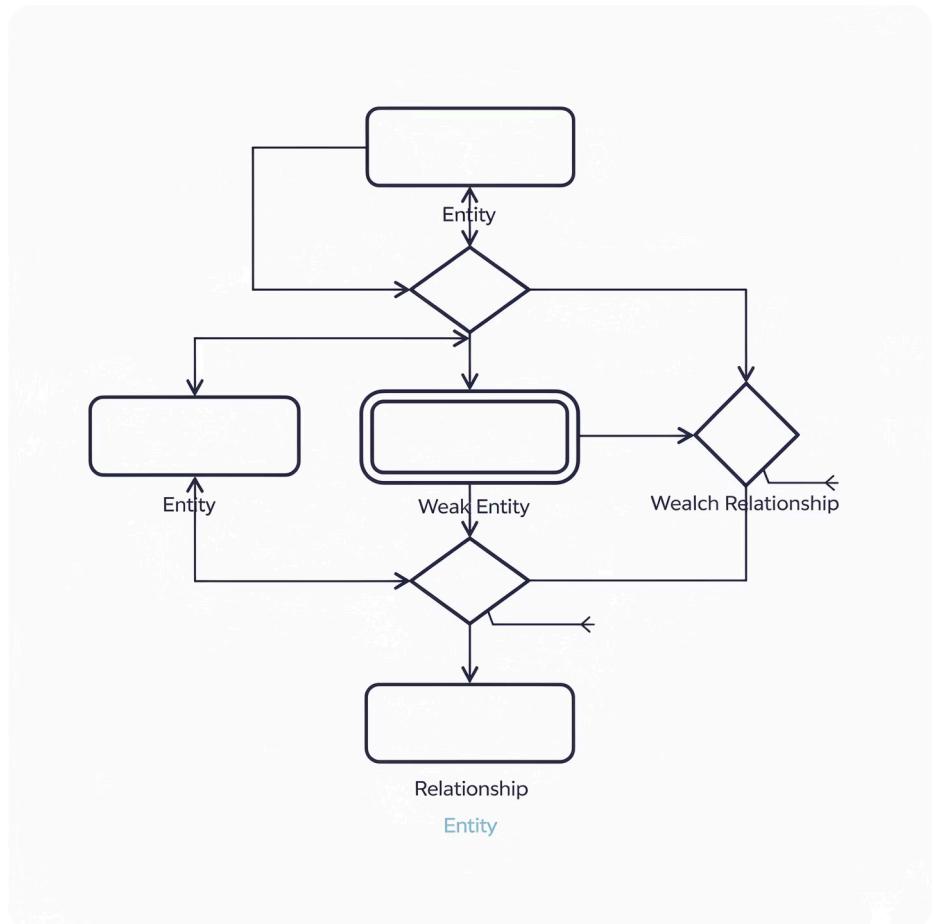
Forgetting to identify unique identifiers for entities, making it impossible to distinguish individual instances

# Weak Entity Sets

A weak entity is an entity that cannot be uniquely identified by its own attributes alone. It depends on another entity (called the owner or strong entity) for its existence and identification.

## Characteristics

- Does not have a primary key of its own
- Has a partial key (discriminator) that uniquely identifies it within the context of its owner
- Existence depends on its relationship with the strong entity
- Represented by a double rectangle in ER diagrams
- Connected through an identifying relationship (double diamond)



## Example

DEPENDENT is a weak entity related to EMPLOYEE. A dependent cannot exist without an employee and is identified by {Employee\_ID + Dependent\_Name}.

# Strong vs Weak Entities

## Strong Entity

Has its own primary key that uniquely identifies each instance independently

Can exist independently of other entities

Represented by single rectangle in ER diagrams

**Examples:** STUDENT, EMPLOYEE, CUSTOMER, PRODUCT

## Weak Entity

Cannot be uniquely identified by its own attributes alone

Depends on a strong entity for existence and identification

Represented by double rectangle in ER diagrams

**Examples:** DEPENDENT, ROOM (in BUILDING), ORDER\_ITEM (in ORDER)

# Weak Entity Examples

## Example 1: Employee and Dependent

An EMPLOYEE (strong entity) has DEPENDENTS (weak entity). Each dependent is identified by the combination of Employee\_ID and Dependent\_Name. If an employee record is deleted, all associated dependent records are also removed.

## Example 2: Building and Room

A BUILDING (strong entity) contains ROOMS (weak entity). Each room is identified by Building\_ID and Room\_Number. Room 101 in Building A is different from Room 101 in Building B. Rooms cannot exist without a building.

## Example 3: Order and Order Item

An ORDER (strong entity) contains ORDER\_ITEMS (weak entity). Each item is identified by Order\_ID and Item\_Number. Order items have no meaning outside the context of their parent order.

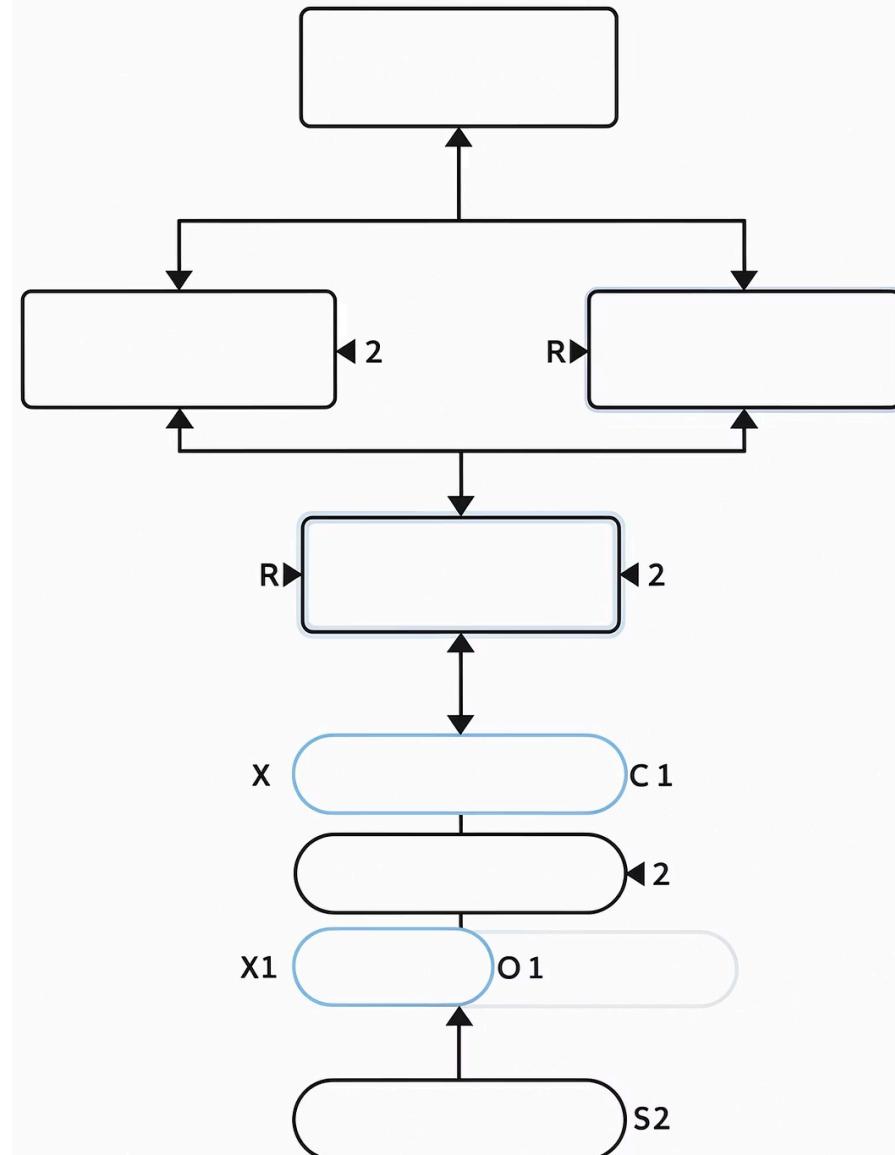
 SECTION 2.6

# Enhanced Entity-Relationship Model

The Enhanced Entity-Relationship (EER) model extends the basic ER model with additional concepts to better represent complex database requirements. It introduces features that capture more semantic information about the database structure.

EER models are particularly useful for large, complex databases where entities share common characteristics or have hierarchical relationships. These advanced concepts help reduce redundancy and improve the clarity of database designs.

The four key concepts in EER modeling are subclass/superclass relationships, generalization, specialization, and aggregation.



# Subclass and Superclass

## Superclass

A superclass (also called parent class or generalized entity) is an entity type that contains common attributes and relationships shared by multiple specialized entities.

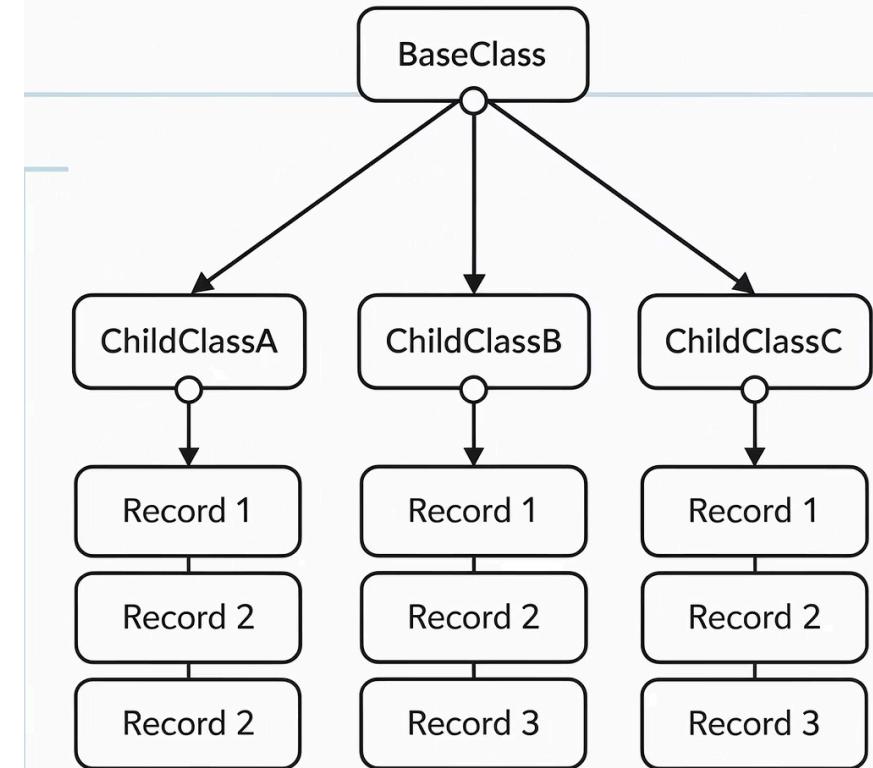
**Example:** PERSON is a superclass containing attributes like Person\_ID, Name, Date\_of\_Birth, and Address that are common to all people.

Subclasses inherit all attributes and relationships from their superclass. This inheritance mechanism promotes reusability and reduces redundancy in the database design.

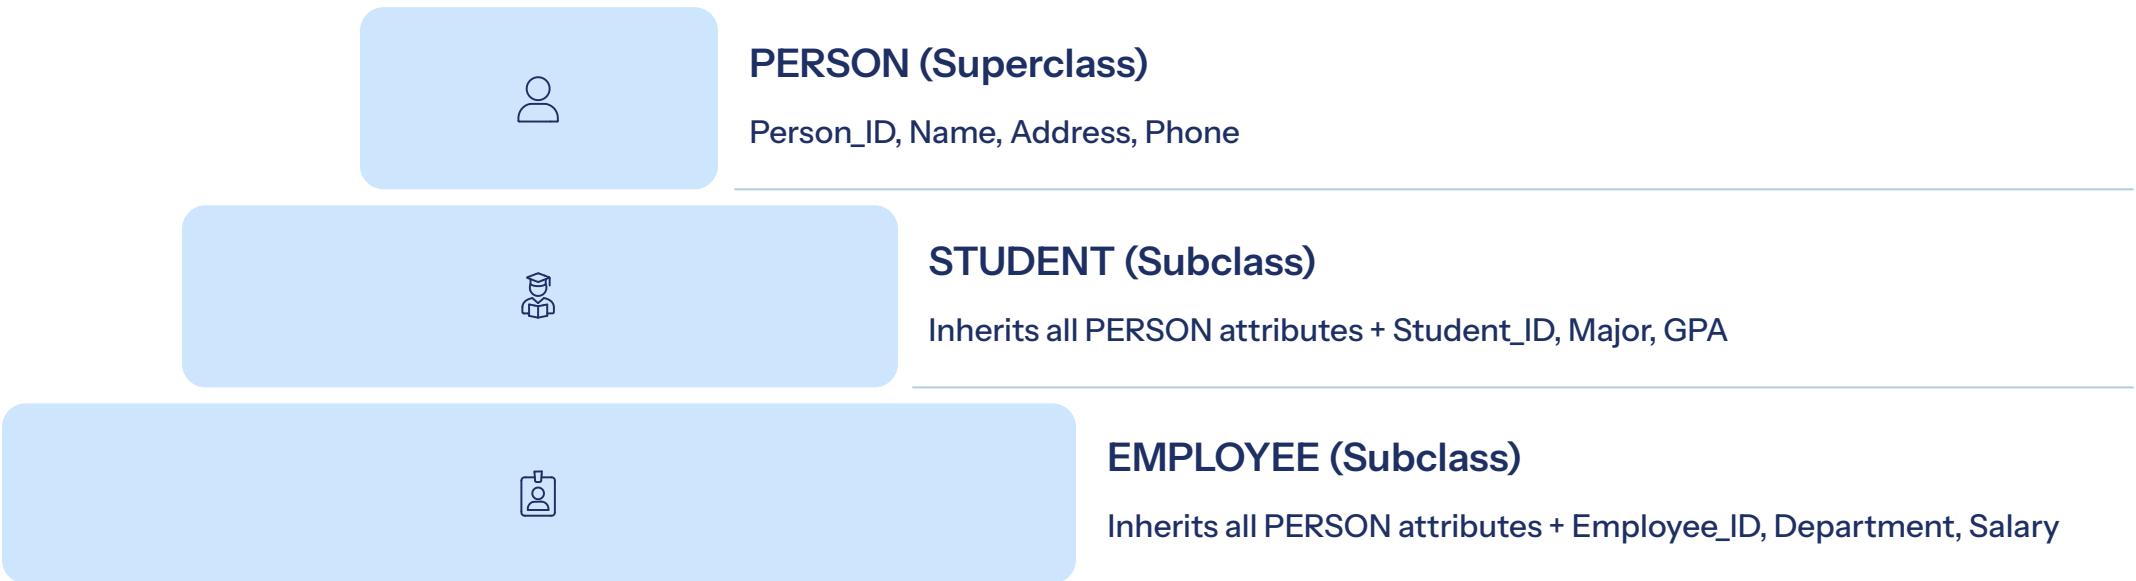
## Subclass

A subclass (also called child class or specialized entity) is an entity type that inherits attributes from a superclass and adds its own specific attributes.

**Examples:** STUDENT and EMPLOYEE are subclasses of PERSON, adding attributes like Student\_ID, GPA or Employee\_ID, Salary respectively.



# Superclass-Subclass Relationships



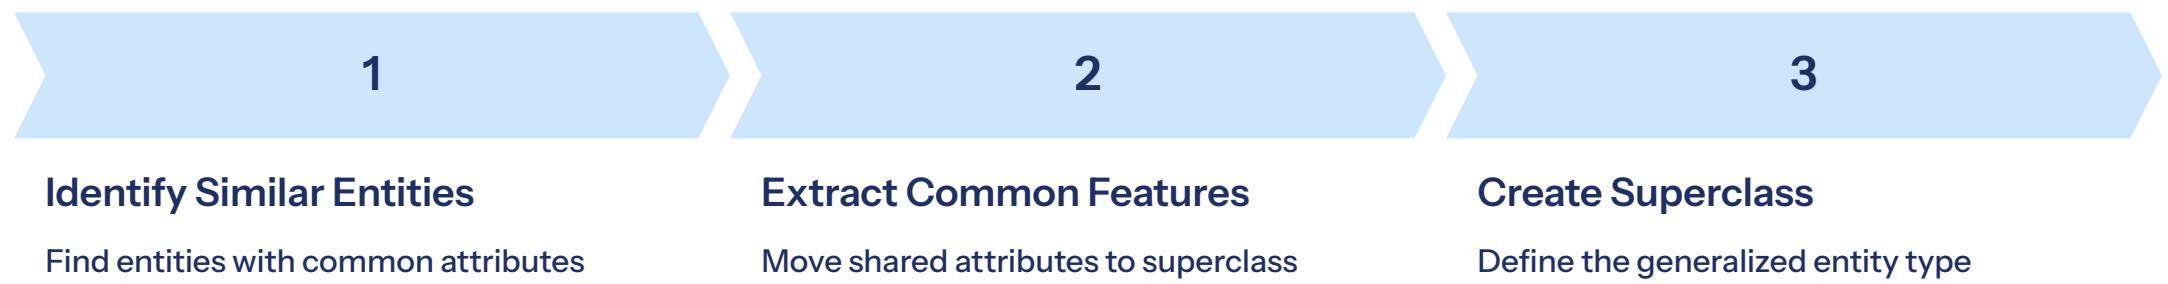
The "IS-A" relationship exists between subclass and superclass. Every STUDENT "is a" PERSON, and every EMPLOYEE "is a" PERSON. This means any instance of a subclass is also an instance of the superclass.

↳ SECTION 2.6.2

# Generalization

Generalization is the bottom-up process of defining a generalized (superclass) entity from multiple specialized (subclass) entities. It identifies common attributes and relationships among existing entities and creates a higher-level abstraction.

## The Generalization Process



**Example:** Given separate CAR, TRUCK, and MOTORCYCLE entities, we can generalize them into a VEHICLE superclass containing common attributes like Vehicle\_ID, Make, Model, and Year, while keeping specific attributes in the subclasses.

# Specialization

Specialization is the top-down process of defining specialized subclass entities from a general superclass entity. It divides a high-level entity into multiple lower-level entities based on distinguishing characteristics.

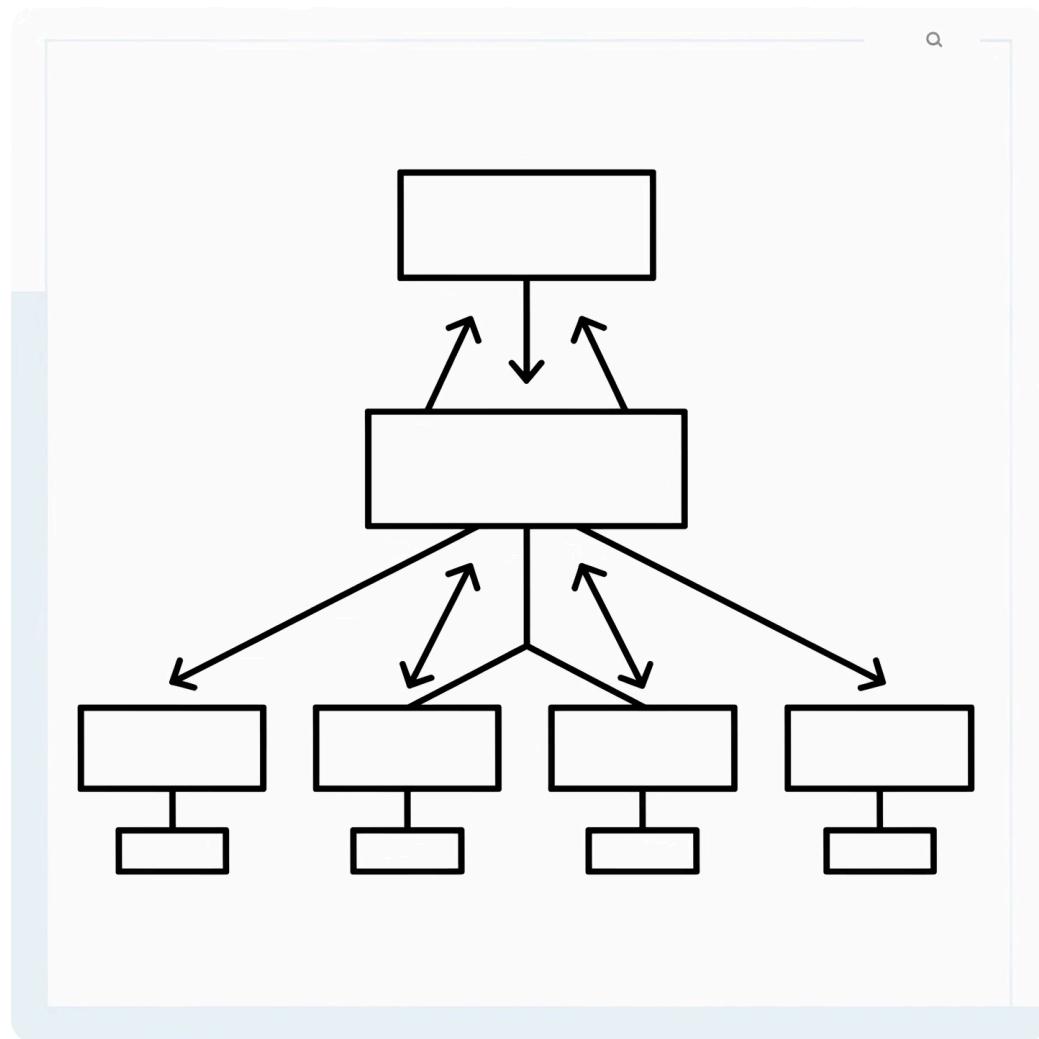
## The Specialization Process



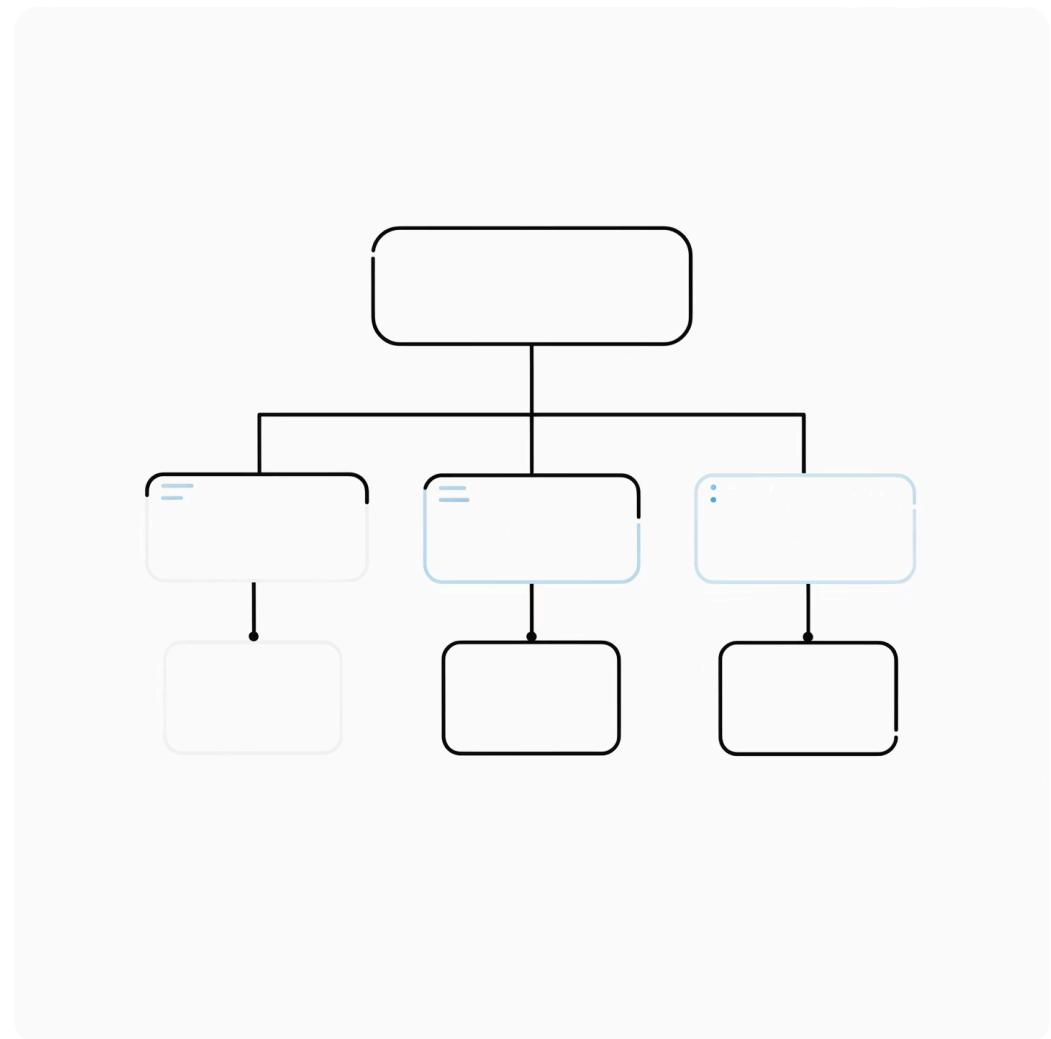
**Example:** Starting with an ACCOUNT entity, we can specialize it into SAVINGS\_ACCOUNT and CHECKING\_ACCOUNT subclasses, each with specific attributes like Interest\_Rate or Overdraft\_Limit respectively.

# Generalization vs Specialization

## Generalization (Bottom-Up)



## Specialization (Top-Down)



- Starts with specific subclasses
- Combines them into a general superclass
- Emphasizes commonality
- Reduces redundancy by eliminating duplicate attributes
- **Think:** "What do these entities have in common?"

- Starts with a general superclass
- Divides it into specific subclasses
- Emphasizes differences
- Adds detail through specialized attributes
- **Think:** "What makes these instances different?"

These are complementary processes that achieve the same result: a hierarchy of entities with inheritance relationships. The choice between them depends on your design approach and how you initially identify entities.

# Specialization Constraints

## Disjoint vs Overlapping

**Disjoint:** A superclass instance can belong to at most one subclass (represented by 'd' in diagrams)

**Example:** An EMPLOYEE is either HOURLY or SALARIED, not both

**Overlapping:** A superclass instance can belong to multiple subclasses (represented by 'o' in diagrams)

**Example:** A PERSON can be both a STUDENT and an EMPLOYEE

## Total vs Partial

**Total Specialization:** Every superclass instance must belong to at least one subclass (double line)

**Example:** Every ACCOUNT must be either SAVINGS or CHECKING

**Partial Specialization:** Some superclass instances may not belong to any subclass (single line)

**Example:** Not all EMPLOYEES are MANAGERS or TECHNICIANS

# Aggregation

Aggregation is an abstraction concept that allows treating a relationship as a higher-level entity. It represents the idea that relationships themselves can participate in other relationships.

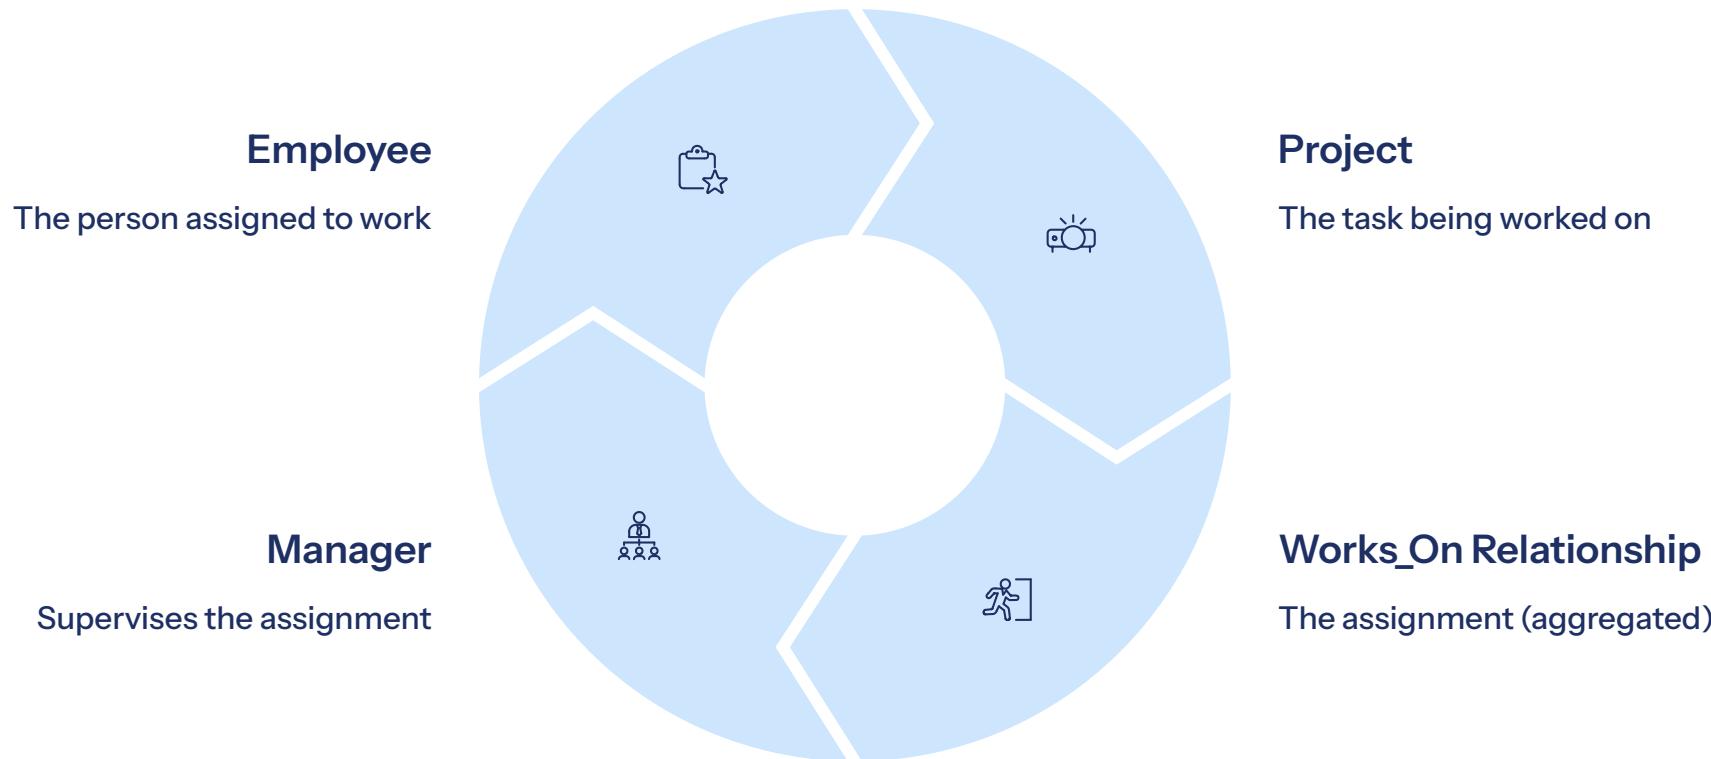
Aggregation is used when you need to create relationships between a relationship and another entity. This is necessary when the relationship has attributes or participates in other relationships.

## When to Use Aggregation

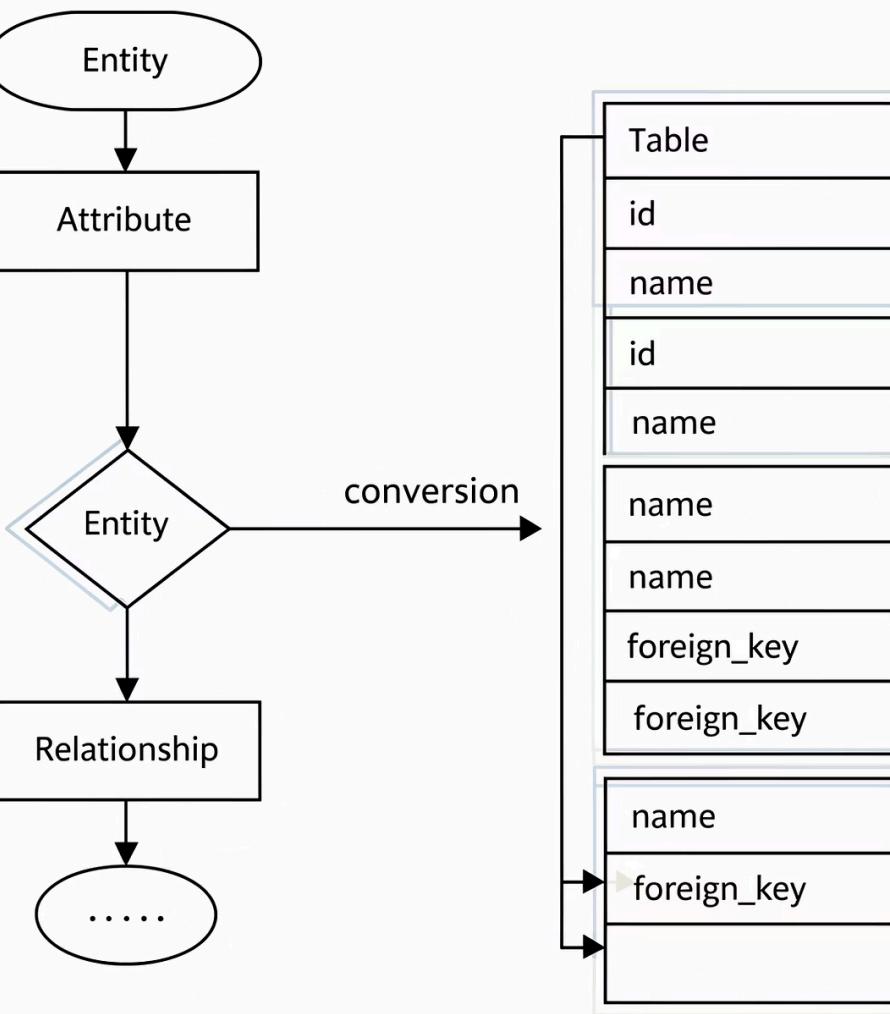
- When a relationship needs to be treated as an entity
- When a relationship participates in another relationship
- When you need to represent relationships between relationships
- To avoid creating redundant entities and relationships

# Aggregation Example

Consider a scenario where EMPLOYEES work on PROJECTs, and we need to track which MANAGER supervises each employee-project assignment.



Without aggregation, we would struggle to properly represent that a manager supervises specific employee-project combinations. Aggregation allows us to treat the "Works\_On" relationship as an entity that can participate in the "Supervises" relationship.



SECTION 2.7

## Converting ER Diagrams to Database Schema

After creating an ER diagram, the next step is converting it into a relational database schema. This transformation process follows systematic rules that map ER components to database tables, columns, and constraints.

The conversion process is crucial because it bridges the gap between conceptual design (ER model) and physical implementation (relational database). Following established conversion rules ensures that the database accurately reflects the designed structure and maintains all specified constraints.

# ER to Relational Mapping Overview

01

## Map Strong Entities

Create a table for each strong entity with columns for all attributes

03

## Map Binary Relationships

Handle 1:1, 1:N, and M:N relationships appropriately

05

## Map Higher-Degree Relationships

Convert ternary and n-ary relationships to tables

02

## Map Weak Entities

Create tables including foreign keys from owner entities

04

## Map Multi-valued Attributes

Create separate tables for attributes with multiple values

06

## Map Specialization Hierarchies

Choose appropriate strategy for superclass/subclass implementation

# Mapping Strong Entity Sets

## Conversion Rule

For each strong entity in the ER diagram:

1. Create a table with the same name
2. Include columns for all simple attributes
3. Expand composite attributes into simple components
4. Designate the primary key
5. Omit derived attributes (calculate when needed)

## Example

### ER Entity: STUDENT

Attributes: Student\_ID (PK), Name, DOB, Address (Street, City, ZIP), Age (derived)

### Resulting Table: STUDENT

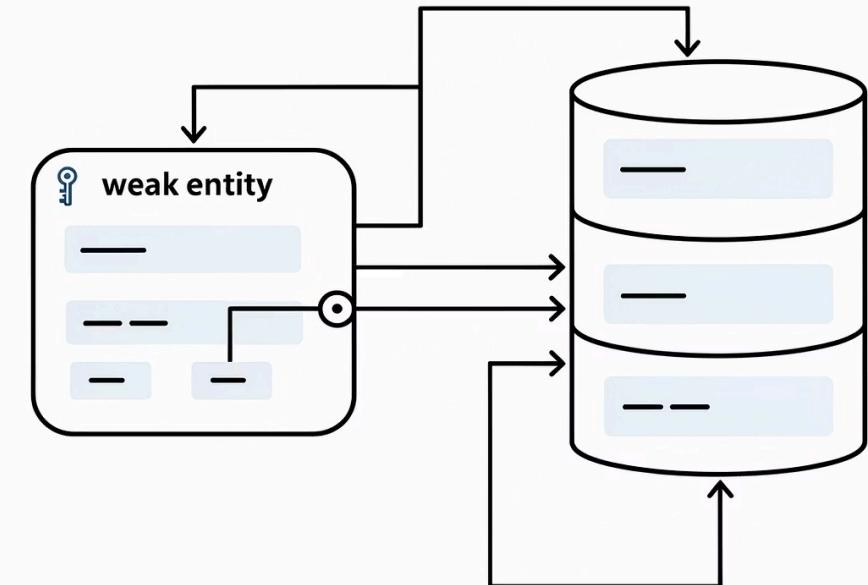
STUDENT  
- Student\_ID (PK)  
- Name  
- DOB  
- Street  
- City  
- ZIP

Note: Age is derived from DOB, so not stored

# Mapping Weak Entity Sets

Weak entities require special handling during conversion because they depend on their owner entity for identification.

## Conversion Rule for Weak Entities



### Create a table for the weak entity

Include all its attributes as columns

### Include the owner's primary key as a foreign key

This establishes the dependency relationship

### Form the primary key

Combine the owner's primary key (FK) + the weak entity's partial key

### Add cascade constraints

Typically use ON DELETE CASCADE to maintain referential integrity

# Weak Entity Mapping Example

## ER Model

### Strong Entity: EMPLOYEE

Employee\_ID (PK), Name, Salary

### Weak Entity: DEPENDENT

Dependent\_Name (Partial Key), Age, Relationship

### Identifying Relationship: Has\_Dependent

## Database Tables

### EMPLOYEE

- Employee\_ID (PK)
- Name
- Salary

### DEPENDENT

- Employee\_ID (FK, PK)
- Dependent\_Name (PK)
- Age
- Relationship

Primary Key = {Employee\_ID, Dependent\_Name}

The dependent cannot exist without an employee. If Employee\_ID 101 is deleted, all dependents with Employee\_ID 101 are automatically removed (CASCADE).

# Mapping Binary Relationships: 1:1

One-to-one relationships can be implemented using several approaches, depending on participation constraints.

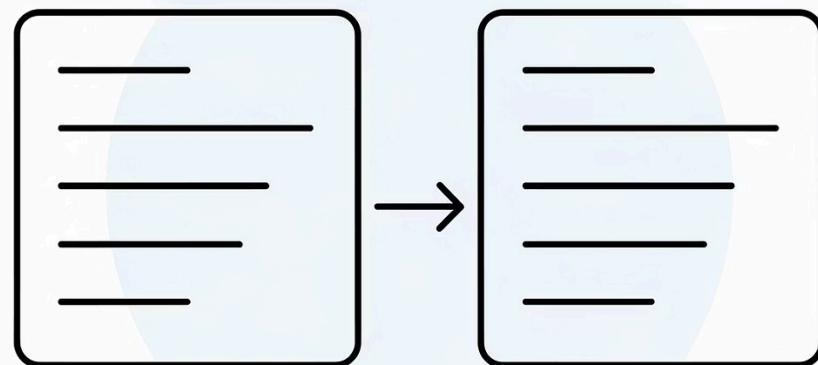
## 1 Foreign Key Approach

Add the primary key of one entity as a foreign key in the other entity's table. Choose the entity with total participation to hold the foreign key to avoid NULL values.

1

**Example:** PERSON (1) - Has - (1) PASSPORT

```
PERSON (Person_ID, Name, Passport_ID FK)  
PASSPORT (Passport_ID, Number, Issue_Date)
```



## 2 Merged Table Approach

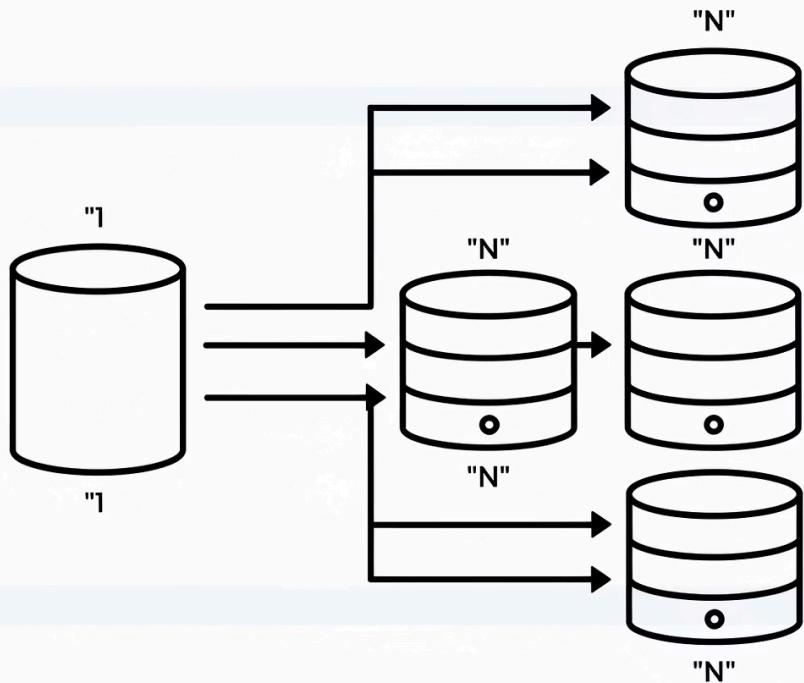
2

If both entities have total participation, merge them into a single table.

**Example:** Combine PERSON and PASSPORT attributes into one CITIZEN table

# Mapping Binary Relationships: 1:N

One-to-many relationships are the most common and straightforward to implement in relational databases.



## Conversion Rule

For each 1:N relationship:

1. Identify the "many" side entity (N)
2. Add the primary key of the "one" side (1) as a foreign key in the "many" side table
3. The foreign key references the primary key of the "one" side
4. No additional table is needed

## Why This Works

The "many" side table can reference the same "one" side record multiple times, but each "many" side record references only one "one" side record.

## Example

**DEPARTMENT (1) - Has - (N)  
EMPLOYEE**

DEPARTMENT  
- Dept\_ID (PK)  
- Dept\_Name  
- Location

EMPLOYEE  
- Emp\_ID (PK)  
- Name  
- Salary  
- Dept\_ID (FK)

# Mapping Binary Relationships: M:N

Many-to-many relationships require a junction table (also called bridge table, associative table, or linking table) because neither entity can hold a simple foreign key reference.

## Conversion Rule



### 1 Create Junction Table

New table with descriptive name

### 2 Add Foreign Keys

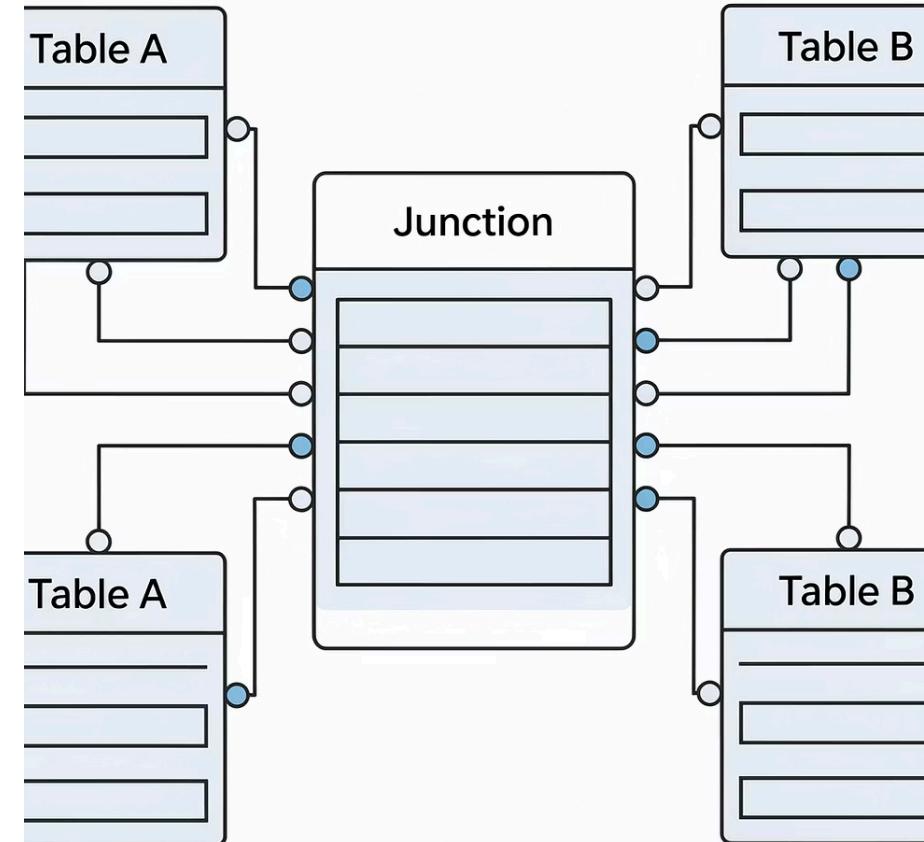
Include PKs from both entities as FPs

### 3 Define Primary Key

Combine both FPs as composite PK

### 4 Add Attributes

Include relationship-specific attributes



# M:N Relationship Mapping Example

## ER Model

### STUDENT (M) - Enrolls\_In - (N) COURSE

Relationship attributes: Enrollment\_Date, Grade

Many students can enroll in many courses, and each course can have many students.

## Database Tables

### STUDENT

- Student\_ID (PK)
- Name
- Major

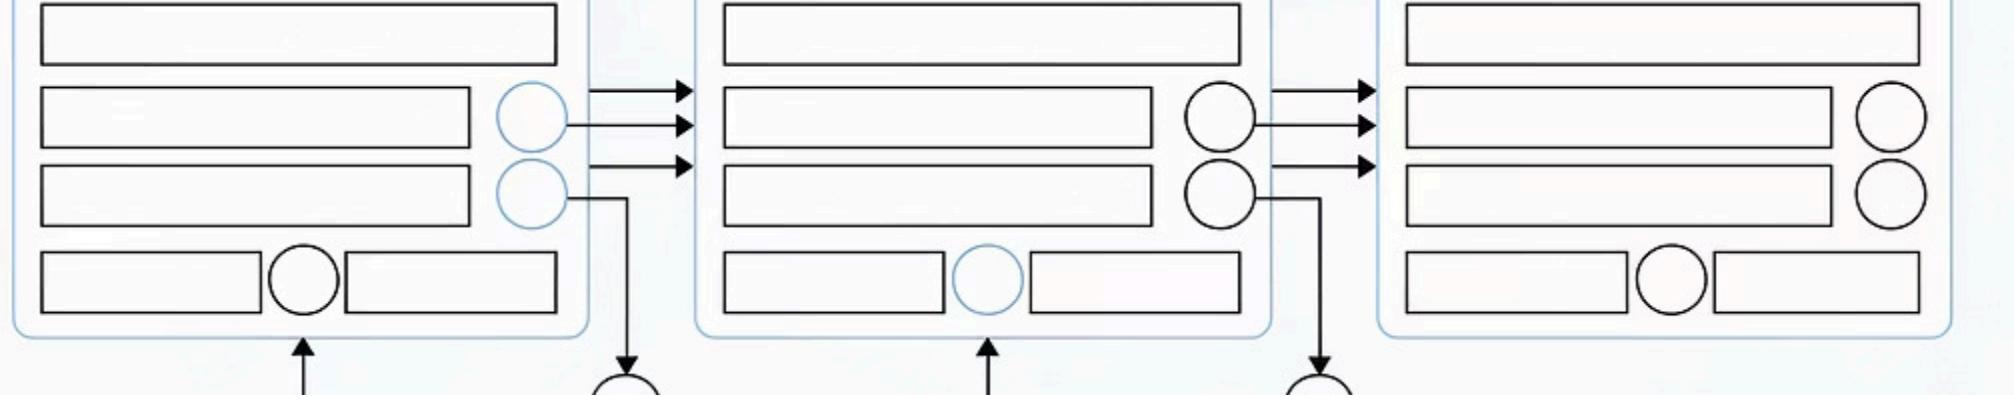
### COURSE

- Course\_ID (PK)
- Course\_Name
- Credits

### ENROLLMENT (Junction Table)

- Student\_ID (FK, PK)
- Course\_ID (FK, PK)
- Enrollment\_Date
- Grade

The ENROLLMENT junction table allows us to represent the many-to-many relationship and store attributes specific to each enrollment. Each row represents one student enrolled in one course.



# Mapping Multi-valued Attributes

Multi-valued attributes (attributes that can have multiple values for a single entity instance) cannot be directly represented in relational databases and must be converted to separate tables.

## Conversion Rule

01

### Create New Table

Create a separate table for the multi-valued attribute

02

### Include Foreign Key

Add the primary key of the parent entity as a foreign key

03

### Add Attribute Column

Include a column for the multi-valued attribute

04

### Define Primary Key

Primary key = {Parent PK, Attribute Value} or add surrogate key

# Multi-valued Attribute Example

## ER Model

### Entity: EMPLOYEE

Attributes:

- Emp\_ID (PK)
- Name
- Department
- Phone\_Numbers (multi-valued)
- Skills (multi-valued)

An employee can have multiple phone numbers and multiple skills.

## Database Tables

### EMPLOYEE

- Emp\_ID (PK)
- Name
- Department

### EMPLOYEE\_PHONE

- Emp\_ID (FK, PK)
- Phone\_Number (PK)

### EMPLOYEE\_SKILL

- Emp\_ID (FK, PK)
- Skill\_Name (PK)
- Proficiency\_Level

This approach allows flexible storage of multiple values while maintaining normalization. An employee with three phone numbers will have three rows in EMPLOYEE\_PHONE.

# Mapping Specialization Hierarchies

Converting superclass/subclass hierarchies to relational tables can be done using multiple strategies. The choice depends on constraints (total/partial, disjoint/overlapping) and query patterns.

## Strategy 1: Single Table

Create one table with all attributes from superclass and all subclasses. Add a discriminator column to identify the subclass type.

**Advantages:** Simple, no joins needed

**Disadvantages:** Many NULL values if subclasses have many unique attributes

## Strategy 2: Table per Subclass

Create separate tables for each subclass containing both inherited and specific attributes.

**Advantages:** No NULL values, clear separation

**Disadvantages:** Data redundancy for shared attributes

## Strategy 3: Table per Hierarchy Level

Create one table for superclass and separate tables for each subclass.

Subclass tables only contain specific attributes + foreign key to superclass.

**Advantages:** No redundancy, no excessive NULLs

**Disadvantages:** Requires joins for queries

# Specialization Mapping Example

**Scenario:** PERSON superclass with STUDENT and EMPLOYEE subclasses (overlapping, partial)

## Strategy 1: Single Table

- PERSON
  - Person\_ID (PK)
  - Name
  - Address
  - Type
  - Student\_ID
  - GPA
  - Major
  - Employee\_ID
  - Dept
  - Salary

*Many NULLs*

## Strategy 2: Per Subclass

- STUDENT
  - Person\_ID (PK)
  - Name
  - Address
  - Student\_ID
  - GPA
  - Major
- EMPLOYEE
  - Person\_ID (PK)
  - Name
  - Address
  - Employee\_ID
  - Dept
  - Salary

*Redundant attributes*

## Strategy 3: Per Level

- PERSON
  - Person\_ID (PK)
  - Name
  - Address
- STUDENT
  - Person\_ID (FK,PK)
  - Student\_ID
  - GPA
  - Major
- EMPLOYEE
  - Person\_ID (FK,PK)
  - Employee\_ID
  - Dept
  - Salary

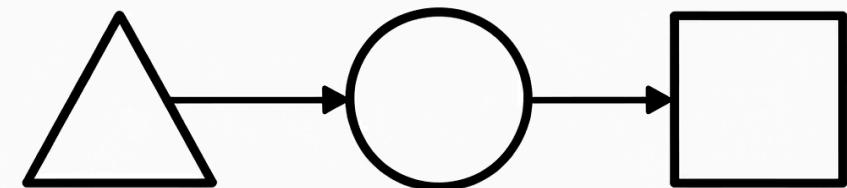
*Best normalized*

# Mapping Ternary Relationships

Ternary relationships (involving three entities) require junction tables similar to many-to-many binary relationships, but with three foreign keys.

## Conversion Rule

- Create Junction Table**  
Create a new table for the ternary relationship
- Add Three Foreign Keys**  
Include primary keys from all three participating entities as foreign keys
- Define Primary Key**  
Typically use all three foreign keys as a composite primary key
- Include Relationship Attributes**  
Add any descriptive attributes of the relationship itself



# Ternary Relationship Example

## ER Model

### Relationship: SUPPLIES

SUPPLIER supplies PART to PROJECT

Attributes: Quantity, Date

This captures which supplier provides which parts for which projects, along with quantity and date information.

## Database Tables

### SUPPLIER

- Supplier\_ID (PK)
- Name

### PART

- Part\_ID (PK)
- Part\_Name

### PROJECT

- Project\_ID (PK)
- Project\_Name

### SUPPLIES

- Supplier\_ID (FK, PK)
- Part\_ID (FK, PK)
- Project\_ID (FK, PK)
- Quantity
- Date

Each row in SUPPLIES represents a specific transaction where one supplier provides one type of part to one project.

# Best Practices for ER to Database Conversion

- **Maintain naming conventions**

Use consistent, descriptive names for tables and columns that match your ER diagram entities and attributes

- **Document foreign key relationships**

Clearly specify all foreign key constraints, including ON DELETE and ON UPDATE rules

- **Choose appropriate data types**

Select data types that efficiently store your data while ensuring accuracy and preventing invalid values

- **Consider indexing**

Create indexes on foreign keys and frequently queried columns to improve performance

- **Normalize appropriately**

Ensure tables are normalized to reduce redundancy, but consider denormalization for performance-critical queries

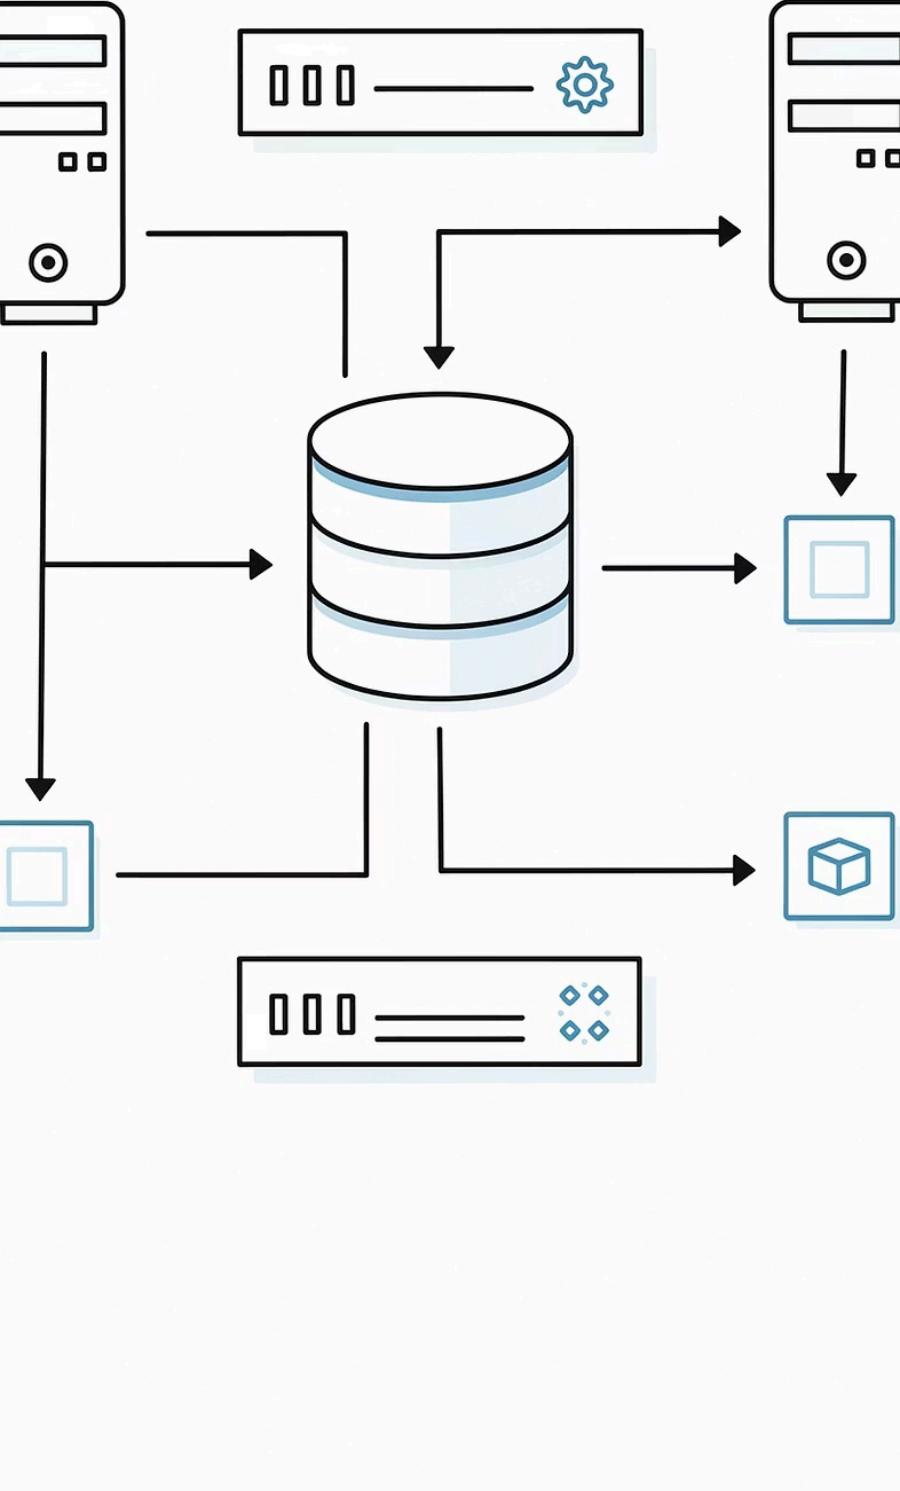
- **Validate with sample data**

Test your schema with realistic data to identify potential issues before full implementation

# Unit II Key Takeaways

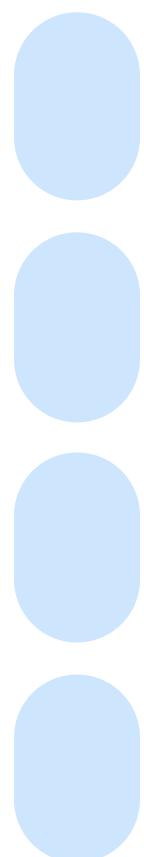
<b>ER Model Fundamentals</b> <p>Entities, attributes, and relationships form the building blocks of database design</p>	<b>Cardinality Constraints</b> <p>1:1, 1:N, and M:N relationships define how entities associate with each other</p>
<b>Keys Are Critical</b> <p>Primary and foreign keys maintain uniqueness and referential integrity</p>	<b>Visual Communication</b> <p>ER diagrams provide clear, standardized representations of database structure</p>
<b>Enhanced Features</b> <p>Specialization, generalization, and aggregation handle complex scenarios</p>	<b>Systematic Conversion</b> <p>Following mapping rules ensures accurate translation from ER to relational schema</p>

# Mastering Database Design



The Entity-Relationship model provides a powerful framework for conceptualizing and designing databases that accurately reflect business requirements. By mastering entities, attributes, relationships, and the conversion process, you build a strong foundation for creating efficient, scalable database systems.

As you continue your journey in database management, remember that good design is iterative. Start with a clear understanding of requirements, create thorough ER diagrams, validate with stakeholders, and systematically convert to relational schemas. These skills will serve you throughout your career in information technology.



**Practice with real scenarios**

**Review and refine designs**

**Study existing databases**

**Collaborate with peers**

Continue exploring database management concepts, and apply these principles to build robust, well-designed database systems.