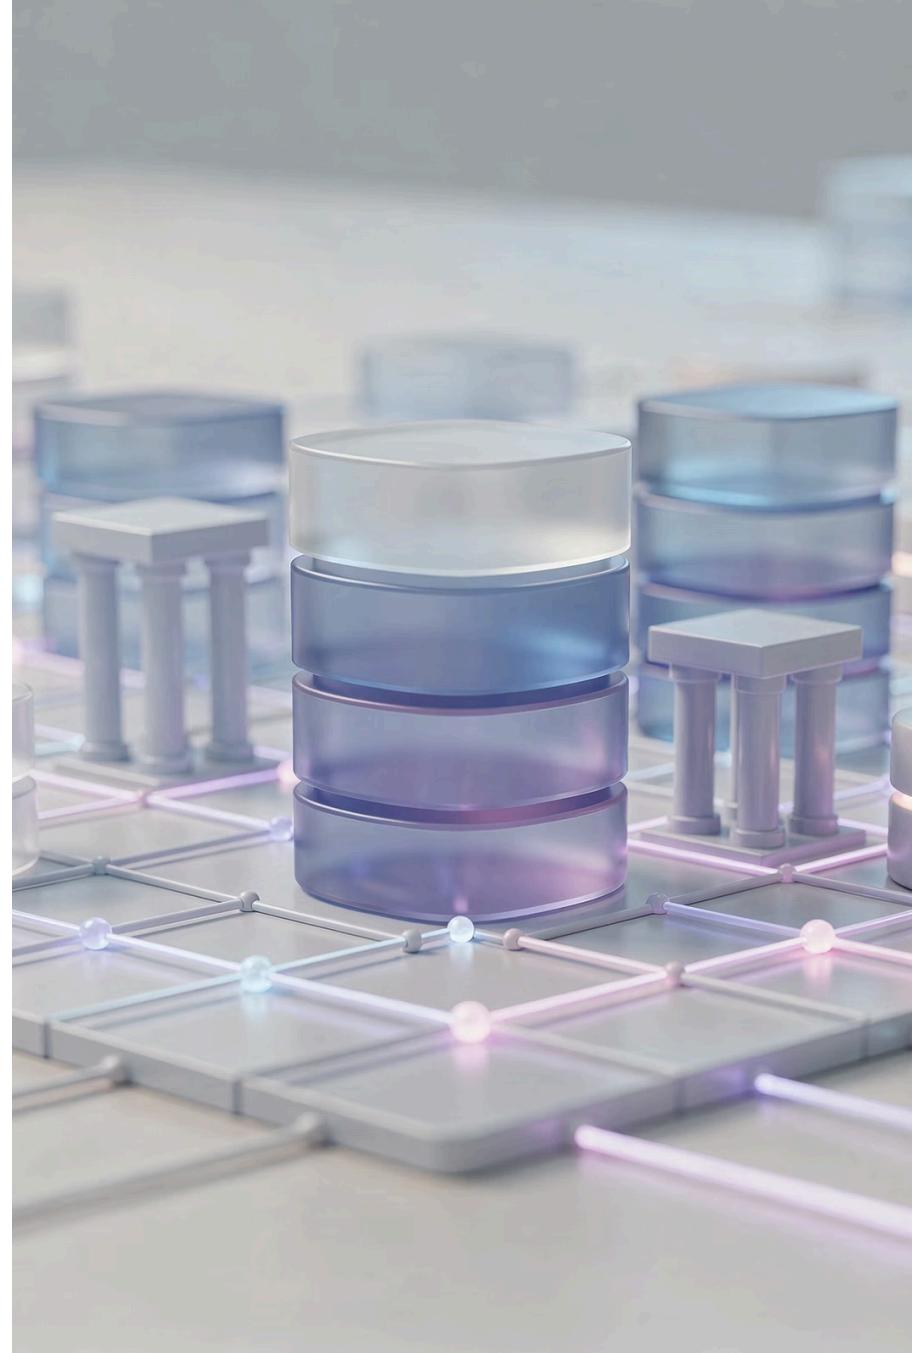


Database Normalization

Refining Database Design Through Systematic Normalization

UNIT IV

DBMS THEORY





COURSE CONTEXT

Course Overview

Program: Diploma Engineering

Branch: Information & Communication Technology

Subject: Database Management System

Code: DI04032011

This unit focuses on the critical process of database normalization, a systematic approach to organizing data that eliminates redundancy and ensures data integrity. Understanding normalization is fundamental for ICT professionals who design and maintain efficient database systems.

Through this comprehensive study, you'll learn how to transform poorly designed databases into robust, scalable solutions that support enterprise operations effectively.

 LEARNING PATH

What You'll Master in Unit IV

01

Foundation Concepts

Understanding why normalization matters in database design and its impact on system performance

02

Functional Dependencies

Exploring the relationships between attributes and how they determine data organization

03

Normal Forms

Applying systematic rules to achieve First, Second, and Third Normal Forms

04

Practical Application

Transforming real-world scenarios into properly normalized database structures



The Database Design Challenge

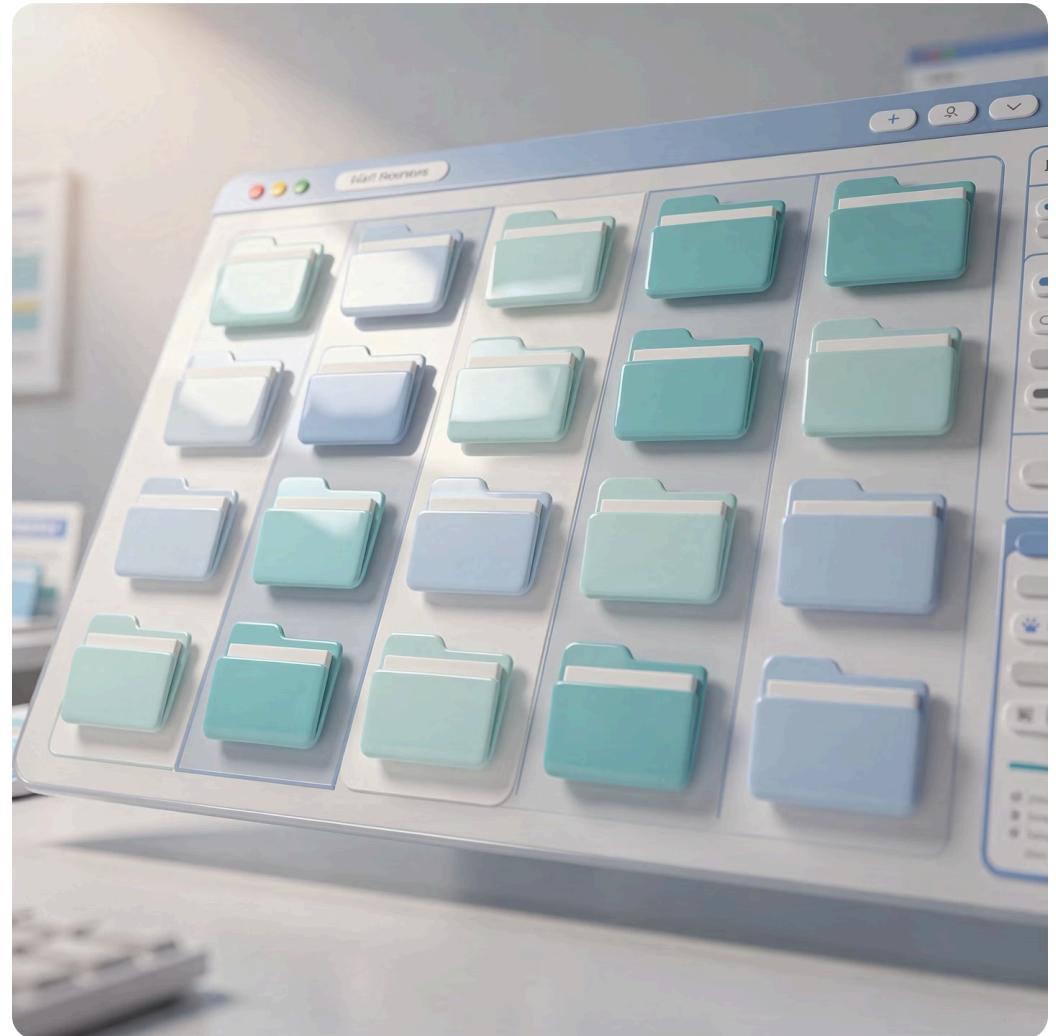
Without proper organization, databases become inefficient, error-prone, and difficult to maintain. Normalization provides the systematic framework to prevent these problems before they occur.

Consider a student registration system where course information is duplicated across thousands of enrollment records. When a course name changes, hundreds of records require updates. This redundancy wastes storage and creates opportunities for inconsistency.

4.1 Importance of Normalization

Normalization is a systematic technique for organizing database tables to minimize data redundancy and dependency problems. It transforms a database schema into a structure that reduces anomalies during data operations.

The process follows specific rules called normal forms, each addressing different types of redundancy and dependency issues. By applying these forms sequentially, designers create databases that are efficient, maintainable, and scalable.



CORE BENEFITS

Why Normalization Matters

Eliminates Redundancy

Removes duplicate data storage, reducing database size and improving efficiency. Each piece of information is stored in exactly one location.

Ensures Data Integrity

Maintains consistency across the database by preventing conflicting or contradictory information from existing simultaneously.

Simplifies Maintenance

Makes updates, insertions, and deletions straightforward by organizing data logically and eliminating complex dependencies.

Improves Query Performance

Optimizes data retrieval operations through better-organized table structures and reduced data scanning requirements.

Data Anomalies: The Problems Normalization Solves

Poorly designed databases suffer from three critical types of anomalies that normalization systematically eliminates. Understanding these problems helps appreciate why normalization is essential.

1

Insertion Anomaly

Inability to add data without the presence of other data. For example, you cannot add a new course to the system until at least one student enrolls in it.

2

Deletion Anomaly

Unintended loss of data when deleting other data. Removing the last student from a course might accidentally delete all course information.

3

Update Anomaly

Inconsistency arising from partial updates. Changing an instructor's name in one record but not others creates conflicting information.

Real-World Impact: Before and After

Unnormalized System



Normalized System



- Wasted storage space from duplicate data
- Inconsistent information across records
- Complex and error-prone update procedures
- Difficulty maintaining referential integrity
- Poor query performance on large datasets

- Efficient storage with minimal redundancy
- Single source of truth for each data element
- Simple, reliable data modification operations
- Automatic consistency through relationships
- Optimized performance for common queries

The Cost-Benefit Perspective

50%

Storage Reduction

Typical savings in database size
after proper normalization

80%

Fewer Errors

Reduction in data inconsistency
issues reported

3x

Faster Updates

Improvement in modification
operation speed

60%

Maintenance Time

Decrease in time spent fixing data
issues



When to Apply Normalization

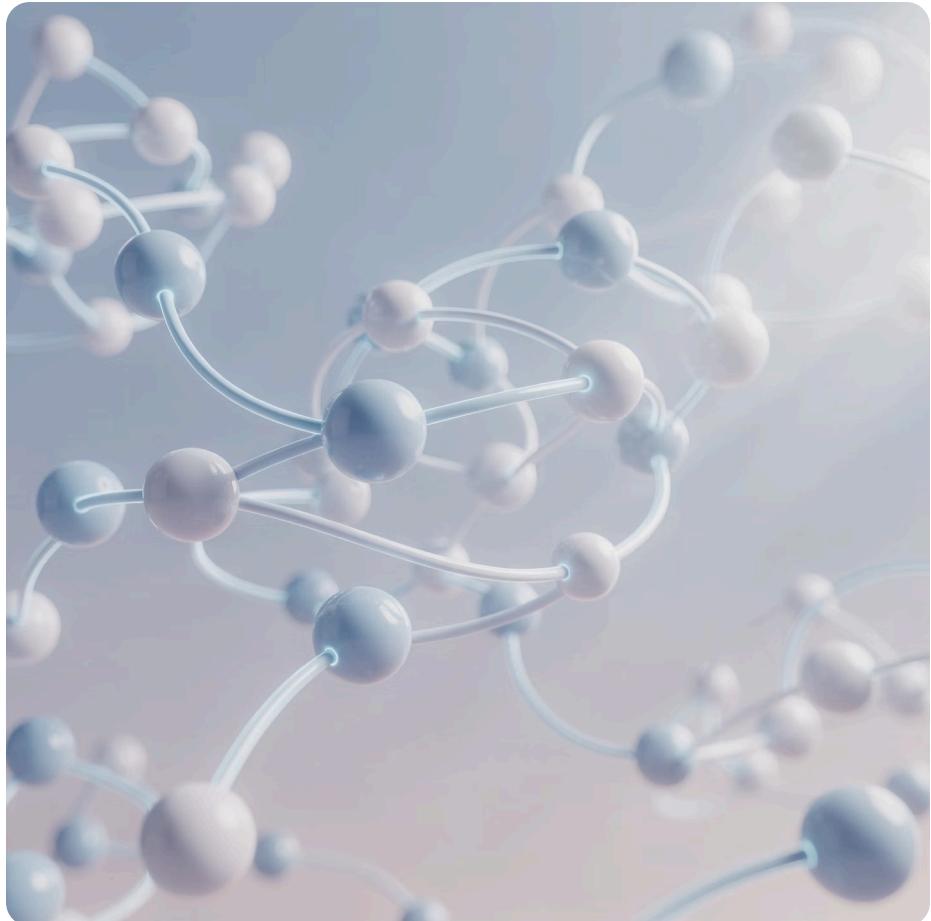
Normalization is most effective when applied during the database design phase, before implementation. However, existing databases can also be normalized through careful refactoring, though this requires more effort and planning.

- 1 Requirements Analysis**
Identify entities and relationships in the problem domain
- 2 Initial Design**
Create preliminary table structures and attributes
- 3 Apply Normalization**
Systematically refine tables through normal forms
- 4 Implementation**
Deploy the optimized database structure

Functional Dependencies

The foundation of normalization theory

Understanding Functional Dependencies



A functional dependency describes a relationship between attributes in a table where one attribute (or set of attributes) uniquely determines another attribute. This concept is fundamental to understanding and applying normalization.

Formally, we say attribute B is functionally dependent on attribute A if each value of A is associated with exactly one value of B. We write this as $A \rightarrow B$, read as "A determines B" or "B depends on A."

Functional Dependency Notation

Understanding the notation helps communicate database design decisions precisely and unambiguously among team members.

Determinant (Left Side)

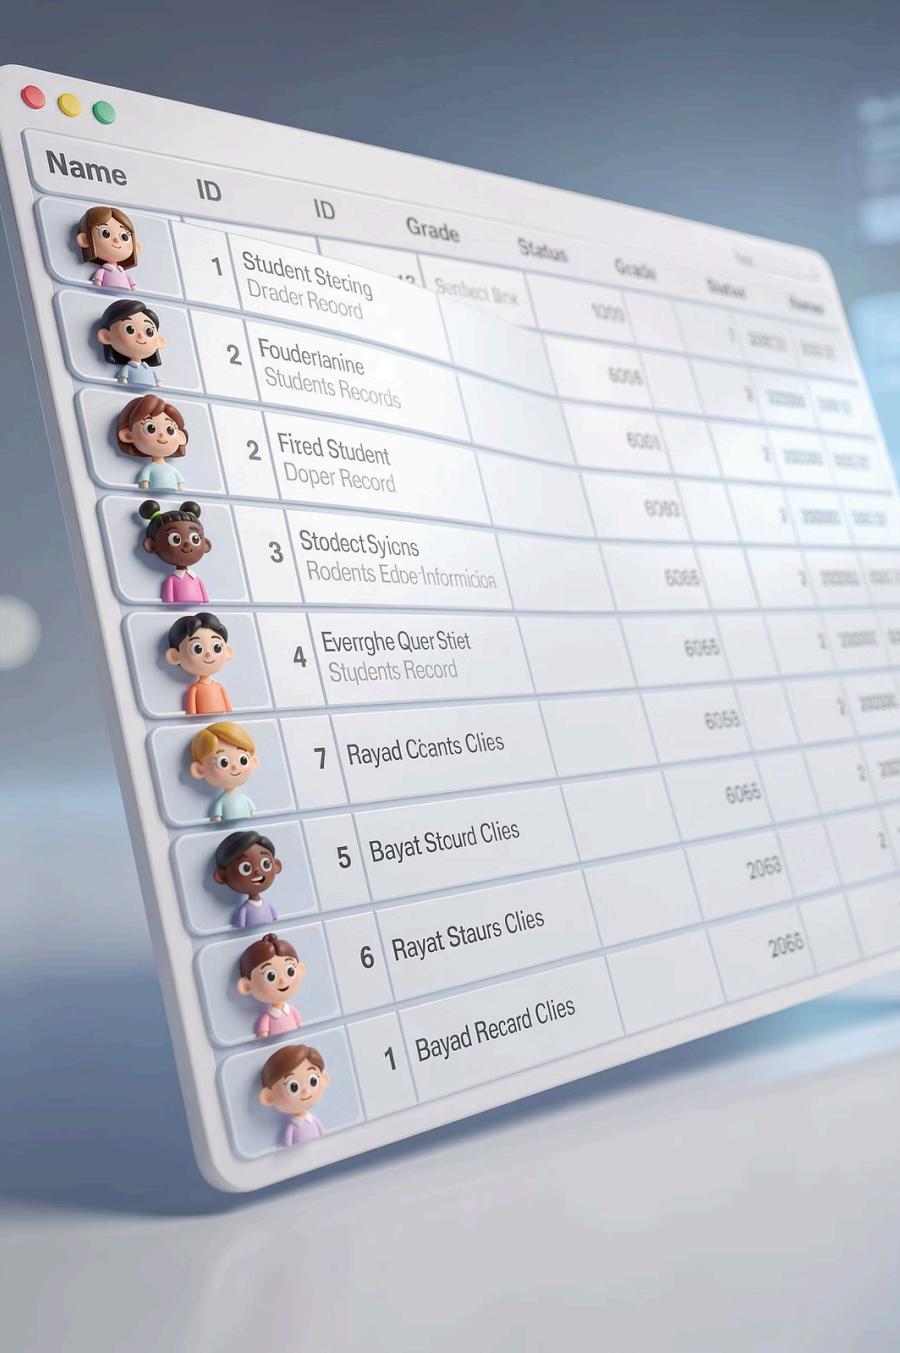
The attribute or set of attributes that determines other attributes. In $\text{StudentID} \rightarrow \text{StudentName}$, StudentID is the determinant.

Dependent (Right Side)

The attribute whose value is determined by the determinant. In $\text{StudentID} \rightarrow \text{StudentName}$, StudentName is dependent.

Arrow Symbol

The \rightarrow symbol represents functional determination, showing the direction of dependency from determinant to dependent.



A digital tablet screen showing a student enrollment database table. The table has columns for Name, ID, Grade, Status, and other details. Each row contains a small cartoon character icon representing a student. The data includes rows for Student Steling Drader Record, Founderianine Students Records, Fired Student Doper Record, StodectSjions Rodents Edbe Informacion, Everrghe Quer Stiet Students Record, Rayad Ccants Clies, Bayat Stourd Clies, Rayat Staurs Clies, and Bayad Recard Clies.

Name	ID	ID	Grade	Status	Grade	Status
Student Steling Drader Record	1	12	Sensect Nine	1079	1	10000 10000
Founderianine Students Records	2			6055	1	10000 10000
Fired Student Doper Record	2		6051		1	10000 10000
StodectSjions Rodents Edbe Informacion	3		6050		1	10000 10000
Everrghe Quer Stiet Students Record	4		6059		1	10000 10000
Rayad Ccants Clies	7		6058		1	10000 10000
Bayat Stourd Clies	5		6055		1	10000 10000
Rayat Staurs Clies	6		2063		1	10000 10000
Bayad Recard Clies	1		2060			

Practical Example: Student Database

Consider a simple student enrollment table with the following attributes: **StudentID**, **StudentName**, **CourseID**, **CourseName**, **InstructorName**.

Several functional dependencies exist in this structure:

- **StudentID → StudentName**: Each student ID uniquely identifies one student name
- **CourseID → CourseName**: Each course ID determines exactly one course name
- **CourseID → InstructorName**: Each course is taught by one specific instructor
- **(StudentID, CourseID) → StudentName, CourseName, InstructorName**: The combination determines all other attributes

Types of Functional Dependencies

Functional dependencies are classified into several categories based on their characteristics. Understanding these types is crucial for applying normalization rules correctly.



Partial Dependency

Dependency on part of a composite key



Full Dependency

Dependency on entire composite key



Transitive Dependency

Indirect dependency through another attribute

4.2.1

Partial Functional Dependency

A partial functional dependency occurs when a non-key attribute depends on only part of a composite primary key, rather than the entire key. This situation violates Second Normal Form and must be eliminated.

Partial dependencies indicate that some data is stored at the wrong level of granularity. They lead to redundancy because the same information is repeated for different combinations of key values.

- ❑ **Key Insight:** Partial dependencies only occur when you have a composite primary key (a key made up of two or more attributes). With a single-attribute key, partial dependencies are impossible.

Partial Dependency Example

Consider an Enrollment table with composite key (StudentID, CourseID):

StudentID	CourseID	StudentName	CourseName
101	CS201	John Smith	Database Systems
101	CS202	John Smith	Networks
102	CS201	Jane Doe	Database Systems

Here, StudentName depends only on StudentID (not the full key), and CourseName depends only on CourseID. These are partial dependencies that cause StudentName and CourseName to be unnecessarily repeated.

Problems Caused by Partial Dependencies



Data Redundancy

Student names and course names are duplicated across multiple enrollment records, wasting storage space.



Update Anomalies

Changing a student's name requires updating multiple records, increasing the risk of inconsistency.



Deletion Anomalies

Removing a student's last enrollment deletes all information about that student from the database.



Insertion Anomalies

Cannot add a student to the database until they enroll in at least one course.

Resolving Partial Dependencies

The solution is to decompose the table into multiple tables, each focused on a single entity. This eliminates redundancy and anomalies.

Students Table

StudentID	StudentName
101	John Smith
102	Jane Doe

Courses Table

CourseID	CourseName
CS201	Database Systems
CS202	Networks

Enrollments Table

StudentID	CourseID
101	CS201
101	CS202
102	CS201

Now each table stores information at the appropriate level, eliminating partial dependencies completely.



4.2.2

Full Functional Dependency

A full functional dependency exists when a non-key attribute depends on the entire composite primary key, not just part of it. This is the desired state after eliminating partial dependencies.

In full functional dependency, removing any attribute from the determinant breaks the dependency. For example, if $(\text{StudentID}, \text{CourseID}) \rightarrow \text{Grade}$, the Grade depends on the specific combination of student and course—neither StudentID alone nor CourseID alone determines the grade.

Full Dependency Characteristics

1

Complete Determinant

The dependent attribute requires all parts of the composite key to be uniquely determined. No subset of the key is sufficient.

2

Proper Granularity

Data is stored at the correct level of detail, matching the natural relationship between entities in the problem domain.

3

No Redundancy

Information appears only once in the database, associated with the full key that uniquely identifies it.

Full Dependency Example

In our enrollment scenario, certain attributes naturally depend on the full composite key:

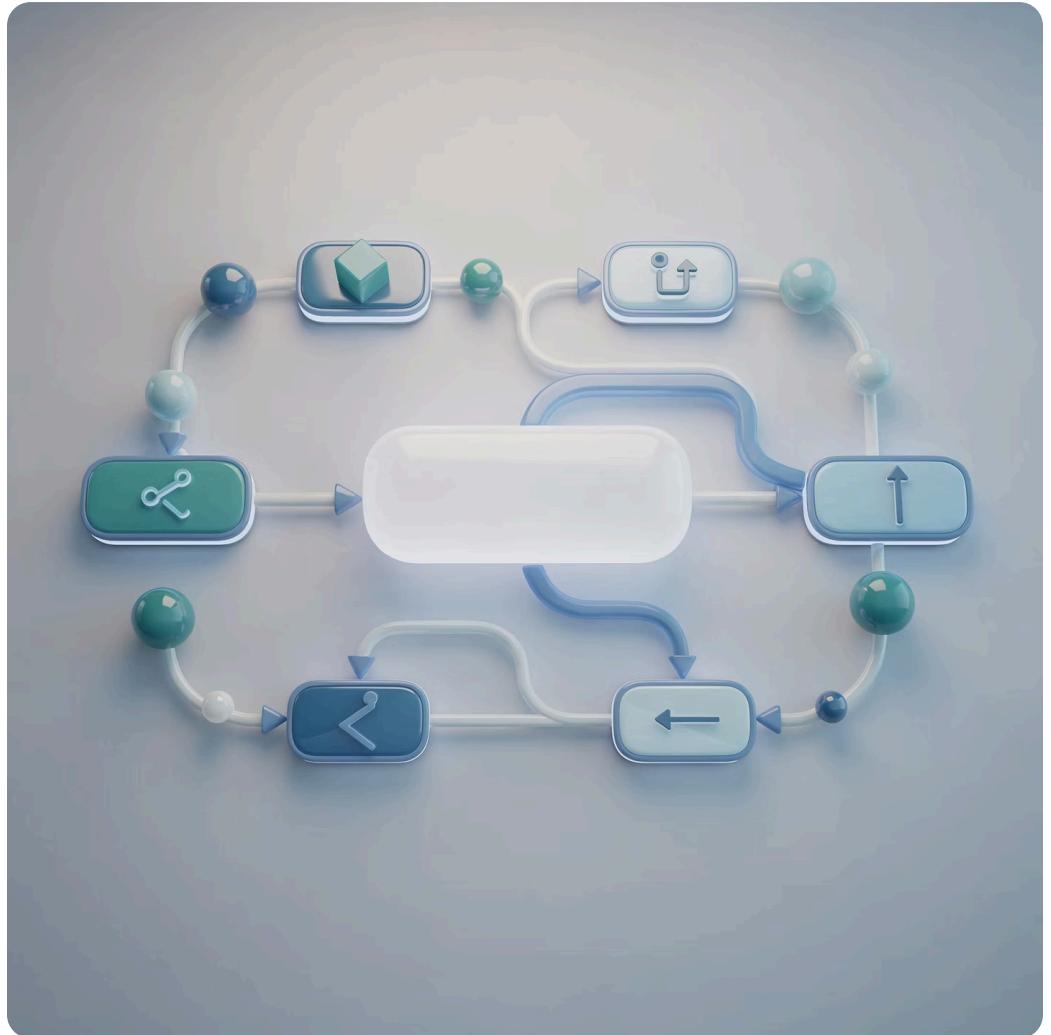
StudentID	CourseID	EnrollmentDate	Grade
101	CS201	2024-01-15	A
101	CS202	2024-01-15	B+
102	CS201	2024-01-18	A-

Both EnrollmentDate and Grade exhibit full functional dependency on (StudentID, CourseID). The enrollment date and grade exist only for a specific student in a specific course—neither attribute makes sense with only StudentID or only CourseID.

Identifying Full Dependencies

To determine if a dependency is full, ask: "Does this attribute make sense without the complete key?"

- **Grade:** Needs both student and course—full dependency
- **StudentName:** Makes sense with just StudentID—partial dependency
- **CourseName:** Makes sense with just CourseID—partial dependency
- **EnrollmentDate:** Needs both student and course—full dependency

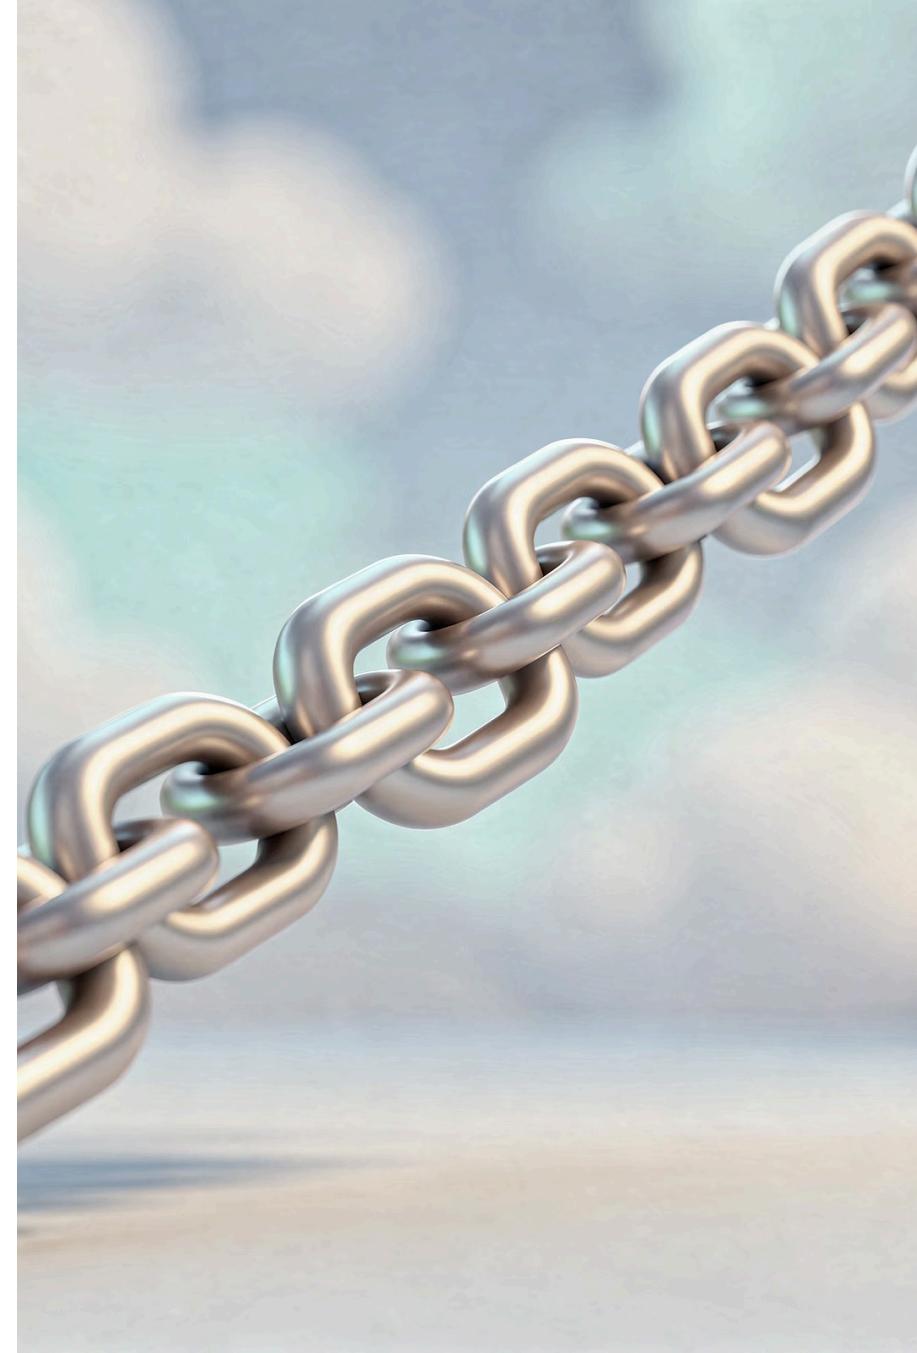


4.2.3

Transitive Dependency

A transitive dependency occurs when a non-key attribute depends on another non-key attribute, which in turn depends on the primary key. This creates an indirect dependency: $A \rightarrow B \rightarrow C$ means C transitively depends on A through B.

Transitive dependencies violate Third Normal Form and introduce redundancy even when partial dependencies have been eliminated. They often represent hierarchical relationships that should be stored in separate tables.



Transitive Dependency Example

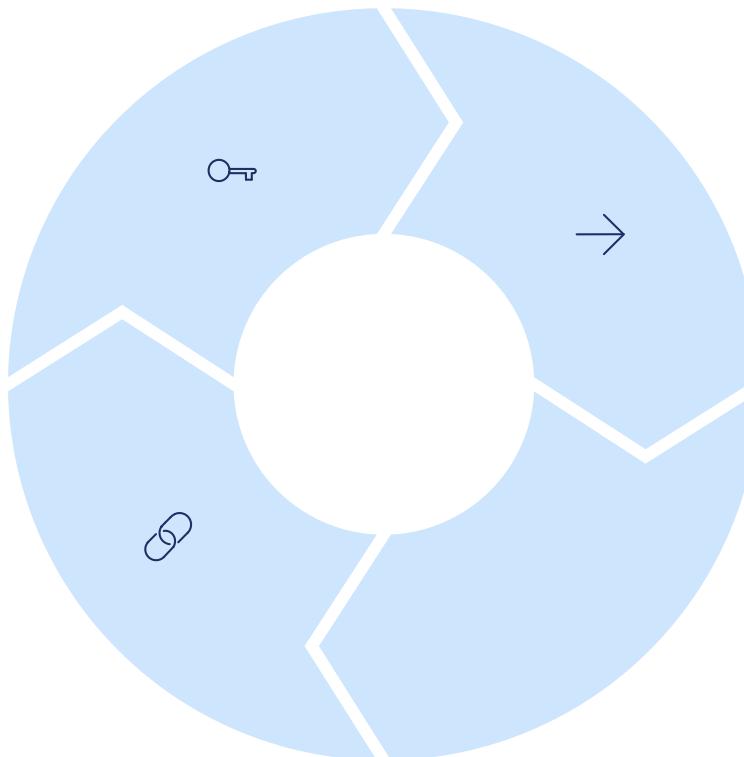
Consider a Student table with advisor information:

StudentID	StudentName	AdvisorID	AdvisorName
101	John Smith	A01	Dr. Anderson
102	Jane Doe	A01	Dr. Anderson
103	Bob Wilson	A02	Dr. Brown

Here, $\text{StudentID} \rightarrow \text{AdvisorID}$ (each student has one advisor), and $\text{AdvisorID} \rightarrow \text{AdvisorName}$ (each advisor has one name). Therefore, AdvisorName transitively depends on StudentID through AdvisorID . Dr. Anderson's name is unnecessarily repeated for each student she advises.

Recognizing Transitive Dependencies

Primary Key
The starting point of the dependency chain



Intermediate Attribute
Non-key attribute that depends on the key

Final Attribute
Non-key attribute that depends on intermediate

Problems from Transitive Dependencies

Before Resolution



After Resolution



- Advisor names repeated for each student
- Updating advisor name requires multiple changes
- Risk of inconsistent advisor information
- Cannot store advisors without students
- Wasted storage space

- Each advisor stored once
- Single-point updates
- Guaranteed consistency
- Advisors independent of students
- Efficient storage utilization

Resolving Transitive Dependencies

The solution is to separate the transitively dependent attributes into a new table:

Student Table

StudentID	StudentName	AdvisorID
101	John Smith	A01
102	Jane Doe	A01
103	Bob Wilson	A02

Students now reference advisors through AdvisorID, eliminating the transitive dependency.

Advisor Table

AdvisorID	AdvisorName
A01	Dr. Anderson
A02	Dr. Brown

Advisor information exists in one place, referenced by students as needed.

Normal Forms

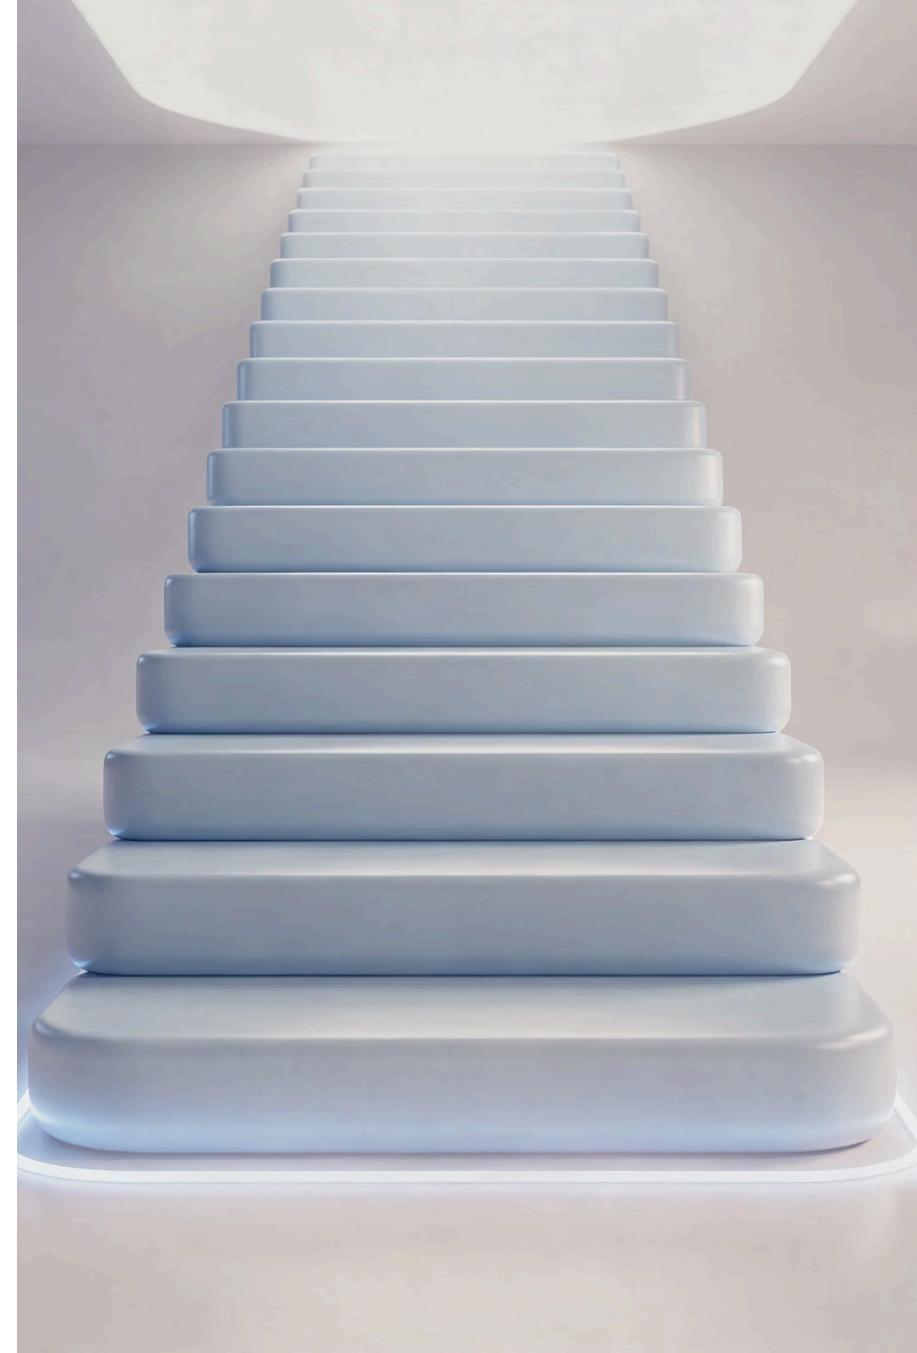
Systematic levels of database organization

◇ SECTION 4.3

Introduction to Normal Forms

Normal forms are progressive levels of database organization, each building upon the previous one. Each normal form addresses specific types of redundancy and dependency problems, creating increasingly refined database structures.

The normalization process typically proceeds through First Normal Form (1NF), Second Normal Form (2NF), and Third Normal Form (3NF). Higher normal forms exist but are less commonly used in practical database design. Most commercial databases target 3NF as the optimal balance between normalization benefits and practical usability.



The Normal Form Hierarchy

1

First Normal Form (1NF)

Eliminates repeating groups and ensures atomic values in all columns

2

Second Normal Form (2NF)

Achieves 1NF and removes partial dependencies on composite keys



Third Normal Form (3NF)

Achieves 2NF and eliminates transitive dependencies between non-key attributes

Each level requires satisfying all previous levels plus additional conditions. You cannot skip levels in the normalization process.

Normal Form Prerequisites



1NF Requirements

1

Atomic values only, no repeating groups

2NF Requirements

2

Must be in 1NF + no partial dependencies

3NF Requirements

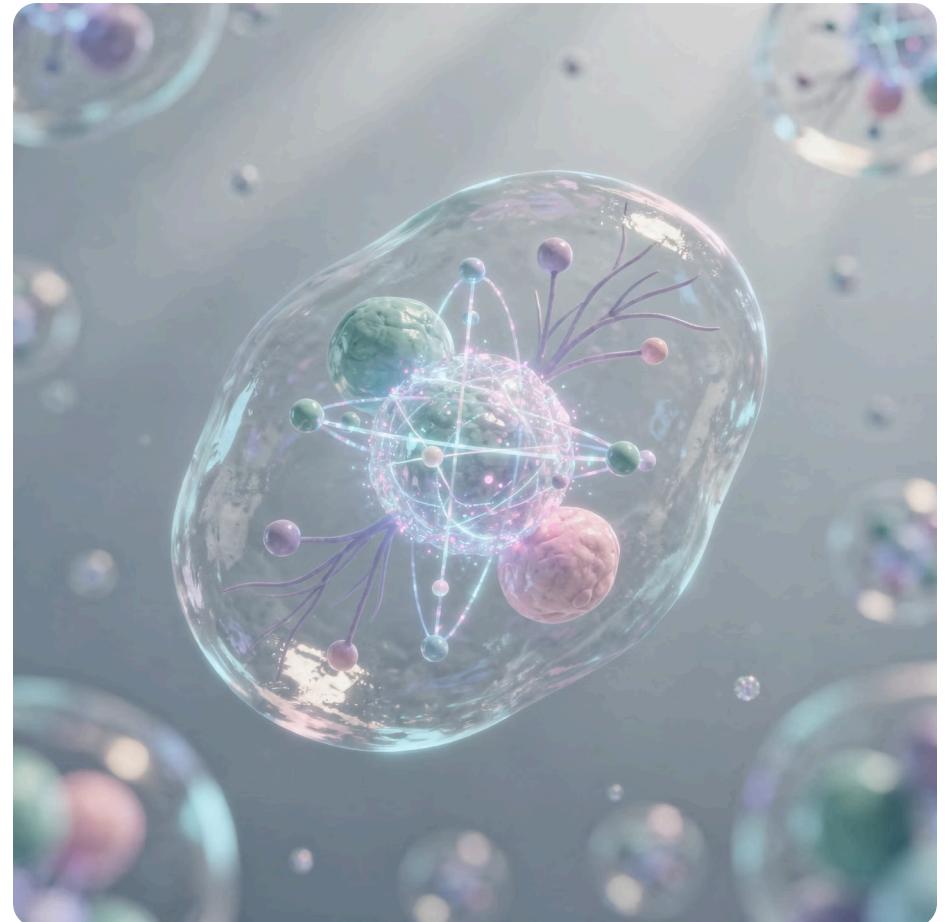
3

Must be in 2NF + no transitive dependencies

First Normal Form (1NF)

First Normal Form is the foundation of all normalization. A table is in 1NF if it meets two critical requirements: every column contains atomic (indivisible) values, and there are no repeating groups of columns.

Atomicity means each cell contains a single value, not a list, set, or composite value. Repeating groups occur when similar information is stored in multiple columns (like Phone1, Phone2, Phone3).



1NF Violation: Non-Atomic Values

This table violates 1NF because the PhoneNumbers column contains multiple values:

StudentID	StudentName	PhoneNumbers
101	John Smith	555-1234, 555-5678
102	Jane Doe	555-8765, 555-4321, 555-1111

- **Problem:** How do you search for a specific phone number? How do you update just one number? How do you know how many phone numbers a student has? These operations become unnecessarily complex.

1NF Violation: Repeating Groups

This table violates 1NF through repeating column groups:

StudentID	Name	Course1	Grade1	Course2	Grade2
101	John	CS201	A	CS202	B
102	Jane	CS201	A-		

This design limits students to a fixed number of courses and wastes space when students take fewer courses. Adding a third course requires schema changes.

Converting to 1NF: Solution 1

Create separate rows for each phone number:

Before (Violates 1NF)

StudentID	Name	Phones
101	John	555-1234, 555-5678

After (Meets 1NF)

StudentID	Name	Phone
101	John	555-1234
101	John	555-5678

Now each cell contains exactly one value. The composite key (StudentID, Phone) uniquely identifies each row.

Converting to 1NF: Solution 2

For repeating groups like courses, create a separate enrollment table:

Student Table

StudentID	StudentName
101	John Smith
102	Jane Doe

Enrollment Table

StudentID	CourseID	Grade
101	CS201	A
101	CS202	B
102	CS201	A-

This structure supports any number of courses per student without schema changes.

1NF Benefits and Characteristics



Atomic Values

Each column contains indivisible data elements that cannot be meaningfully broken down further



Uniform Structure

All rows have the same structure with consistent column types throughout the table



Queryable Data

Simple SQL operations can search, filter, and aggregate data reliably and efficiently



Unlimited Flexibility

No artificial limits on the number of related items each record can have



4.3.2

Second Normal Form (2NF)

A table is in Second Normal Form if it is in First Normal Form and all non-key attributes are fully functionally dependent on the entire primary key. This means eliminating partial dependencies that exist when using composite keys.

2NF only applies to tables with composite primary keys. If your table has a single-column primary key, it automatically satisfies 2NF (assuming it already meets 1NF). The focus is on ensuring that every piece of data depends on the complete key, not just part of it.

2NF Violation Example

Consider this enrollment table with partial dependencies:

StudentID	CourseID	StudentName	CourseName	Grade
101	CS201	John Smith	Database Systems	A
101	CS202	John Smith	Computer Networks	B+
102	CS201	Jane Doe	Database Systems	A-

The composite key is (StudentID, CourseID). However, StudentName depends only on StudentID, and CourseName depends only on CourseID. These are partial dependencies that violate 2NF.

Identifying 2NF Violations

01

Identify the Primary Key

Determine if the table has a composite primary key. If not, 2NF is automatically satisfied.

03

Test Dependency Completeness

Ask: Does this attribute require the entire key, or just part of it?

02

Examine Each Non-Key Attribute

For every attribute that's not part of the key, determine what it depends on.

04

Flag Partial Dependencies

Any attribute depending on only part of the key violates 2NF.

Converting to 2NF

Decompose the table into multiple tables where each non-key attribute depends on the complete primary key:

Student Table

StudentID	StudentName
101	John Smith
102	Jane Doe

StudentName now depends on the complete key (StudentID only).

Course Table

CourseID	CourseName
CS201	Database Systems
CS202	Computer Networks

CourseName depends on the complete key (CourseID only).

Enrollment Table

StudentID	CourseID	Grade
101	CS201	A
101	CS202	B+
102	CS201	A-

Grade depends on the complete composite key.

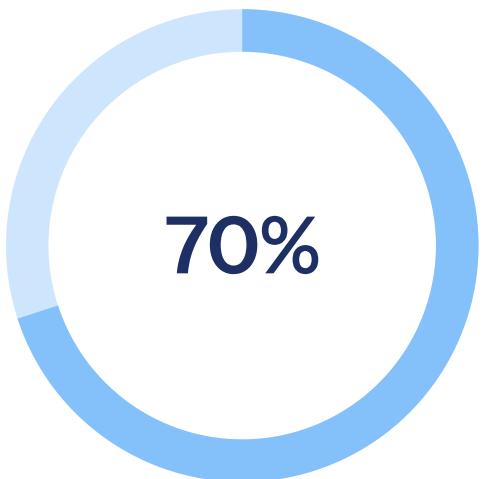


2NF Decomposition Strategy

Follow this systematic approach to achieve Second Normal Form:

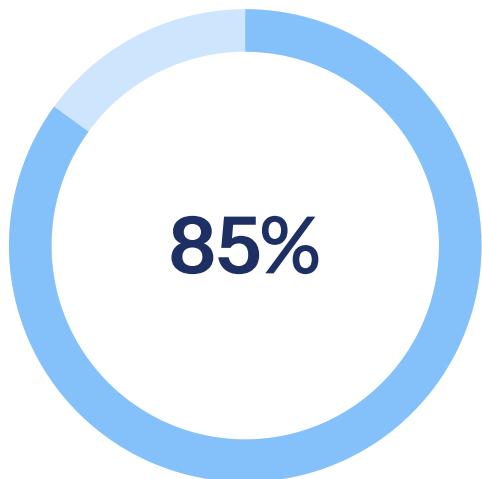
1. **Identify partial dependencies:** Find all non-key attributes that depend on only part of the composite key
2. **Group by determinant:** Collect attributes that share the same partial key dependency
3. **Create new tables:** Make a separate table for each group with its determinant as the primary key
4. **Maintain relationships:** Keep foreign keys in the original table to preserve connections
5. **Verify completeness:** Ensure the decomposition preserves all original data and relationships

2NF Impact on Database Quality



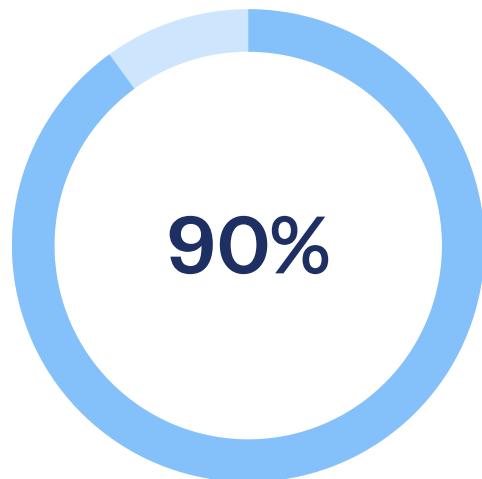
Redundancy Reduction

Typical decrease in duplicate data after achieving 2NF



Update Efficiency

Improvement in modification operation simplicity



Data Consistency

Increase in maintaining accurate information

Third Normal Form (3NF)



A table is in Third Normal Form if it is in Second Normal Form and no non-key attribute depends on another non-key attribute. This eliminates transitive dependencies where data depends indirectly on the primary key through another attribute.

3NF represents the standard target for most commercial database designs, providing an excellent balance between normalization benefits and practical usability without excessive complexity.

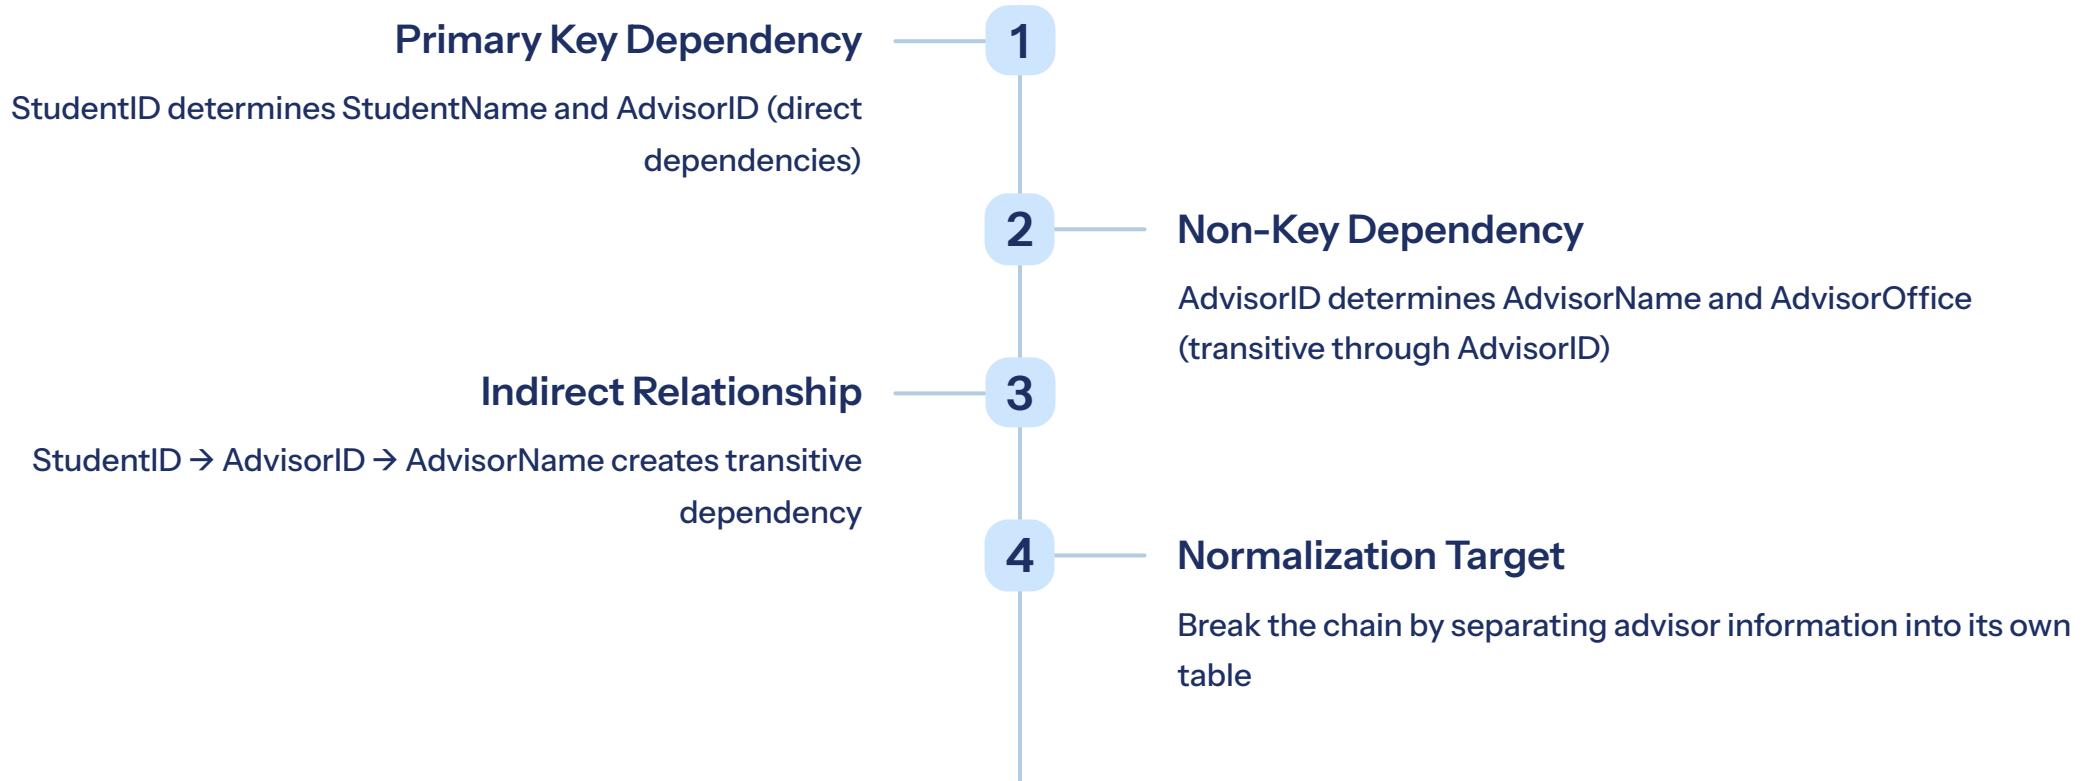
3NF Violation Example

Consider a Student table with advisor information:

StudentID	StudentName	AdvisorID	AdvisorName	AdvisorOffice
101	John Smith	A01	Dr. Anderson	Room 301
102	Jane Doe	A01	Dr. Anderson	Room 301
103	Bob Wilson	A02	Dr. Brown	Room 305

While this table is in 2NF (single-column key), it violates 3NF. AdvisorName and AdvisorOffice depend on AdvisorID, which is not the primary key. These transitive dependencies cause redundancy.

Recognizing Transitive Dependencies in 3NF



Converting to 3NF

Separate the transitively dependent attributes into a new table:

Student Table (No Transitive Dependencies)

StudentID	StudentName	AdvisorID
101	John Smith	A01
102	Jane Doe	A01
103	Bob Wilson	A02

Students reference advisors through AdvisorID foreign key.

Advisor Table

AdvisorID	Name	Office
A01	Dr. Anderson	Rm 301
A02	Dr. Brown	Rm 305

Advisor details stored once, eliminating redundancy.

3NF Decomposition Process

1

Identify Transitive Dependencies

Find non-key attributes that depend on other non-key attributes rather than directly on the primary key

2

Create Independent Tables

Extract each group of transitively dependent attributes into a new table with appropriate primary key

3

Establish Foreign Keys

Maintain relationships between tables using foreign key references to connect related data

4

Verify 3NF Compliance

Ensure all non-key attributes depend directly and only on primary keys in each table

Benefits of Achieving 3NF



Minimal Redundancy

Data stored in single locations eliminates duplication and wasted space. Each fact appears once and only once in the database structure.



Data Integrity

Consistent information guaranteed through proper table structure. Updates affect single records, preventing contradictory data.



Update Efficiency

Modifications touch minimal records, improving performance. Single-point updates reduce processing overhead significantly.

Easier Maintenance

Logical organization simplifies troubleshooting and enhancement. Clear relationships between tables aid understanding.

Complete Normalization Example

Let's trace a table through all three normal forms to see the complete transformation:

Original Unnormalized Table

OrderID	Customer	Products	Total
1001	John Smith	Laptop, Mouse, Keyboard	\$1,250
1002	Jane Doe	Monitor, Cable	\$450

Problems: Multiple products in one field (not atomic), repeating customer data, no clear relationship structure.

After 1NF: Atomic Values

OrderID	CustomerName	CustomerEmail	ProductID	ProductName	Price
1001	John Smith	john@email.com	P01	Laptop	\$1000
1001	John Smith	john@email.com	P02	Mouse	\$25
1001	John Smith	john@email.com	P03	Keyboard	\$75

Each cell contains atomic values, but significant redundancy remains with partial and transitive dependencies still present.

After 2NF: Removing Partial Dependencies

Orders Table

OrderID	CustomerID	OrderDate
1001	C01	2024-01-15
1002	C02	2024-01-16

Products Table

ProductID	Name	Price
P01	Laptop	\$1000
P02	Mouse	\$25

OrderDetails Table

OrderID	ProductID	Quantity
1001	P01	1
1001	P02	1

Customers Table

CustomerID	Name	Email
C01	John Smith	john@email.com

Customer and product data no longer repeat for each order line. However, transitive dependencies may still exist.

After 3NF: Final Normalized Structure

If customers have addresses stored at the city level, separate location information:

Customers

Customer ID	Name	CityID
C01	John Smith	CT01
C02	Jane Doe	CT02

Cities

CityID	CityName	State
CT01	Mumbai	MH
CT02	Delhi	DL

Now State depends directly on CityID (its primary key), not transitively through Customer. All tables are in 3NF.

Comparing Normal Forms: Summary

First Normal Form (1NF)

- Eliminates repeating groups
- Ensures atomic values in all columns
- Creates uniform row structure
- Foundation for further normalization

Second Normal Form (2NF)

- Builds on 1NF requirements
- Removes partial dependencies
- Applies to composite key tables
- Reduces redundancy significantly

Third Normal Form (3NF)

- Builds on 2NF requirements
- Eliminates transitive dependencies
- Creates fully normalized structure
- Industry standard for most databases

When to Stop Normalizing

While higher normal forms exist (BCNF, 4NF, 5NF), most practical databases stop at Third Normal Form. The decision involves balancing theoretical purity against practical considerations.

Reasons to Normalize Further

- Complex data relationships
- Critical data integrity requirements
- Frequent update operations
- Academic or theoretical projects

Reasons to Stop at 3NF

- Read-heavy applications
- Performance concerns with many joins
- Simpler application code
- Industry standard practice



Practical Normalization Guidelines

Start with Requirements

Understand your data relationships and business rules thoroughly before beginning normalization. Document all functional dependencies clearly.

Apply Forms Sequentially

Always progress through normal forms in order: 1NF, then 2NF, then 3NF. Each level builds on the previous one—you cannot skip steps.

Test Your Design

Verify the normalized structure handles all required operations correctly. Check for completeness and performance with realistic data volumes.

Document Decisions

Record why you made specific normalization choices. This helps future developers understand and maintain the database structure.

Common Normalization Mistakes to Avoid

Over-Normalization

Creating too many tables with excessive joins that harm query performance. Balance normalization benefits against practical usability needs.

Incomplete Analysis

Missing functional dependencies during initial analysis. Thoroughly document all relationships before starting the normalization process.

Skipping Levels

Jumping directly to 3NF without ensuring 1NF and 2NF. Each normal form builds on previous ones and must be satisfied sequentially.

Ignoring Performance

Normalizing without considering query patterns. Some denormalization may be appropriate for read-heavy applications with performance requirements.

Key Takeaways: Unit IV Summary

Normalization Purpose

Systematic technique to eliminate redundancy, prevent anomalies, and ensure data integrity through progressive refinement

Functional Dependencies

Understanding partial, full, and transitive dependencies is essential for identifying normalization opportunities

Three Normal Forms

1NF ensures atomic values, 2NF eliminates partial dependencies, 3NF removes transitive dependencies

Practical Application

Most commercial databases target 3NF as the optimal balance between theory and real-world performance

Thank You

You've completed Unit IV: Refining Database Design Through Normalization. You now understand how to transform poorly designed databases into efficient, maintainable structures that support robust application development.

"Good database design is invisible to users but invaluable to developers.
Normalization is the discipline that makes both possible."

UNIT COMPLETE

READY FOR ASSESSMENT

